

# Humanoid Final Project

December 2, 2025

## 1 Intro To AI Final Project

### 1.1 Libraries from Open AI Gymnasium

```
[1]: import gymnasium as gym #gym environment so I can run the humanoid environment
import numpy as np
from stable_baselines3 import PPO #built in reinforcement learning model
from stable_baselines3.common.vec_env import DummyVecEnv #additional library
    ↳for PPO to work correctly
from stable_baselines3.common.monitor import Monitor #keeps track of the
    ↳rewards and lengths for training
import imageio #to save a video of the humanoid
```

### 1.2 Random Agent

```
[2]: def run_random_agent(num_episodes=5, max_steps=1000):
    env = gym.make("Humanoid-v5") #creates the humanoid environment from
    ↳openAI gym
    episode_rewards = [] #these listses store the results of how the random
    ↳agent does each turn through
    episode_lengths = []

    for ep in range(num_episodes): #runs the random agent amount of episodes
    ↳(aka attempts) we set
        obs, info = env.reset() #sets the environment to the starting state
        total_reward = 0.0 #keeps track of the agents points
        length = 0 #counts the amound of steps the agent succeeded

        for t in range(max_steps): #loops through each step of the episodes
            action = env.action_space.sample() #picks a randome action from
            ↳the action space

            result = env.step(action) #stores the information for what happens
            ↳each step
            obs = result[0] #the state of the humanoid after doing the action
            reward = result[1] #the amount of points the action earned
            terminated = result[2] #true if the humanoid dint fal
```

```

        truncated = result[3]  #true if the episode hit the max steps
        info = result[4]

        total_reward += reward  #adding reward from this step to total
        ↪reward

        length += 1  #adds a count to the amount of steps taken

        if terminated or truncated: #stop if the humanoid runs out of time
        ↪or fails
            break

        episode_rewards.append(total_reward)  #stores total rewards for this epi
        episode_lengths.append(length)  #stores how long the epi lasted

    env.close()  #closes the humanoid environment

    print("\nRandom agent summary:")  #prints the average of the random attempts
    print(f"    Mean reward: {np.mean(episode_rewards):.2f}")
    print(f"    Mean length: {np.mean(episode_lengths):.1f}")

run_random_agent()

```

Random agent summary:

Mean reward: 114.43

Mean length: 24.4

### 1.3 Humanoid Environment using PPO

```

[3]: def make_env():
    def _init():
        env = gym.make("Humanoid-v5")  #creates the humanoid environment
        env = Monitor(env)  #keeps track of the episode rewards & lengths
        return env
    return _init

vec_env = DummyVecEnv([make_env()])  #using a specific environment bc "PPO is
    ↪meant to be run primarily on the CPU, especially when you are not using a
    ↪CNN" from resource 3

model = PPO(  #proximal policy optimization model from stable baselines
    "MlpPolicy", #using a mlp neural network as its the default for the
    ↪gymnasium humanoid
    vec_env,
    verbose=1,

```

```

    #these are the parameters for this model, haven't gotten into
    ↪hyperparameter tuning much so will stick with the default
    n_steps=2048,
    batch_size=64,
    learning_rate=3e-4,
    gamma=0.99,
)

model.learn(total_timesteps=1_500_000) #this is how we control hoow many
    ↪training steps the model takes, more is better for the goal
model.save("ppo_humanoid_v5") #saves the trained humanoid model

```

Using cpu device

```

-----
| rollout/          |          |
|   ep_len_mean    | 21.7     |
|   ep_rew_mean    | 96       |
| time/            |          |
|   fps            | 577      |
|   iterations     | 1        |
|   time_elapsed   | 3        |
|   total_timesteps | 2048     |
-----

```

```

-----
| rollout/          |          |
|   ep_len_mean    | 22.9     |
|   ep_rew_mean    | 102      |
| time/            |          |
|   fps            | 503      |
|   iterations     | 2        |
|   time_elapsed   | 8        |
|   total_timesteps | 4096     |
| train/           |          |
|   approx_kl      | 0.020275684 |
|   clip_fraction  | 0.237    |
|   clip_range     | 0.2      |
|   entropy_loss   | -24.1    |
|   explained_variance | -0.00268 |
|   learning_rate  | 0.0003   |
|   loss           | 411      |
|   n_updates      | 10       |
|   policy_gradient_loss | -0.0626  |
|   std            | 0.998    |
|   value_loss     | 1.21e+03 |
-----

```

```

-----
| rollout/          |          |

```

	ep_len_mean		22.9	
	ep_rew_mean		102	
	time/			
	fps		495	
	iterations		3	
	time_elapsed		12	
	total_timesteps		6144	
	train/			
	approx_kl		0.02417365	
	clip_fraction		0.25	
	clip_range		0.2	
	entropy_loss		-24.1	
	explained_variance		0.0149	
	learning_rate		0.0003	
	loss		413	
	n_updates		20	
	policy_gradient_loss		-0.0672	
	std		0.996	
	value_loss		971	

	rollout/			
	ep_len_mean		21.9	
	ep_rew_mean		97.4	
	time/			
	fps		494	
	iterations		4	
	time_elapsed		16	
	total_timesteps		8192	
	train/			
	approx_kl		0.028060088	
	clip_fraction		0.318	
	clip_range		0.2	
	entropy_loss		-24	
	explained_variance		0.019	
	learning_rate		0.0003	
	loss		310	
	n_updates		30	
	policy_gradient_loss		-0.0803	
	std		0.992	
	value_loss		855	

	rollout/			
	ep_len_mean		24.3	
	ep_rew_mean		109	
	time/			
	fps		486	

iterations	5
time_elapsed	21
total_timesteps	10240
train/	
approx_kl	0.0362327
clip_fraction	0.367
clip_range	0.2
entropy_loss	-24
explained_variance	0.0296
learning_rate	0.0003
loss	307
n_updates	40
policy_gradient_loss	-0.0901
std	0.992
value_loss	716

rollout/	
ep_len_mean	25.8
ep_rew_mean	116
time/	
fps	480
iterations	6
time_elapsed	25
total_timesteps	12288
train/	
approx_kl	0.03207309
clip_fraction	0.331
clip_range	0.2
entropy_loss	-23.9
explained_variance	0.0224
learning_rate	0.0003
loss	273
n_updates	50
policy_gradient_loss	-0.0802
std	0.989
value_loss	910

rollout/	
ep_len_mean	25.9
ep_rew_mean	116
time/	
fps	482
iterations	7
time_elapsed	29
total_timesteps	14336
train/	

approx_kl	0.028811298	
clip_fraction	0.306	
clip_range	0.2	
entropy_loss	-23.9	
explained_variance	0.0327	
learning_rate	0.0003	
loss	466	
n_updates	60	
policy_gradient_loss	-0.0802	
std	0.987	
value_loss	939	

rollout/		
ep_len_mean	26.7	
ep_rew_mean	121	
time/		
fps	484	
iterations	8	
time_elapsed	33	
total_timesteps	16384	
train/		
approx_kl	0.02825661	
clip_fraction	0.306	
clip_range	0.2	
entropy_loss	-23.9	
explained_variance	0.0559	
learning_rate	0.0003	
loss	435	
n_updates	70	
policy_gradient_loss	-0.0817	
std	0.99	
value_loss	871	

rollout/		
ep_len_mean	29.5	
ep_rew_mean	133	
time/		
fps	483	
iterations	9	
time_elapsed	38	
total_timesteps	18432	
train/		
approx_kl	0.02437363	
clip_fraction	0.252	
clip_range	0.2	
entropy_loss	-23.9	

	explained_variance	0.034	
	learning_rate	0.0003	
	loss	441	
	n_updates	80	
	policy_gradient_loss	-0.0753	
	std	0.988	
	value_loss	1.06e+03	

	rollout/		
	ep_len_mean	29.4	
	ep_rew_mean	133	
	time/		
	fps	486	
	iterations	10	
	time_elapsed	42	
	total_timesteps	20480	
	train/		
	approx_kl	0.021425027	
	clip_fraction	0.221	
	clip_range	0.2	
	entropy_loss	-23.9	
	explained_variance	0.0802	
	learning_rate	0.0003	
	loss	534	
	n_updates	90	
	policy_gradient_loss	-0.0704	
	std	0.989	
	value_loss	1.04e+03	

	rollout/		
	ep_len_mean	29.8	
	ep_rew_mean	134	
	time/		
	fps	489	
	iterations	11	
	time_elapsed	45	
	total_timesteps	22528	
	train/		
	approx_kl	0.022110213	
	clip_fraction	0.217	
	clip_range	0.2	
	entropy_loss	-23.9	
	explained_variance	0.121	
	learning_rate	0.0003	
	loss	607	
	n_updates	100	

	policy_gradient_loss		-0.0695	
	std		0.988	
	value_loss		1.08e+03	

	rollout/			
	ep_len_mean		30.3	
	ep_rew_mean		138	
	time/			
	fps		485	
	iterations		12	
	time_elapsed		50	
	total_timesteps		24576	
	train/			
	approx_kl		0.018699765	
	clip_fraction		0.193	
	clip_range		0.2	
	entropy_loss		-23.9	
	explained_variance		0.172	
	learning_rate		0.0003	
	loss		457	
	n_updates		110	
	policy_gradient_loss		-0.0663	
	std		0.985	
	value_loss		914	

	rollout/			
	ep_len_mean		31.5	
	ep_rew_mean		145	
	time/			
	fps		485	
	iterations		13	
	time_elapsed		54	
	total_timesteps		26624	
	train/			
	approx_kl		0.020765267	
	clip_fraction		0.201	
	clip_range		0.2	
	entropy_loss		-23.9	
	explained_variance		0.288	
	learning_rate		0.0003	
	loss		476	
	n_updates		120	
	policy_gradient_loss		-0.0669	
	std		0.985	
	value_loss		945	



-----		
rollout/		
ep_len_mean	32.9	
ep_rew_mean	150	
time/		
fps	486	
iterations	14	
time_elapsed	58	
total_timesteps	28672	
train/		
approx_kl	0.02377629	
clip_fraction	0.225	
clip_range	0.2	
entropy_loss	-23.8	
explained_variance	0.336	
learning_rate	0.0003	
loss	509	
n_updates	130	
policy_gradient_loss	-0.0705	
std	0.983	
value_loss	1.05e+03	
-----		

-----		
rollout/		
ep_len_mean	33.8	
ep_rew_mean	153	
time/		
fps	486	
iterations	15	
time_elapsed	63	
total_timesteps	30720	
train/		
approx_kl	0.020942189	
clip_fraction	0.199	
clip_range	0.2	
entropy_loss	-23.8	
explained_variance	0.275	
learning_rate	0.0003	
loss	507	
n_updates	140	
policy_gradient_loss	-0.0648	
std	0.981	
value_loss	1.07e+03	
-----		

-----		
rollout/		
ep_len_mean	34.9	
ep_rew_mean	159	

time/		
fps	486	
iterations	16	
time_elapsed	67	
total_timesteps	32768	
train/		
approx_kl	0.023075894	
clip_fraction	0.223	
clip_range	0.2	
entropy_loss	-23.8	
explained_variance	0.319	
learning_rate	0.0003	
loss	480	
n_updates	150	
policy_gradient_loss	-0.0706	
std	0.98	
value_loss	1.06e+03	

---

rollout/		
ep_len_mean	34.8	
ep_rew_mean	159	
time/		
fps	485	
iterations	17	
time_elapsed	71	
total_timesteps	34816	
train/		
approx_kl	0.021670304	
clip_fraction	0.202	
clip_range	0.2	
entropy_loss	-23.8	
explained_variance	0.332	
learning_rate	0.0003	
loss	511	
n_updates	160	
policy_gradient_loss	-0.0655	
std	0.979	
value_loss	1.09e+03	

---

rollout/		
ep_len_mean	35.6	
ep_rew_mean	161	
time/		
fps	483	
iterations	18	
time_elapsed	76	

	total_timesteps		36864	
	train/			
	approx_kl		0.019386131	
	clip_fraction		0.18	
	clip_range		0.2	
	entropy_loss		-23.7	
	explained_variance		0.363	
	learning_rate		0.0003	
	loss		454	
	n_updates		170	
	policy_gradient_loss		-0.0631	
	std		0.976	
	value_loss		975	

	rollout/			
	ep_len_mean		39	
	ep_rew_mean		178	
	time/			
	fps		481	
	iterations		19	
	time_elapsed		80	
	total_timesteps		38912	
	train/			
	approx_kl		0.021443568	
	clip_fraction		0.194	
	clip_range		0.2	
	entropy_loss		-23.7	
	explained_variance		0.338	
	learning_rate		0.0003	
	loss		705	
	n_updates		180	
	policy_gradient_loss		-0.0647	
	std		0.975	
	value_loss		1.11e+03	

	rollout/			
	ep_len_mean		42	
	ep_rew_mean		193	
	time/			
	fps		481	
	iterations		20	
	time_elapsed		85	
	total_timesteps		40960	
	train/			
	approx_kl		0.019670099	
	clip_fraction		0.18	

clip_range	0.2	
entropy_loss	-23.7	
explained_variance	0.352	
learning_rate	0.0003	
loss	621	
n_updates	190	
policy_gradient_loss	-0.0608	
std	0.975	
value_loss	1.26e+03	

rollout/		
ep_len_mean	43.4	
ep_rew_mean	200	
time/		
fps	479	
iterations	21	
time_elapsed	89	
total_timesteps	43008	
train/		
approx_kl	0.020847611	
clip_fraction	0.203	
clip_range	0.2	
entropy_loss	-23.7	
explained_variance	0.343	
learning_rate	0.0003	
loss	588	
n_updates	200	
policy_gradient_loss	-0.0667	
std	0.973	
value_loss	1.21e+03	

rollout/		
ep_len_mean	44.3	
ep_rew_mean	204	
time/		
fps	478	
iterations	22	
time_elapsed	94	
total_timesteps	45056	
train/		
approx_kl	0.022172535	
clip_fraction	0.2	
clip_range	0.2	
entropy_loss	-23.6	
explained_variance	0.367	
learning_rate	0.0003	

	loss		691	
	n_updates		210	
	policy_gradient_loss		-0.063	
	std		0.971	
	value_loss		1.31e+03	

	rollout/			
	ep_len_mean		45.4	
	ep_rew_mean		208	
	time/			
	fps		478	
	iterations		23	
	time_elapsed		98	
	total_timesteps		47104	
	train/			
	approx_kl		0.021648861	
	clip_fraction		0.202	
	clip_range		0.2	
	entropy_loss		-23.6	
	explained_variance		0.356	
	learning_rate		0.0003	
	loss		512	
	n_updates		220	
	policy_gradient_loss		-0.0642	
	std		0.971	
	value_loss		1.2e+03	

	rollout/			
	ep_len_mean		46.8	
	ep_rew_mean		214	
	time/			
	fps		478	
	iterations		24	
	time_elapsed		102	
	total_timesteps		49152	
	train/			
	approx_kl		0.020830233	
	clip_fraction		0.204	
	clip_range		0.2	
	entropy_loss		-23.6	
	explained_variance		0.427	
	learning_rate		0.0003	
	loss		489	
	n_updates		230	
	policy_gradient_loss		-0.0636	
	std		0.968	

	value_loss		1.25e+03	
-----				
	rollout/			
	ep_len_mean		45.6	
	ep_rew_mean		208	
	time/			
	fps		479	
	iterations		25	
	time_elapsed		106	
	total_timesteps		51200	
	train/			
	approx_kl		0.021650236	
	clip_fraction		0.202	
	clip_range		0.2	
	entropy_loss		-23.5	
	explained_variance		0.364	
	learning_rate		0.0003	
	loss		536	
	n_updates		240	
	policy_gradient_loss		-0.0649	
	std		0.966	
	value_loss		1.17e+03	
-----				
	rollout/			
	ep_len_mean		45.5	
	ep_rew_mean		210	
	time/			
	fps		477	
	iterations		26	
	time_elapsed		111	
	total_timesteps		53248	
	train/			
	approx_kl		0.022814177	
	clip_fraction		0.201	
	clip_range		0.2	
	entropy_loss		-23.5	
	explained_variance		0.425	
	learning_rate		0.0003	
	loss		455	
	n_updates		250	
	policy_gradient_loss		-0.0615	
	std		0.966	
	value_loss		1.11e+03	
-----				
	rollout/			

	ep_len_mean		47.1	
	ep_rew_mean		219	
	time/			
	fps		478	
	iterations		27	
	time_elapsed		115	
	total_timesteps		55296	
	train/			
	approx_kl		0.022531692	
	clip_fraction		0.204	
	clip_range		0.2	
	entropy_loss		-23.5	
	explained_variance		0.454	
	learning_rate		0.0003	
	loss		554	
	n_updates		260	
	policy_gradient_loss		-0.0614	
	std		0.965	
	value_loss		1.11e+03	

	rollout/			
	ep_len_mean		49.3	
	ep_rew_mean		229	
	time/			
	fps		477	
	iterations		28	
	time_elapsed		120	
	total_timesteps		57344	
	train/			
	approx_kl		0.023429487	
	clip_fraction		0.202	
	clip_range		0.2	
	entropy_loss		-23.5	
	explained_variance		0.452	
	learning_rate		0.0003	
	loss		570	
	n_updates		270	
	policy_gradient_loss		-0.062	
	std		0.964	
	value_loss		1.23e+03	

	rollout/			
	ep_len_mean		49.7	
	ep_rew_mean		230	
	time/			
	fps		475	

iterations	29
time_elapsed	124
total_timesteps	59392
train/	
approx_kl	0.023060175
clip_fraction	0.212
clip_range	0.2
entropy_loss	-23.5
explained_variance	0.503
learning_rate	0.0003
loss	460
n_updates	280
policy_gradient_loss	-0.0643
std	0.963
value_loss	1.02e+03

rollout/	
ep_len_mean	53.3
ep_rew_mean	246
time/	
fps	475
iterations	30
time_elapsed	129
total_timesteps	61440
train/	
approx_kl	0.021096101
clip_fraction	0.183
clip_range	0.2
entropy_loss	-23.5
explained_variance	0.486
learning_rate	0.0003
loss	494
n_updates	290
policy_gradient_loss	-0.0576
std	0.962
value_loss	1.03e+03

rollout/	
ep_len_mean	53.3
ep_rew_mean	247
time/	
fps	475
iterations	31
time_elapsed	133
total_timesteps	63488
train/	



approx_kl	0.023991114	
clip_fraction	0.205	
clip_range	0.2	
entropy_loss	-23.4	
explained_variance	0.505	
learning_rate	0.0003	
loss	502	
n_updates	300	
policy_gradient_loss	-0.0628	
std	0.96	
value_loss	1.1e+03	

rollout/		
ep_len_mean	54.1	
ep_rew_mean	252	
time/		
fps	475	
iterations	32	
time_elapsed	137	
total_timesteps	65536	
train/		
approx_kl	0.020747574	
clip_fraction	0.201	
clip_range	0.2	
entropy_loss	-23.4	
explained_variance	0.466	
learning_rate	0.0003	
loss	438	
n_updates	310	
policy_gradient_loss	-0.0618	
std	0.959	
value_loss	1.16e+03	

rollout/		
ep_len_mean	54.9	
ep_rew_mean	256	
time/		
fps	475	
iterations	33	
time_elapsed	142	
total_timesteps	67584	
train/		
approx_kl	0.018756542	
clip_fraction	0.165	
clip_range	0.2	
entropy_loss	-23.4	

	explained_variance		0.508	
	learning_rate		0.0003	
	loss		418	
	n_updates		320	
	policy_gradient_loss		-0.0545	
	std		0.958	
	value_loss		1.15e+03	

	rollout/			
	ep_len_mean		56.4	
	ep_rew_mean		263	
	time/			
	fps		475	
	iterations		34	
	time_elapsed		146	
	total_timesteps		69632	
	train/			
	approx_kl		0.021349886	
	clip_fraction		0.19	
	clip_range		0.2	
	entropy_loss		-23.4	
	explained_variance		0.532	
	learning_rate		0.0003	
	loss		507	
	n_updates		330	
	policy_gradient_loss		-0.0569	
	std		0.959	
	value_loss		1.14e+03	

	rollout/			
	ep_len_mean		56.8	
	ep_rew_mean		266	
	time/			
	fps		473	
	iterations		35	
	time_elapsed		151	
	total_timesteps		71680	
	train/			
	approx_kl		0.019452553	
	clip_fraction		0.153	
	clip_range		0.2	
	entropy_loss		-23.4	
	explained_variance		0.529	
	learning_rate		0.0003	
	loss		390	
	n_updates		340	

	policy_gradient_loss		-0.0519	
	std		0.958	
	value_loss		1e+03	

---

	rollout/			
	ep_len_mean		57.7	
	ep_rew_mean		268	
	time/			
	fps		473	
	iterations		36	
	time_elapsed		155	
	total_timesteps		73728	
	train/			
	approx_kl		0.02216762	
	clip_fraction		0.191	
	clip_range		0.2	
	entropy_loss		-23.4	
	explained_variance		0.534	
	learning_rate		0.0003	
	loss		518	
	n_updates		350	
	policy_gradient_loss		-0.0582	
	std		0.958	
	value_loss		1.16e+03	

---

	rollout/			
	ep_len_mean		58.4	
	ep_rew_mean		272	
	time/			
	fps		473	
	iterations		37	
	time_elapsed		159	
	total_timesteps		75776	
	train/			
	approx_kl		0.02501797	
	clip_fraction		0.219	
	clip_range		0.2	
	entropy_loss		-23.4	
	explained_variance		0.568	
	learning_rate		0.0003	
	loss		494	
	n_updates		360	
	policy_gradient_loss		-0.061	
	std		0.957	
	value_loss		1.13e+03	

---

---

rollout/		
ep_len_mean	56.5	
ep_rew_mean	261	
time/		
fps	474	
iterations	38	
time_elapsed	164	
total_timesteps	77824	
train/		
approx_kl	0.022883646	
clip_fraction	0.191	
clip_range	0.2	
entropy_loss	-23.3	
explained_variance	0.573	
learning_rate	0.0003	
loss	359	
n_updates	370	
policy_gradient_loss	-0.0569	
std	0.955	
value_loss	916	

---



---

rollout/		
ep_len_mean	55.7	
ep_rew_mean	258	
time/		
fps	475	
iterations	39	
time_elapsed	168	
total_timesteps	79872	
train/		
approx_kl	0.017879618	
clip_fraction	0.169	
clip_range	0.2	
entropy_loss	-23.3	
explained_variance	0.586	
learning_rate	0.0003	
loss	363	
n_updates	380	
policy_gradient_loss	-0.0551	
std	0.953	
value_loss	888	

---



---

rollout/		
ep_len_mean	58.2	
ep_rew_mean	270	

---

time/		
fps	475	
iterations	40	
time_elapsed	172	
total_timesteps	81920	
train/		
approx_kl	0.02333469	
clip_fraction	0.193	
clip_range	0.2	
entropy_loss	-23.3	
explained_variance	0.643	
learning_rate	0.0003	
loss	293	
n_updates	390	
policy_gradient_loss	-0.059	
std	0.953	
value_loss	818	

rollout/		
ep_len_mean	60.1	
ep_rew_mean	280	
time/		
fps	474	
iterations	41	
time_elapsed	176	
total_timesteps	83968	
train/		
approx_kl	0.022446707	
clip_fraction	0.178	
clip_range	0.2	
entropy_loss	-23.3	
explained_variance	0.613	
learning_rate	0.0003	
loss	326	
n_updates	400	
policy_gradient_loss	-0.0541	
std	0.952	
value_loss	935	

rollout/		
ep_len_mean	62.6	
ep_rew_mean	290	
time/		
fps	475	
iterations	42	
time_elapsed	181	

	total_timesteps		86016	
	train/			
	approx_kl		0.022968117	
	clip_fraction		0.181	
	clip_range		0.2	
	entropy_loss		-23.3	
	explained_variance		0.639	
	learning_rate		0.0003	
	loss		451	
	n_updates		410	
	policy_gradient_loss		-0.0549	
	std		0.951	
	value_loss		879	

	rollout/			
	ep_len_mean		60	
	ep_rew_mean		278	
	time/			
	fps		474	
	iterations		43	
	time_elapsed		185	
	total_timesteps		88064	
	train/			
	approx_kl		0.019765958	
	clip_fraction		0.158	
	clip_range		0.2	
	entropy_loss		-23.2	
	explained_variance		0.643	
	learning_rate		0.0003	
	loss		308	
	n_updates		420	
	policy_gradient_loss		-0.0524	
	std		0.95	
	value_loss		931	

	rollout/			
	ep_len_mean		59.9	
	ep_rew_mean		280	
	time/			
	fps		474	
	iterations		44	
	time_elapsed		189	
	total_timesteps		90112	
	train/			
	approx_kl		0.021166435	
	clip_fraction		0.168	

clip_range	0.2	
entropy_loss	-23.2	
explained_variance	0.646	
learning_rate	0.0003	
loss	488	
n_updates	430	
policy_gradient_loss	-0.0533	
std	0.949	
value_loss	814	

rollout/		
ep_len_mean	60.1	
ep_rew_mean	283	
time/		
fps	475	
iterations	45	
time_elapsed	193	
total_timesteps	92160	
train/		
approx_kl	0.02053089	
clip_fraction	0.173	
clip_range	0.2	
entropy_loss	-23.2	
explained_variance	0.675	
learning_rate	0.0003	
loss	400	
n_updates	440	
policy_gradient_loss	-0.0542	
std	0.948	
value_loss	779	

rollout/		
ep_len_mean	62.8	
ep_rew_mean	295	
time/		
fps	475	
iterations	46	
time_elapsed	198	
total_timesteps	94208	
train/		
approx_kl	0.022271642	
clip_fraction	0.168	
clip_range	0.2	
entropy_loss	-23.2	
explained_variance	0.654	
learning_rate	0.0003	

	loss		403	
	n_updates		450	
	policy_gradient_loss		-0.0504	
	std		0.948	
	value_loss		925	

	rollout/			
	ep_len_mean		65.5	
	ep_rew_mean		310	
	time/			
	fps		475	
	iterations		47	
	time_elapsed		202	
	total_timesteps		96256	
	train/			
	approx_kl		0.018243058	
	clip_fraction		0.149	
	clip_range		0.2	
	entropy_loss		-23.2	
	explained_variance		0.659	
	learning_rate		0.0003	
	loss		247	
	n_updates		460	
	policy_gradient_loss		-0.0516	
	std		0.948	
	value_loss		894	

	rollout/			
	ep_len_mean		65.2	
	ep_rew_mean		308	
	time/			
	fps		475	
	iterations		48	
	time_elapsed		206	
	total_timesteps		98304	
	train/			
	approx_kl		0.019605339	
	clip_fraction		0.169	
	clip_range		0.2	
	entropy_loss		-23.2	
	explained_variance		0.719	
	learning_rate		0.0003	
	loss		303	
	n_updates		470	
	policy_gradient_loss		-0.0525	
	std		0.946	



	value_loss	758	
-----			
	rollout/		
	ep_len_mean	66.4	
	ep_rew_mean	313	
	time/		
	fps	475	
	iterations	49	
	time_elapsed	210	
	total_timesteps	100352	
	train/		
	approx_kl	0.019888582	
	clip_fraction	0.141	
	clip_range	0.2	
	entropy_loss	-23.1	
	explained_variance	0.641	
	learning_rate	0.0003	
	loss	444	
	n_updates	480	
	policy_gradient_loss	-0.0475	
	std	0.944	
	value_loss	908	
-----			
	rollout/		
	ep_len_mean	65.5	
	ep_rew_mean	306	
	time/		
	fps	474	
	iterations	50	
	time_elapsed	215	
	total_timesteps	102400	
	train/		
	approx_kl	0.023531526	
	clip_fraction	0.185	
	clip_range	0.2	
	entropy_loss	-23.1	
	explained_variance	0.709	
	learning_rate	0.0003	
	loss	318	
	n_updates	490	
	policy_gradient_loss	-0.0571	
	std	0.944	
	value_loss	806	
-----			
	rollout/		

	ep_len_mean	65.3	
	ep_rew_mean	306	
	time/		
	fps	476	
	iterations	51	
	time_elapsed	219	
	total_timesteps	104448	
	train/		
	approx_kl	0.016149659	
	clip_fraction	0.13	
	clip_range	0.2	
	entropy_loss	-23.1	
	explained_variance	0.669	
	learning_rate	0.0003	
	loss	300	
	n_updates	500	
	policy_gradient_loss	-0.0459	
	std	0.944	
	value_loss	869	

	rollout/		
	ep_len_mean	64	
	ep_rew_mean	300	
	time/		
	fps	476	
	iterations	52	
	time_elapsed	223	
	total_timesteps	106496	
	train/		
	approx_kl	0.018896911	
	clip_fraction	0.153	
	clip_range	0.2	
	entropy_loss	-23.1	
	explained_variance	0.633	
	learning_rate	0.0003	
	loss	397	
	n_updates	510	
	policy_gradient_loss	-0.051	
	std	0.943	
	value_loss	939	

	rollout/		
	ep_len_mean	65.8	
	ep_rew_mean	310	
	time/		
	fps	476	

iterations	53
time_elapsed	227
total_timesteps	108544
train/	
approx_kl	0.019792873
clip_fraction	0.159
clip_range	0.2
entropy_loss	-23.1
explained_variance	0.716
learning_rate	0.0003
loss	296
n_updates	520
policy_gradient_loss	-0.0506
std	0.943
value_loss	686

rollout/	
ep_len_mean	65.9
ep_rew_mean	310
time/	
fps	477
iterations	54
time_elapsed	231
total_timesteps	110592
train/	
approx_kl	0.022327265
clip_fraction	0.164
clip_range	0.2
entropy_loss	-23.1
explained_variance	0.746
learning_rate	0.0003
loss	296
n_updates	530
policy_gradient_loss	-0.0526
std	0.942
value_loss	644

rollout/	
ep_len_mean	66
ep_rew_mean	312
time/	
fps	476
iterations	55
time_elapsed	236
total_timesteps	112640
train/	

approx_kl	0.021175595	
clip_fraction	0.152	
clip_range	0.2	
entropy_loss	-23.1	
explained_variance	0.751	
learning_rate	0.0003	
loss	218	
n_updates	540	
policy_gradient_loss	-0.0476	
std	0.941	
value_loss	664	

rollout/		
ep_len_mean	66.5	
ep_rew_mean	315	
time/		
fps	477	
iterations	56	
time_elapsed	240	
total_timesteps	114688	
train/		
approx_kl	0.02090808	
clip_fraction	0.153	
clip_range	0.2	
entropy_loss	-23.1	
explained_variance	0.771	
learning_rate	0.0003	
loss	187	
n_updates	550	
policy_gradient_loss	-0.0494	
std	0.94	
value_loss	617	

rollout/		
ep_len_mean	66.3	
ep_rew_mean	314	
time/		
fps	477	
iterations	57	
time_elapsed	244	
total_timesteps	116736	
train/		
approx_kl	0.02316078	
clip_fraction	0.169	
clip_range	0.2	
entropy_loss	-23	

	explained_variance	0.747	
	learning_rate	0.0003	
	loss	193	
	n_updates	560	
	policy_gradient_loss	-0.0503	
	std	0.938	
	value_loss	580	

	rollout/		
	ep_len_mean	66.6	
	ep_rew_mean	315	
	time/		
	fps	477	
	iterations	58	
	time_elapsed	248	
	total_timesteps	118784	
	train/		
	approx_kl	0.019586178	
	clip_fraction	0.166	
	clip_range	0.2	
	entropy_loss	-23	
	explained_variance	0.73	
	learning_rate	0.0003	
	loss	323	
	n_updates	570	
	policy_gradient_loss	-0.0515	
	std	0.939	
	value_loss	701	

	rollout/		
	ep_len_mean	66.7	
	ep_rew_mean	316	
	time/		
	fps	477	
	iterations	59	
	time_elapsed	253	
	total_timesteps	120832	
	train/		
	approx_kl	0.019020766	
	clip_fraction	0.155	
	clip_range	0.2	
	entropy_loss	-23	
	explained_variance	0.76	
	learning_rate	0.0003	
	loss	291	
	n_updates	580	

	policy_gradient_loss		-0.0477	
	std		0.939	
	value_loss		684	

	rollout/			
	ep_len_mean		67.8	
	ep_rew_mean		320	
	time/			
	fps		477	
	iterations		60	
	time_elapsed		257	
	total_timesteps		122880	
	train/			
	approx_kl		0.021294706	
	clip_fraction		0.152	
	clip_range		0.2	
	entropy_loss		-23	
	explained_variance		0.777	
	learning_rate		0.0003	
	loss		245	
	n_updates		590	
	policy_gradient_loss		-0.0483	
	std		0.935	
	value_loss		602	

	rollout/			
	ep_len_mean		68.6	
	ep_rew_mean		325	
	time/			
	fps		477	
	iterations		61	
	time_elapsed		261	
	total_timesteps		124928	
	train/			
	approx_kl		0.018613134	
	clip_fraction		0.141	
	clip_range		0.2	
	entropy_loss		-22.9	
	explained_variance		0.742	
	learning_rate		0.0003	
	loss		277	
	n_updates		600	
	policy_gradient_loss		-0.0439	
	std		0.934	
	value_loss		695	

---

rollout/		
ep_len_mean	70.7	
ep_rew_mean	333	
time/		
fps	476	
iterations	62	
time_elapsed	266	
total_timesteps	126976	
train/		
approx_kl	0.02136048	
clip_fraction	0.163	
clip_range	0.2	
entropy_loss	-22.9	
explained_variance	0.798	
learning_rate	0.0003	
loss	217	
n_updates	610	
policy_gradient_loss	-0.0483	
std	0.935	
value_loss	621	

---



---

rollout/		
ep_len_mean	69.3	
ep_rew_mean	326	
time/		
fps	475	
iterations	63	
time_elapsed	271	
total_timesteps	129024	
train/		
approx_kl	0.020181686	
clip_fraction	0.166	
clip_range	0.2	
entropy_loss	-22.9	
explained_variance	0.801	
learning_rate	0.0003	
loss	212	
n_updates	620	
policy_gradient_loss	-0.051	
std	0.934	
value_loss	577	

---



---

rollout/		
ep_len_mean	67.6	
ep_rew_mean	318	

---

time/		
fps	475	
iterations	64	
time_elapsed	275	
total_timesteps	131072	
train/		
approx_kl	0.020851977	
clip_fraction	0.169	
clip_range	0.2	
entropy_loss	-22.9	
explained_variance	0.807	
learning_rate	0.0003	
loss	225	
n_updates	630	
policy_gradient_loss	-0.0498	
std	0.933	
value_loss	536	

---

rollout/		
ep_len_mean	66.8	
ep_rew_mean	315	
time/		
fps	475	
iterations	65	
time_elapsed	279	
total_timesteps	133120	
train/		
approx_kl	0.019043177	
clip_fraction	0.156	
clip_range	0.2	
entropy_loss	-22.9	
explained_variance	0.813	
learning_rate	0.0003	
loss	155	
n_updates	640	
policy_gradient_loss	-0.0488	
std	0.933	
value_loss	512	

---

rollout/		
ep_len_mean	65.4	
ep_rew_mean	311	
time/		
fps	475	
iterations	66	
time_elapsed	284	



	total_timesteps		135168	
	train/			
	approx_kl		0.019931728	
	clip_fraction		0.157	
	clip_range		0.2	
	entropy_loss		-22.9	
	explained_variance		0.806	
	learning_rate		0.0003	
	loss		159	
	n_updates		650	
	policy_gradient_loss		-0.0454	
	std		0.932	
	value_loss		474	

---

	rollout/			
	ep_len_mean		66.6	
	ep_rew_mean		316	
	time/			
	fps		475	
	iterations		67	
	time_elapsed		288	
	total_timesteps		137216	
	train/			
	approx_kl		0.0228469	
	clip_fraction		0.185	
	clip_range		0.2	
	entropy_loss		-22.9	
	explained_variance		0.852	
	learning_rate		0.0003	
	loss		278	
	n_updates		660	
	policy_gradient_loss		-0.0552	
	std		0.932	
	value_loss		452	

---

	rollout/			
	ep_len_mean		64.6	
	ep_rew_mean		305	
	time/			
	fps		476	
	iterations		68	
	time_elapsed		292	
	total_timesteps		139264	
	train/			
	approx_kl		0.019598013	
	clip_fraction		0.155	

clip_range	0.2	
entropy_loss	-22.9	
explained_variance	0.828	
learning_rate	0.0003	
loss	144	
n_updates	670	
policy_gradient_loss	-0.0482	
std	0.932	
value_loss	493	

rollout/		
ep_len_mean	65.5	
ep_rew_mean	307	
time/		
fps	475	
iterations	69	
time_elapsed	297	
total_timesteps	141312	
train/		
approx_kl	0.024314534	
clip_fraction	0.2	
clip_range	0.2	
entropy_loss	-22.9	
explained_variance	0.828	
learning_rate	0.0003	
loss	130	
n_updates	680	
policy_gradient_loss	-0.0529	
std	0.931	
value_loss	392	

rollout/		
ep_len_mean	68.5	
ep_rew_mean	322	
time/		
fps	475	
iterations	70	
time_elapsed	301	
total_timesteps	143360	
train/		
approx_kl	0.024317527	
clip_fraction	0.192	
clip_range	0.2	
entropy_loss	-22.9	
explained_variance	0.808	
learning_rate	0.0003	

	loss		158	
	n_updates		690	
	policy_gradient_loss		-0.053	
	std		0.931	
	value_loss		447	

	rollout/			
	ep_len_mean		71.2	
	ep_rew_mean		336	
	time/			
	fps		475	
	iterations		71	
	time_elapsed		305	
	total_timesteps		145408	
	train/			
	approx_kl		0.021015283	
	clip_fraction		0.166	
	clip_range		0.2	
	entropy_loss		-22.9	
	explained_variance		0.786	
	learning_rate		0.0003	
	loss		129	
	n_updates		700	
	policy_gradient_loss		-0.049	
	std		0.93	
	value_loss		449	

	rollout/			
	ep_len_mean		72.2	
	ep_rew_mean		342	
	time/			
	fps		476	
	iterations		72	
	time_elapsed		309	
	total_timesteps		147456	
	train/			
	approx_kl		0.024852328	
	clip_fraction		0.176	
	clip_range		0.2	
	entropy_loss		-22.9	
	explained_variance		0.825	
	learning_rate		0.0003	
	loss		148	
	n_updates		710	
	policy_gradient_loss		-0.0514	
	std		0.931	

	value_loss	472	
-----			
	rollout/		
	ep_len_mean	73.7	
	ep_rew_mean	349	
	time/		
	fps	475	
	iterations	73	
	time_elapsed	314	
	total_timesteps	149504	
	train/		
	approx_kl	0.018236108	
	clip_fraction	0.149	
	clip_range	0.2	
	entropy_loss	-22.9	
	explained_variance	0.825	
	learning_rate	0.0003	
	loss	215	
	n_updates	720	
	policy_gradient_loss	-0.0475	
	std	0.931	
	value_loss	506	
-----			
	rollout/		
	ep_len_mean	70	
	ep_rew_mean	331	
	time/		
	fps	474	
	iterations	74	
	time_elapsed	319	
	total_timesteps	151552	
	train/		
	approx_kl	0.01886107	
	clip_fraction	0.157	
	clip_range	0.2	
	entropy_loss	-22.9	
	explained_variance	0.835	
	learning_rate	0.0003	
	loss	203	
	n_updates	730	
	policy_gradient_loss	-0.0469	
	std	0.931	
	value_loss	495	
-----			
	rollout/		

	ep_len_mean		71.2	
	ep_rew_mean		338	
	time/			
	fps		474	
	iterations		75	
	time_elapsed		323	
	total_timesteps		153600	
	train/			
	approx_kl		0.021296185	
	clip_fraction		0.172	
	clip_range		0.2	
	entropy_loss		-22.9	
	explained_variance		0.847	
	learning_rate		0.0003	
	loss		158	
	n_updates		740	
	policy_gradient_loss		-0.0482	
	std		0.93	
	value_loss		387	

	rollout/			
	ep_len_mean		72.1	
	ep_rew_mean		340	
	time/			
	fps		474	
	iterations		76	
	time_elapsed		328	
	total_timesteps		155648	
	train/			
	approx_kl		0.01888765	
	clip_fraction		0.141	
	clip_range		0.2	
	entropy_loss		-22.8	
	explained_variance		0.823	
	learning_rate		0.0003	
	loss		280	
	n_updates		750	
	policy_gradient_loss		-0.0459	
	std		0.929	
	value_loss		523	

	rollout/			
	ep_len_mean		71.3	
	ep_rew_mean		339	
	time/			
	fps		473	

iterations	77
time_elapsed	333
total_timesteps	157696
train/	
approx_kl	0.019331139
clip_fraction	0.143
clip_range	0.2
entropy_loss	-22.8
explained_variance	0.803
learning_rate	0.0003
loss	145
n_updates	760
policy_gradient_loss	-0.0472
std	0.927
value_loss	526

rollout/	
ep_len_mean	71.9
ep_rew_mean	342
time/	
fps	472
iterations	78
time_elapsed	338
total_timesteps	159744
train/	
approx_kl	0.02581295
clip_fraction	0.192
clip_range	0.2
entropy_loss	-22.8
explained_variance	0.876
learning_rate	0.0003
loss	173
n_updates	770
policy_gradient_loss	-0.0509
std	0.926
value_loss	372

rollout/	
ep_len_mean	73.6
ep_rew_mean	348
time/	
fps	472
iterations	79
time_elapsed	342
total_timesteps	161792
train/	

approx_kl	0.020753756	
clip_fraction	0.177	
clip_range	0.2	
entropy_loss	-22.8	
explained_variance	0.865	
learning_rate	0.0003	
loss	138	
n_updates	780	
policy_gradient_loss	-0.0478	
std	0.925	
value_loss	361	

rollout/		
ep_len_mean	74.2	
ep_rew_mean	351	
time/		
fps	471	
iterations	80	
time_elapsed	347	
total_timesteps	163840	
train/		
approx_kl	0.020470867	
clip_fraction	0.166	
clip_range	0.2	
entropy_loss	-22.8	
explained_variance	0.828	
learning_rate	0.0003	
loss	128	
n_updates	790	
policy_gradient_loss	-0.0475	
std	0.925	
value_loss	494	

rollout/		
ep_len_mean	76.6	
ep_rew_mean	360	
time/		
fps	471	
iterations	81	
time_elapsed	352	
total_timesteps	165888	
train/		
approx_kl	0.018120013	
clip_fraction	0.152	
clip_range	0.2	
entropy_loss	-22.8	

	explained_variance	0.827	
	learning_rate	0.0003	
	loss	182	
	n_updates	800	
	policy_gradient_loss	-0.0474	
	std	0.925	
	value_loss	562	

	rollout/		
	ep_len_mean	74.3	
	ep_rew_mean	351	
	time/		
	fps	470	
	iterations	82	
	time_elapsed	356	
	total_timesteps	167936	
	train/		
	approx_kl	0.019318711	
	clip_fraction	0.15	
	clip_range	0.2	
	entropy_loss	-22.8	
	explained_variance	0.852	
	learning_rate	0.0003	
	loss	210	
	n_updates	810	
	policy_gradient_loss	-0.0465	
	std	0.926	
	value_loss	518	

	rollout/		
	ep_len_mean	73.6	
	ep_rew_mean	348	
	time/		
	fps	470	
	iterations	83	
	time_elapsed	361	
	total_timesteps	169984	
	train/		
	approx_kl	0.021066088	
	clip_fraction	0.163	
	clip_range	0.2	
	entropy_loss	-22.8	
	explained_variance	0.816	
	learning_rate	0.0003	
	loss	252	
	n_updates	820	



	policy_gradient_loss		-0.0482	
	std		0.925	
	value_loss		514	

---

	rollout/			
	ep_len_mean		72.5	
	ep_rew_mean		341	
	time/			
	fps		470	
	iterations		84	
	time_elapsed		365	
	total_timesteps		172032	
	train/			
	approx_kl		0.022356246	
	clip_fraction		0.171	
	clip_range		0.2	
	entropy_loss		-22.8	
	explained_variance		0.848	
	learning_rate		0.0003	
	loss		149	
	n_updates		830	
	policy_gradient_loss		-0.0497	
	std		0.924	
	value_loss		439	

---



---

	rollout/			
	ep_len_mean		72.2	
	ep_rew_mean		342	
	time/			
	fps		469	
	iterations		85	
	time_elapsed		370	
	total_timesteps		174080	
	train/			
	approx_kl		0.01929266	
	clip_fraction		0.147	
	clip_range		0.2	
	entropy_loss		-22.7	
	explained_variance		0.857	
	learning_rate		0.0003	
	loss		142	
	n_updates		840	
	policy_gradient_loss		-0.046	
	std		0.923	
	value_loss		468	

---

-----		
rollout/		
ep_len_mean	71.2	
ep_rew_mean	336	
time/		
fps	469	
iterations	86	
time_elapsed	374	
total_timesteps	176128	
train/		
approx_kl	0.020821612	
clip_fraction	0.176	
clip_range	0.2	
entropy_loss	-22.7	
explained_variance	0.856	
learning_rate	0.0003	
loss	159	
n_updates	850	
policy_gradient_loss	-0.0491	
std	0.922	
value_loss	415	
-----		

-----		
rollout/		
ep_len_mean	70.8	
ep_rew_mean	337	
time/		
fps	469	
iterations	87	
time_elapsed	379	
total_timesteps	178176	
train/		
approx_kl	0.025869755	
clip_fraction	0.193	
clip_range	0.2	
entropy_loss	-22.7	
explained_variance	0.853	
learning_rate	0.0003	
loss	116	
n_updates	860	
policy_gradient_loss	-0.052	
std	0.921	
value_loss	409	
-----		

-----		
rollout/		
ep_len_mean	71	
ep_rew_mean	338	

time/		
fps	469	
iterations	88	
time_elapsed	383	
total_timesteps	180224	
train/		
approx_kl	0.022976631	
clip_fraction	0.187	
clip_range	0.2	
entropy_loss	-22.7	
explained_variance	0.869	
learning_rate	0.0003	
loss	176	
n_updates	870	
policy_gradient_loss	-0.0501	
std	0.92	
value_loss	408	

---

rollout/		
ep_len_mean	71	
ep_rew_mean	337	
time/		
fps	469	
iterations	89	
time_elapsed	388	
total_timesteps	182272	
train/		
approx_kl	0.023173138	
clip_fraction	0.178	
clip_range	0.2	
entropy_loss	-22.7	
explained_variance	0.88	
learning_rate	0.0003	
loss	125	
n_updates	880	
policy_gradient_loss	-0.0505	
std	0.918	
value_loss	351	

---

rollout/		
ep_len_mean	70.1	
ep_rew_mean	333	
time/		
fps	469	
iterations	90	
time_elapsed	392	

	total_timesteps		184320	
	train/			
	approx_kl		0.023492534	
	clip_fraction		0.191	
	clip_range		0.2	
	entropy_loss		-22.6	
	explained_variance		0.888	
	learning_rate		0.0003	
	loss		145	
	n_updates		890	
	policy_gradient_loss		-0.0531	
	std		0.917	
	value_loss		336	

	rollout/			
	ep_len_mean		70.8	
	ep_rew_mean		336	
	time/			
	fps		469	
	iterations		91	
	time_elapsed		397	
	total_timesteps		186368	
	train/			
	approx_kl		0.023987236	
	clip_fraction		0.202	
	clip_range		0.2	
	entropy_loss		-22.6	
	explained_variance		0.867	
	learning_rate		0.0003	
	loss		142	
	n_updates		900	
	policy_gradient_loss		-0.0555	
	std		0.916	
	value_loss		374	

	rollout/			
	ep_len_mean		71.2	
	ep_rew_mean		335	
	time/			
	fps		468	
	iterations		92	
	time_elapsed		402	
	total_timesteps		188416	
	train/			
	approx_kl		0.025471434	
	clip_fraction		0.185	

clip_range	0.2	
entropy_loss	-22.6	
explained_variance	0.856	
learning_rate	0.0003	
loss	101	
n_updates	910	
policy_gradient_loss	-0.0514	
std	0.915	
value_loss	408	

rollout/		
ep_len_mean	73.3	
ep_rew_mean	345	
time/		
fps	467	
iterations	93	
time_elapsed	407	
total_timesteps	190464	
train/		
approx_kl	0.02463483	
clip_fraction	0.181	
clip_range	0.2	
entropy_loss	-22.6	
explained_variance	0.832	
learning_rate	0.0003	
loss	269	
n_updates	920	
policy_gradient_loss	-0.05	
std	0.913	
value_loss	458	

rollout/		
ep_len_mean	72.7	
ep_rew_mean	345	
time/		
fps	467	
iterations	94	
time_elapsed	411	
total_timesteps	192512	
train/		
approx_kl	0.023497637	
clip_fraction	0.18	
clip_range	0.2	
entropy_loss	-22.5	
explained_variance	0.86	
learning_rate	0.0003	

	loss		162	
	n_updates		930	
	policy_gradient_loss		-0.0519	
	std		0.912	
	value_loss		440	

	rollout/			
	ep_len_mean		73.8	
	ep_rew_mean		351	
	time/			
	fps		467	
	iterations		95	
	time_elapsed		416	
	total_timesteps		194560	
	train/			
	approx_kl		0.026575403	
	clip_fraction		0.213	
	clip_range		0.2	
	entropy_loss		-22.5	
	explained_variance		0.877	
	learning_rate		0.0003	
	loss		127	
	n_updates		940	
	policy_gradient_loss		-0.0559	
	std		0.91	
	value_loss		368	

	rollout/			
	ep_len_mean		74.1	
	ep_rew_mean		352	
	time/			
	fps		467	
	iterations		96	
	time_elapsed		420	
	total_timesteps		196608	
	train/			
	approx_kl		0.027708221	
	clip_fraction		0.199	
	clip_range		0.2	
	entropy_loss		-22.5	
	explained_variance		0.867	
	learning_rate		0.0003	
	loss		119	
	n_updates		950	
	policy_gradient_loss		-0.0502	
	std		0.909	

	value_loss		376	
-----				
	rollout/			
	ep_len_mean		73.4	
	ep_rew_mean		349	
	time/			
	fps		467	
	iterations		97	
	time_elapsed		425	
	total_timesteps		198656	
	train/			
	approx_kl		0.023997769	
	clip_fraction		0.189	
	clip_range		0.2	
	entropy_loss		-22.5	
	explained_variance		0.868	
	learning_rate		0.0003	
	loss		149	
	n_updates		960	
	policy_gradient_loss		-0.0496	
	std		0.909	
	value_loss		394	
-----				
	rollout/			
	ep_len_mean		72.4	
	ep_rew_mean		342	
	time/			
	fps		467	
	iterations		98	
	time_elapsed		429	
	total_timesteps		200704	
	train/			
	approx_kl		0.022676608	
	clip_fraction		0.188	
	clip_range		0.2	
	entropy_loss		-22.5	
	explained_variance		0.858	
	learning_rate		0.0003	
	loss		119	
	n_updates		970	
	policy_gradient_loss		-0.0514	
	std		0.909	
	value_loss		469	
-----				
	rollout/			

	ep_len_mean	71.9	
	ep_rew_mean	341	
	time/		
	fps	467	
	iterations	99	
	time_elapsed	434	
	total_timesteps	202752	
	train/		
	approx_kl	0.028271373	
	clip_fraction	0.215	
	clip_range	0.2	
	entropy_loss	-22.5	
	explained_variance	0.863	
	learning_rate	0.0003	
	loss	85	
	n_updates	980	
	policy_gradient_loss	-0.0575	
	std	0.908	
	value_loss	379	

---

	rollout/		
	ep_len_mean	74.3	
	ep_rew_mean	350	
	time/		
	fps	466	
	iterations	100	
	time_elapsed	438	
	total_timesteps	204800	
	train/		
	approx_kl	0.02827181	
	clip_fraction	0.214	
	clip_range	0.2	
	entropy_loss	-22.4	
	explained_variance	0.862	
	learning_rate	0.0003	
	loss	111	
	n_updates	990	
	policy_gradient_loss	-0.0547	
	std	0.906	
	value_loss	373	

---

	rollout/		
	ep_len_mean	75.3	
	ep_rew_mean	352	
	time/		
	fps	466	



iterations	101
time_elapsed	443
total_timesteps	206848
train/	
approx_kl	0.023341492
clip_fraction	0.204
clip_range	0.2
entropy_loss	-22.4
explained_variance	0.85
learning_rate	0.0003
loss	83.3
n_updates	1000
policy_gradient_loss	-0.0528
std	0.905
value_loss	367

rollout/	
ep_len_mean	74.6
ep_rew_mean	347
time/	
fps	466
iterations	102
time_elapsed	447
total_timesteps	208896
train/	
approx_kl	0.024424857
clip_fraction	0.203
clip_range	0.2
entropy_loss	-22.4
explained_variance	0.88
learning_rate	0.0003
loss	160
n_updates	1010
policy_gradient_loss	-0.055
std	0.904
value_loss	383

rollout/	
ep_len_mean	74.5
ep_rew_mean	347
time/	
fps	466
iterations	103
time_elapsed	452
total_timesteps	210944
train/	

approx_kl	0.026520483	
clip_fraction	0.198	
clip_range	0.2	
entropy_loss	-22.4	
explained_variance	0.882	
learning_rate	0.0003	
loss	172	
n_updates	1020	
policy_gradient_loss	-0.052	
std	0.903	
value_loss	364	

rollout/		
ep_len_mean	73.4	
ep_rew_mean	343	
time/		
fps	466	
iterations	104	
time_elapsed	456	
total_timesteps	212992	
train/		
approx_kl	0.03129163	
clip_fraction	0.212	
clip_range	0.2	
entropy_loss	-22.3	
explained_variance	0.898	
learning_rate	0.0003	
loss	98.7	
n_updates	1030	
policy_gradient_loss	-0.0527	
std	0.903	
value_loss	271	

rollout/		
ep_len_mean	73.2	
ep_rew_mean	342	
time/		
fps	465	
iterations	105	
time_elapsed	461	
total_timesteps	215040	
train/		
approx_kl	0.023996562	
clip_fraction	0.183	
clip_range	0.2	
entropy_loss	-22.3	

	explained_variance	0.85	
	learning_rate	0.0003	
	loss	141	
	n_updates	1040	
	policy_gradient_loss	-0.0484	
	std	0.902	
	value_loss	427	

	rollout/		
	ep_len_mean	72.6	
	ep_rew_mean	340	
	time/		
	fps	465	
	iterations	106	
	time_elapsed	466	
	total_timesteps	217088	
	train/		
	approx_kl	0.021450441	
	clip_fraction	0.161	
	clip_range	0.2	
	entropy_loss	-22.3	
	explained_variance	0.849	
	learning_rate	0.0003	
	loss	140	
	n_updates	1050	
	policy_gradient_loss	-0.043	
	std	0.901	
	value_loss	483	

	rollout/		
	ep_len_mean	73.6	
	ep_rew_mean	349	
	time/		
	fps	465	
	iterations	107	
	time_elapsed	471	
	total_timesteps	219136	
	train/		
	approx_kl	0.023194268	
	clip_fraction	0.166	
	clip_range	0.2	
	entropy_loss	-22.3	
	explained_variance	0.881	
	learning_rate	0.0003	
	loss	142	
	n_updates	1060	

	policy_gradient_loss		-0.0486	
	std		0.9	
	value_loss		385	

	rollout/			
	ep_len_mean		73.6	
	ep_rew_mean		349	
	time/			
	fps		465	
	iterations		108	
	time_elapsed		475	
	total_timesteps		221184	
	train/			
	approx_kl		0.025292404	
	clip_fraction		0.211	
	clip_range		0.2	
	entropy_loss		-22.3	
	explained_variance		0.895	
	learning_rate		0.0003	
	loss		112	
	n_updates		1070	
	policy_gradient_loss		-0.0538	
	std		0.899	
	value_loss		356	

	rollout/			
	ep_len_mean		74.2	
	ep_rew_mean		355	
	time/			
	fps		465	
	iterations		109	
	time_elapsed		479	
	total_timesteps		223232	
	train/			
	approx_kl		0.022817308	
	clip_fraction		0.157	
	clip_range		0.2	
	entropy_loss		-22.3	
	explained_variance		0.832	
	learning_rate		0.0003	
	loss		244	
	n_updates		1080	
	policy_gradient_loss		-0.046	
	std		0.898	
	value_loss		523	

-----		
rollout/		
ep_len_mean	75.2	
ep_rew_mean	362	
time/		
fps	464	
iterations	110	
time_elapsed	484	
total_timesteps	225280	
train/		
approx_kl	0.022698823	
clip_fraction	0.174	
clip_range	0.2	
entropy_loss	-22.3	
explained_variance	0.88	
learning_rate	0.0003	
loss	183	
n_updates	1090	
policy_gradient_loss	-0.045	
std	0.898	
value_loss	421	
-----		

-----		
rollout/		
ep_len_mean	76.2	
ep_rew_mean	366	
time/		
fps	464	
iterations	111	
time_elapsed	489	
total_timesteps	227328	
train/		
approx_kl	0.025699953	
clip_fraction	0.175	
clip_range	0.2	
entropy_loss	-22.3	
explained_variance	0.843	
learning_rate	0.0003	
loss	124	
n_updates	1100	
policy_gradient_loss	-0.0513	
std	0.898	
value_loss	402	
-----		

-----		
rollout/		
ep_len_mean	74.3	
ep_rew_mean	359	

time/		
fps	464	
iterations	112	
time_elapsed	493	
total_timesteps	229376	
train/		
approx_kl	0.025716085	
clip_fraction	0.194	
clip_range	0.2	
entropy_loss	-22.2	
explained_variance	0.879	
learning_rate	0.0003	
loss	97.6	
n_updates	1110	
policy_gradient_loss	-0.0488	
std	0.897	
value_loss	384	

---

rollout/		
ep_len_mean	75.2	
ep_rew_mean	360	
time/		
fps	464	
iterations	113	
time_elapsed	497	
total_timesteps	231424	
train/		
approx_kl	0.026015269	
clip_fraction	0.201	
clip_range	0.2	
entropy_loss	-22.2	
explained_variance	0.897	
learning_rate	0.0003	
loss	161	
n_updates	1120	
policy_gradient_loss	-0.0545	
std	0.896	
value_loss	356	

---

rollout/		
ep_len_mean	75.7	
ep_rew_mean	357	
time/		
fps	464	
iterations	114	
time_elapsed	502	

	total_timesteps		233472	
	train/			
	approx_kl		0.025203353	
	clip_fraction		0.183	
	clip_range		0.2	
	entropy_loss		-22.2	
	explained_variance		0.861	
	learning_rate		0.0003	
	loss		218	
	n_updates		1130	
	policy_gradient_loss		-0.05	
	std		0.896	
	value_loss		461	

	rollout/			
	ep_len_mean		72.9	
	ep_rew_mean		342	
	time/			
	fps		464	
	iterations		115	
	time_elapsed		507	
	total_timesteps		235520	
	train/			
	approx_kl		0.022531696	
	clip_fraction		0.181	
	clip_range		0.2	
	entropy_loss		-22.2	
	explained_variance		0.878	
	learning_rate		0.0003	
	loss		112	
	n_updates		1140	
	policy_gradient_loss		-0.0501	
	std		0.896	
	value_loss		357	

	rollout/			
	ep_len_mean		73.7	
	ep_rew_mean		346	
	time/			
	fps		464	
	iterations		116	
	time_elapsed		511	
	total_timesteps		237568	
	train/			
	approx_kl		0.027084976	
	clip_fraction		0.18	

clip_range	0.2	
entropy_loss	-22.2	
explained_variance	0.878	
learning_rate	0.0003	
loss	145	
n_updates	1150	
policy_gradient_loss	-0.0502	
std	0.895	
value_loss	398	

rollout/		
ep_len_mean	73.1	
ep_rew_mean	346	
time/		
fps	464	
iterations	117	
time_elapsed	516	
total_timesteps	239616	
train/		
approx_kl	0.026010264	
clip_fraction	0.195	
clip_range	0.2	
entropy_loss	-22.2	
explained_variance	0.881	
learning_rate	0.0003	
loss	95.3	
n_updates	1160	
policy_gradient_loss	-0.0504	
std	0.893	
value_loss	317	

rollout/		
ep_len_mean	71.8	
ep_rew_mean	339	
time/		
fps	464	
iterations	118	
time_elapsed	520	
total_timesteps	241664	
train/		
approx_kl	0.025635201	
clip_fraction	0.199	
clip_range	0.2	
entropy_loss	-22.1	
explained_variance	0.872	
learning_rate	0.0003	



	loss		103	
	n_updates		1170	
	policy_gradient_loss		-0.0512	
	std		0.892	
	value_loss		367	

	rollout/			
	ep_len_mean		71.5	
	ep_rew_mean		338	
	time/			
	fps		464	
	iterations		119	
	time_elapsed		524	
	total_timesteps		243712	
	train/			
	approx_kl		0.023276102	
	clip_fraction		0.186	
	clip_range		0.2	
	entropy_loss		-22.1	
	explained_variance		0.894	
	learning_rate		0.0003	
	loss		108	
	n_updates		1180	
	policy_gradient_loss		-0.0509	
	std		0.893	
	value_loss		335	

	rollout/			
	ep_len_mean		70.6	
	ep_rew_mean		333	
	time/			
	fps		463	
	iterations		120	
	time_elapsed		529	
	total_timesteps		245760	
	train/			
	approx_kl		0.02408496	
	clip_fraction		0.19	
	clip_range		0.2	
	entropy_loss		-22.1	
	explained_variance		0.879	
	learning_rate		0.0003	
	loss		86.2	
	n_updates		1190	
	policy_gradient_loss		-0.0512	
	std		0.893	

	value_loss		361	
-----				
	rollout/			
	ep_len_mean		72.5	
	ep_rew_mean		344	
	time/			
	fps		463	
	iterations		121	
	time_elapsed		534	
	total_timesteps		247808	
	train/			
	approx_kl		0.028162234	
	clip_fraction		0.227	
	clip_range		0.2	
	entropy_loss		-22.1	
	explained_variance		0.887	
	learning_rate		0.0003	
	loss		75.7	
	n_updates		1200	
	policy_gradient_loss		-0.0548	
	std		0.892	
	value_loss		316	
-----				
	rollout/			
	ep_len_mean		74.7	
	ep_rew_mean		352	
	time/			
	fps		463	
	iterations		122	
	time_elapsed		538	
	total_timesteps		249856	
	train/			
	approx_kl		0.03242269	
	clip_fraction		0.216	
	clip_range		0.2	
	entropy_loss		-22.1	
	explained_variance		0.886	
	learning_rate		0.0003	
	loss		132	
	n_updates		1210	
	policy_gradient_loss		-0.0508	
	std		0.892	
	value_loss		334	
-----				
	rollout/			

	ep_len_mean		74.5	
	ep_rew_mean		352	
	time/			
	fps		463	
	iterations		123	
	time_elapsed		543	
	total_timesteps		251904	
	train/			
	approx_kl		0.029214488	
	clip_fraction		0.197	
	clip_range		0.2	
	entropy_loss		-22.1	
	explained_variance		0.894	
	learning_rate		0.0003	
	loss		86.8	
	n_updates		1220	
	policy_gradient_loss		-0.0508	
	std		0.891	
	value_loss		335	

	rollout/			
	ep_len_mean		76.1	
	ep_rew_mean		362	
	time/			
	fps		463	
	iterations		124	
	time_elapsed		547	
	total_timesteps		253952	
	train/			
	approx_kl		0.027607955	
	clip_fraction		0.201	
	clip_range		0.2	
	entropy_loss		-22.1	
	explained_variance		0.899	
	learning_rate		0.0003	
	loss		87.3	
	n_updates		1230	
	policy_gradient_loss		-0.0548	
	std		0.889	
	value_loss		308	

	rollout/			
	ep_len_mean		75.6	
	ep_rew_mean		361	
	time/			
	fps		463	

iterations	125
time_elapsed	552
total_timesteps	256000
train/	
approx_kl	0.02644234
clip_fraction	0.202
clip_range	0.2
entropy_loss	-22.1
explained_variance	0.901
learning_rate	0.0003
loss	170
n_updates	1240
policy_gradient_loss	-0.054
std	0.889
value_loss	310

rollout/	
ep_len_mean	74.9
ep_rew_mean	357
time/	
fps	463
iterations	126
time_elapsed	557
total_timesteps	258048
train/	
approx_kl	0.02795106
clip_fraction	0.19
clip_range	0.2
entropy_loss	-22.1
explained_variance	0.897
learning_rate	0.0003
loss	143
n_updates	1250
policy_gradient_loss	-0.0472
std	0.889
value_loss	325

rollout/	
ep_len_mean	72.5
ep_rew_mean	346
time/	
fps	463
iterations	127
time_elapsed	561
total_timesteps	260096
train/	

approx_kl	0.023835853	
clip_fraction	0.181	
clip_range	0.2	
entropy_loss	-22.1	
explained_variance	0.897	
learning_rate	0.0003	
loss	118	
n_updates	1260	
policy_gradient_loss	-0.0511	
std	0.888	
value_loss	360	

rollout/		
ep_len_mean	73.1	
ep_rew_mean	346	
time/		
fps	463	
iterations	128	
time_elapsed	566	
total_timesteps	262144	
train/		
approx_kl	0.027307792	
clip_fraction	0.204	
clip_range	0.2	
entropy_loss	-22	
explained_variance	0.907	
learning_rate	0.0003	
loss	90.9	
n_updates	1270	
policy_gradient_loss	-0.0518	
std	0.887	
value_loss	309	

rollout/		
ep_len_mean	73.2	
ep_rew_mean	348	
time/		
fps	463	
iterations	129	
time_elapsed	570	
total_timesteps	264192	
train/		
approx_kl	0.026529191	
clip_fraction	0.205	
clip_range	0.2	
entropy_loss	-22	

	explained_variance	0.893	
	learning_rate	0.0003	
	loss	77.3	
	n_updates	1280	
	policy_gradient_loss	-0.0525	
	std	0.886	
	value_loss	310	

	rollout/		
	ep_len_mean	75.3	
	ep_rew_mean	358	
	time/		
	fps	463	
	iterations	130	
	time_elapsed	574	
	total_timesteps	266240	
	train/		
	approx_kl	0.027898839	
	clip_fraction	0.222	
	clip_range	0.2	
	entropy_loss	-22	
	explained_variance	0.904	
	learning_rate	0.0003	
	loss	112	
	n_updates	1290	
	policy_gradient_loss	-0.0522	
	std	0.886	
	value_loss	336	

	rollout/		
	ep_len_mean	77.5	
	ep_rew_mean	368	
	time/		
	fps	463	
	iterations	131	
	time_elapsed	578	
	total_timesteps	268288	
	train/		
	approx_kl	0.028017942	
	clip_fraction	0.204	
	clip_range	0.2	
	entropy_loss	-22	
	explained_variance	0.912	
	learning_rate	0.0003	
	loss	72.7	
	n_updates	1300	

	policy_gradient_loss		-0.0548	
	std		0.885	
	value_loss		306	

---

	rollout/			
	ep_len_mean		78.3	
	ep_rew_mean		374	
	time/			
	fps		463	
	iterations		132	
	time_elapsed		583	
	total_timesteps		270336	
	train/			
	approx_kl		0.02853347	
	clip_fraction		0.213	
	clip_range		0.2	
	entropy_loss		-22	
	explained_variance		0.909	
	learning_rate		0.0003	
	loss		116	
	n_updates		1310	
	policy_gradient_loss		-0.0519	
	std		0.885	
	value_loss		322	

---



---

	rollout/			
	ep_len_mean		79.9	
	ep_rew_mean		381	
	time/			
	fps		463	
	iterations		133	
	time_elapsed		587	
	total_timesteps		272384	
	train/			
	approx_kl		0.024572883	
	clip_fraction		0.191	
	clip_range		0.2	
	entropy_loss		-22	
	explained_variance		0.877	
	learning_rate		0.0003	
	loss		169	
	n_updates		1320	
	policy_gradient_loss		-0.0502	
	std		0.883	
	value_loss		395	

---

-----		
rollout/		
ep_len_mean	77.4	
ep_rew_mean	371	
time/		
fps	463	
iterations	134	
time_elapsed	592	
total_timesteps	274432	
train/		
approx_kl	0.025752977	
clip_fraction	0.212	
clip_range	0.2	
entropy_loss	-22	
explained_variance	0.915	
learning_rate	0.0003	
loss	118	
n_updates	1330	
policy_gradient_loss	-0.0546	
std	0.882	
value_loss	297	
-----		

-----		
rollout/		
ep_len_mean	74.1	
ep_rew_mean	354	
time/		
fps	463	
iterations	135	
time_elapsed	596	
total_timesteps	276480	
train/		
approx_kl	0.027920296	
clip_fraction	0.216	
clip_range	0.2	
entropy_loss	-21.9	
explained_variance	0.878	
learning_rate	0.0003	
loss	92.8	
n_updates	1340	
policy_gradient_loss	-0.0497	
std	0.881	
value_loss	350	
-----		

-----		
rollout/		
ep_len_mean	73.8	
ep_rew_mean	357	



time/		
fps	463	
iterations	136	
time_elapsed	601	
total_timesteps	278528	
train/		
approx_kl	0.03202445	
clip_fraction	0.222	
clip_range	0.2	
entropy_loss	-21.9	
explained_variance	0.902	
learning_rate	0.0003	
loss	103	
n_updates	1350	
policy_gradient_loss	-0.0551	
std	0.88	
value_loss	297	

rollout/		
ep_len_mean	72.3	
ep_rew_mean	348	
time/		
fps	463	
iterations	137	
time_elapsed	605	
total_timesteps	280576	
train/		
approx_kl	0.030003266	
clip_fraction	0.217	
clip_range	0.2	
entropy_loss	-21.9	
explained_variance	0.899	
learning_rate	0.0003	
loss	94	
n_updates	1360	
policy_gradient_loss	-0.0541	
std	0.878	
value_loss	309	

rollout/		
ep_len_mean	74.3	
ep_rew_mean	358	
time/		
fps	462	
iterations	138	
time_elapsed	610	

	total_timesteps		282624	
	train/			
	approx_kl		0.029054984	
	clip_fraction		0.217	
	clip_range		0.2	
	entropy_loss		-21.9	
	explained_variance		0.872	
	learning_rate		0.0003	
	loss		152	
	n_updates		1370	
	policy_gradient_loss		-0.053	
	std		0.878	
	value_loss		374	

	rollout/			
	ep_len_mean		74.5	
	ep_rew_mean		357	
	time/			
	fps		463	
	iterations		139	
	time_elapsed		614	
	total_timesteps		284672	
	train/			
	approx_kl		0.030210804	
	clip_fraction		0.232	
	clip_range		0.2	
	entropy_loss		-21.9	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		81.2	
	n_updates		1380	
	policy_gradient_loss		-0.0544	
	std		0.879	
	value_loss		264	

	rollout/			
	ep_len_mean		75.3	
	ep_rew_mean		360	
	time/			
	fps		462	
	iterations		140	
	time_elapsed		619	
	total_timesteps		286720	
	train/			
	approx_kl		0.0318602	
	clip_fraction		0.238	

clip_range	0.2	
entropy_loss	-21.9	
explained_variance	0.925	
learning_rate	0.0003	
loss	118	
n_updates	1390	
policy_gradient_loss	-0.0584	
std	0.877	
value_loss	283	

rollout/		
ep_len_mean	75.4	
ep_rew_mean	363	
time/		
fps	462	
iterations	141	
time_elapsed	623	
total_timesteps	288768	
train/		
approx_kl	0.03190822	
clip_fraction	0.216	
clip_range	0.2	
entropy_loss	-21.8	
explained_variance	0.9	
learning_rate	0.0003	
loss	140	
n_updates	1400	
policy_gradient_loss	-0.0531	
std	0.877	
value_loss	360	

rollout/		
ep_len_mean	73.1	
ep_rew_mean	351	
time/		
fps	462	
iterations	142	
time_elapsed	628	
total_timesteps	290816	
train/		
approx_kl	0.027093336	
clip_fraction	0.22	
clip_range	0.2	
entropy_loss	-21.8	
explained_variance	0.899	
learning_rate	0.0003	

	loss		87	
	n_updates		1410	
	policy_gradient_loss		-0.0551	
	std		0.877	
	value_loss		322	

	rollout/			
	ep_len_mean		75.3	
	ep_rew_mean		365	
	time/			
	fps		462	
	iterations		143	
	time_elapsed		632	
	total_timesteps		292864	
	train/			
	approx_kl		0.030916108	
	clip_fraction		0.221	
	clip_range		0.2	
	entropy_loss		-21.8	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		148	
	n_updates		1420	
	policy_gradient_loss		-0.0553	
	std		0.875	
	value_loss		304	

	rollout/			
	ep_len_mean		73.4	
	ep_rew_mean		355	
	time/			
	fps		462	
	iterations		144	
	time_elapsed		637	
	total_timesteps		294912	
	train/			
	approx_kl		0.030198023	
	clip_fraction		0.203	
	clip_range		0.2	
	entropy_loss		-21.8	
	explained_variance		0.89	
	learning_rate		0.0003	
	loss		156	
	n_updates		1430	
	policy_gradient_loss		-0.0521	
	std		0.876	

	value_loss	350	
-----			
	rollout/		
	ep_len_mean	74.6	
	ep_rew_mean	360	
	time/		
	fps	462	
	iterations	145	
	time_elapsed	642	
	total_timesteps	296960	
	train/		
	approx_kl	0.028422989	
	clip_fraction	0.223	
	clip_range	0.2	
	entropy_loss	-21.8	
	explained_variance	0.902	
	learning_rate	0.0003	
	loss	93.4	
	n_updates	1440	
	policy_gradient_loss	-0.0551	
	std	0.876	
	value_loss	309	
-----			
	rollout/		
	ep_len_mean	73.3	
	ep_rew_mean	354	
	time/		
	fps	462	
	iterations	146	
	time_elapsed	646	
	total_timesteps	299008	
	train/		
	approx_kl	0.026583541	
	clip_fraction	0.212	
	clip_range	0.2	
	entropy_loss	-21.8	
	explained_variance	0.891	
	learning_rate	0.0003	
	loss	84.6	
	n_updates	1450	
	policy_gradient_loss	-0.0522	
	std	0.877	
	value_loss	307	
-----			
	rollout/		

	ep_len_mean	75	
	ep_rew_mean	362	
	time/		
	fps	462	
	iterations	147	
	time_elapsed	650	
	total_timesteps	301056	
	train/		
	approx_kl	0.032587834	
	clip_fraction	0.226	
	clip_range	0.2	
	entropy_loss	-21.8	
	explained_variance	0.912	
	learning_rate	0.0003	
	loss	87.5	
	n_updates	1460	
	policy_gradient_loss	-0.0581	
	std	0.876	
	value_loss	272	

---

	rollout/		
	ep_len_mean	77.4	
	ep_rew_mean	371	
	time/		
	fps	462	
	iterations	148	
	time_elapsed	655	
	total_timesteps	303104	
	train/		
	approx_kl	0.03018652	
	clip_fraction	0.216	
	clip_range	0.2	
	entropy_loss	-21.8	
	explained_variance	0.904	
	learning_rate	0.0003	
	loss	127	
	n_updates	1470	
	policy_gradient_loss	-0.0535	
	std	0.875	
	value_loss	349	

---

	rollout/		
	ep_len_mean	77.9	
	ep_rew_mean	376	
	time/		
	fps	462	

iterations	149
time_elapsed	659
total_timesteps	305152
train/	
approx_kl	0.025048396
clip_fraction	0.185
clip_range	0.2
entropy_loss	-21.8
explained_variance	0.874
learning_rate	0.0003
loss	167
n_updates	1480
policy_gradient_loss	-0.0487
std	0.875
value_loss	427

rollout/	
ep_len_mean	79.3
ep_rew_mean	383
time/	
fps	462
iterations	150
time_elapsed	664
total_timesteps	307200
train/	
approx_kl	0.029722432
clip_fraction	0.24
clip_range	0.2
entropy_loss	-21.8
explained_variance	0.915
learning_rate	0.0003
loss	117
n_updates	1490
policy_gradient_loss	-0.0537
std	0.873
value_loss	301

rollout/	
ep_len_mean	79.5
ep_rew_mean	384
time/	
fps	462
iterations	151
time_elapsed	668
total_timesteps	309248
train/	

approx_kl	0.022992592	
clip_fraction	0.193	
clip_range	0.2	
entropy_loss	-21.8	
explained_variance	0.896	
learning_rate	0.0003	
loss	119	
n_updates	1500	
policy_gradient_loss	-0.0492	
std	0.873	
value_loss	364	

rollout/		
ep_len_mean	77.2	
ep_rew_mean	374	
time/		
fps	462	
iterations	152	
time_elapsed	673	
total_timesteps	311296	
train/		
approx_kl	0.030298412	
clip_fraction	0.23	
clip_range	0.2	
entropy_loss	-21.7	
explained_variance	0.908	
learning_rate	0.0003	
loss	129	
n_updates	1510	
policy_gradient_loss	-0.0559	
std	0.872	
value_loss	321	

rollout/		
ep_len_mean	76	
ep_rew_mean	369	
time/		
fps	461	
iterations	153	
time_elapsed	678	
total_timesteps	313344	
train/		
approx_kl	0.03204535	
clip_fraction	0.213	
clip_range	0.2	
entropy_loss	-21.7	



	explained_variance	0.914	
	learning_rate	0.0003	
	loss	111	
	n_updates	1520	
	policy_gradient_loss	-0.0535	
	std	0.873	
	value_loss	309	

	rollout/		
	ep_len_mean	77.8	
	ep_rew_mean	377	
	time/		
	fps	461	
	iterations	154	
	time_elapsed	683	
	total_timesteps	315392	
	train/		
	approx_kl	0.029432988	
	clip_fraction	0.209	
	clip_range	0.2	
	entropy_loss	-21.7	
	explained_variance	0.913	
	learning_rate	0.0003	
	loss	120	
	n_updates	1530	
	policy_gradient_loss	-0.0518	
	std	0.873	
	value_loss	269	

	rollout/		
	ep_len_mean	75.8	
	ep_rew_mean	367	
	time/		
	fps	461	
	iterations	155	
	time_elapsed	687	
	total_timesteps	317440	
	train/		
	approx_kl	0.032193825	
	clip_fraction	0.218	
	clip_range	0.2	
	entropy_loss	-21.7	
	explained_variance	0.897	
	learning_rate	0.0003	
	loss	74.6	
	n_updates	1540	

	policy_gradient_loss		-0.0528	
	std		0.873	
	value_loss		343	

	rollout/			
	ep_len_mean		77	
	ep_rew_mean		370	
	time/			
	fps		461	
	iterations		156	
	time_elapsed		692	
	total_timesteps		319488	
	train/			
	approx_kl		0.025263827	
	clip_fraction		0.206	
	clip_range		0.2	
	entropy_loss		-21.7	
	explained_variance		0.902	
	learning_rate		0.0003	
	loss		71.2	
	n_updates		1550	
	policy_gradient_loss		-0.0543	
	std		0.871	
	value_loss		296	

	rollout/			
	ep_len_mean		73.6	
	ep_rew_mean		351	
	time/			
	fps		460	
	iterations		157	
	time_elapsed		697	
	total_timesteps		321536	
	train/			
	approx_kl		0.034066264	
	clip_fraction		0.234	
	clip_range		0.2	
	entropy_loss		-21.7	
	explained_variance		0.908	
	learning_rate		0.0003	
	loss		72.5	
	n_updates		1560	
	policy_gradient_loss		-0.0548	
	std		0.872	
	value_loss		308	

-----		
rollout/		
ep_len_mean	74.5	
ep_rew_mean	356	
time/		
fps	460	
iterations	158	
time_elapsed	702	
total_timesteps	323584	
train/		
approx_kl	0.029259939	
clip_fraction	0.228	
clip_range	0.2	
entropy_loss	-21.7	
explained_variance	0.916	
learning_rate	0.0003	
loss	68	
n_updates	1570	
policy_gradient_loss	-0.0544	
std	0.871	
value_loss	282	
-----		

-----		
rollout/		
ep_len_mean	76.2	
ep_rew_mean	366	
time/		
fps	460	
iterations	159	
time_elapsed	707	
total_timesteps	325632	
train/		
approx_kl	0.026571747	
clip_fraction	0.212	
clip_range	0.2	
entropy_loss	-21.7	
explained_variance	0.905	
learning_rate	0.0003	
loss	104	
n_updates	1580	
policy_gradient_loss	-0.0539	
std	0.871	
value_loss	304	
-----		

-----		
rollout/		
ep_len_mean	75.7	
ep_rew_mean	366	

time/		
fps	460	
iterations	160	
time_elapsed	711	
total_timesteps	327680	
train/		
approx_kl	0.028968835	
clip_fraction	0.217	
clip_range	0.2	
entropy_loss	-21.7	
explained_variance	0.906	
learning_rate	0.0003	
loss	140	
n_updates	1590	
policy_gradient_loss	-0.0531	
std	0.871	
value_loss	335	

---

rollout/		
ep_len_mean	79.2	
ep_rew_mean	384	
time/		
fps	460	
iterations	161	
time_elapsed	716	
total_timesteps	329728	
train/		
approx_kl	0.02946002	
clip_fraction	0.228	
clip_range	0.2	
entropy_loss	-21.7	
explained_variance	0.911	
learning_rate	0.0003	
loss	136	
n_updates	1600	
policy_gradient_loss	-0.0533	
std	0.871	
value_loss	314	

---

rollout/		
ep_len_mean	76.1	
ep_rew_mean	370	
time/		
fps	460	
iterations	162	
time_elapsed	721	

	total_timesteps		331776	
	train/			
	approx_kl		0.034206618	
	clip_fraction		0.224	
	clip_range		0.2	
	entropy_loss		-21.7	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		115	
	n_updates		1610	
	policy_gradient_loss		-0.0566	
	std		0.87	
	value_loss		288	

	rollout/			
	ep_len_mean		76.2	
	ep_rew_mean		367	
	time/			
	fps		460	
	iterations		163	
	time_elapsed		725	
	total_timesteps		333824	
	train/			
	approx_kl		0.031575322	
	clip_fraction		0.233	
	clip_range		0.2	
	entropy_loss		-21.7	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		94.8	
	n_updates		1620	
	policy_gradient_loss		-0.0532	
	std		0.87	
	value_loss		315	

	rollout/			
	ep_len_mean		76.9	
	ep_rew_mean		370	
	time/			
	fps		460	
	iterations		164	
	time_elapsed		730	
	total_timesteps		335872	
	train/			
	approx_kl		0.030834097	
	clip_fraction		0.216	

clip_range	0.2	
entropy_loss	-21.7	
explained_variance	0.9	
learning_rate	0.0003	
loss	98.7	
n_updates	1630	
policy_gradient_loss	-0.0526	
std	0.869	
value_loss	287	

rollout/		
ep_len_mean	77.3	
ep_rew_mean	374	
time/		
fps	459	
iterations	165	
time_elapsed	735	
total_timesteps	337920	
train/		
approx_kl	0.032992564	
clip_fraction	0.238	
clip_range	0.2	
entropy_loss	-21.7	
explained_variance	0.908	
learning_rate	0.0003	
loss	98.6	
n_updates	1640	
policy_gradient_loss	-0.0559	
std	0.868	
value_loss	299	

rollout/		
ep_len_mean	79.1	
ep_rew_mean	381	
time/		
fps	459	
iterations	166	
time_elapsed	739	
total_timesteps	339968	
train/		
approx_kl	0.032717634	
clip_fraction	0.209	
clip_range	0.2	
entropy_loss	-21.6	
explained_variance	0.899	
learning_rate	0.0003	

	loss		135	
	n_updates		1650	
	policy_gradient_loss		-0.0516	
	std		0.868	
	value_loss		319	

---

	rollout/			
	ep_len_mean		77.4	
	ep_rew_mean		376	
	time/			
	fps		459	
	iterations		167	
	time_elapsed		743	
	total_timesteps		342016	
	train/			
	approx_kl		0.02996528	
	clip_fraction		0.246	
	clip_range		0.2	
	entropy_loss		-21.6	
	explained_variance		0.923	
	learning_rate		0.0003	
	loss		135	
	n_updates		1660	
	policy_gradient_loss		-0.0568	
	std		0.867	
	value_loss		294	

---

	rollout/			
	ep_len_mean		76.8	
	ep_rew_mean		374	
	time/			
	fps		459	
	iterations		168	
	time_elapsed		748	
	total_timesteps		344064	
	train/			
	approx_kl		0.031151615	
	clip_fraction		0.226	
	clip_range		0.2	
	entropy_loss		-21.6	
	explained_variance		0.885	
	learning_rate		0.0003	
	loss		91.8	
	n_updates		1670	
	policy_gradient_loss		-0.0538	
	std		0.867	

	value_loss		379	
-----				
	rollout/			
	ep_len_mean		77.7	
	ep_rew_mean		376	
	time/			
	fps		459	
	iterations		169	
	time_elapsed		752	
	total_timesteps		346112	
	train/			
	approx_kl		0.03397999	
	clip_fraction		0.243	
	clip_range		0.2	
	entropy_loss		-21.6	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		80.5	
	n_updates		1680	
	policy_gradient_loss		-0.055	
	std		0.867	
	value_loss		266	
-----				
	rollout/			
	ep_len_mean		78.2	
	ep_rew_mean		378	
	time/			
	fps		459	
	iterations		170	
	time_elapsed		757	
	total_timesteps		348160	
	train/			
	approx_kl		0.029604629	
	clip_fraction		0.233	
	clip_range		0.2	
	entropy_loss		-21.6	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		154	
	n_updates		1690	
	policy_gradient_loss		-0.0526	
	std		0.866	
	value_loss		322	
-----				
	rollout/			



	ep_len_mean		78.9	
	ep_rew_mean		377	
	time/			
	fps		459	
	iterations		171	
	time_elapsed		762	
	total_timesteps		350208	
	train/			
	approx_kl		0.028935228	
	clip_fraction		0.227	
	clip_range		0.2	
	entropy_loss		-21.6	
	explained_variance		0.911	
	learning_rate		0.0003	
	loss		97.7	
	n_updates		1700	
	policy_gradient_loss		-0.0551	
	std		0.867	
	value_loss		338	

	rollout/			
	ep_len_mean		78.8	
	ep_rew_mean		375	
	time/			
	fps		459	
	iterations		172	
	time_elapsed		766	
	total_timesteps		352256	
	train/			
	approx_kl		0.033061966	
	clip_fraction		0.221	
	clip_range		0.2	
	entropy_loss		-21.6	
	explained_variance		0.899	
	learning_rate		0.0003	
	loss		137	
	n_updates		1710	
	policy_gradient_loss		-0.0516	
	std		0.867	
	value_loss		322	

	rollout/			
	ep_len_mean		78	
	ep_rew_mean		372	
	time/			
	fps		459	

iterations	173
time_elapsed	770
total_timesteps	354304
train/	
approx_kl	0.036346354
clip_fraction	0.234
clip_range	0.2
entropy_loss	-21.6
explained_variance	0.917
learning_rate	0.0003
loss	130
n_updates	1720
policy_gradient_loss	-0.0549
std	0.866
value_loss	284

rollout/	
ep_len_mean	79.3
ep_rew_mean	379
time/	
fps	459
iterations	174
time_elapsed	776
total_timesteps	356352
train/	
approx_kl	0.028896514
clip_fraction	0.218
clip_range	0.2
entropy_loss	-21.6
explained_variance	0.904
learning_rate	0.0003
loss	82.6
n_updates	1730
policy_gradient_loss	-0.0501
std	0.865
value_loss	344

rollout/	
ep_len_mean	79.2
ep_rew_mean	382
time/	
fps	458
iterations	175
time_elapsed	781
total_timesteps	358400
train/	

approx_kl	0.028346056	
clip_fraction	0.223	
clip_range	0.2	
entropy_loss	-21.6	
explained_variance	0.909	
learning_rate	0.0003	
loss	127	
n_updates	1740	
policy_gradient_loss	-0.0526	
std	0.863	
value_loss	325	

rollout/		
ep_len_mean	79.6	
ep_rew_mean	385	
time/		
fps	458	
iterations	176	
time_elapsed	786	
total_timesteps	360448	
train/		
approx_kl	0.030258374	
clip_fraction	0.224	
clip_range	0.2	
entropy_loss	-21.6	
explained_variance	0.909	
learning_rate	0.0003	
loss	131	
n_updates	1750	
policy_gradient_loss	-0.0521	
std	0.863	
value_loss	298	

rollout/		
ep_len_mean	79	
ep_rew_mean	380	
time/		
fps	457	
iterations	177	
time_elapsed	792	
total_timesteps	362496	
train/		
approx_kl	0.031018503	
clip_fraction	0.237	
clip_range	0.2	
entropy_loss	-21.5	

	explained_variance	0.918	
	learning_rate	0.0003	
	loss	93.8	
	n_updates	1760	
	policy_gradient_loss	-0.0519	
	std	0.862	
	value_loss	307	

	rollout/		
	ep_len_mean	76.8	
	ep_rew_mean	371	
	time/		
	fps	457	
	iterations	178	
	time_elapsed	797	
	total_timesteps	364544	
	train/		
	approx_kl	0.031189049	
	clip_fraction	0.216	
	clip_range	0.2	
	entropy_loss	-21.5	
	explained_variance	0.918	
	learning_rate	0.0003	
	loss	98.9	
	n_updates	1770	
	policy_gradient_loss	-0.0529	
	std	0.862	
	value_loss	308	

	rollout/		
	ep_len_mean	77.1	
	ep_rew_mean	371	
	time/		
	fps	456	
	iterations	179	
	time_elapsed	802	
	total_timesteps	366592	
	train/		
	approx_kl	0.031276084	
	clip_fraction	0.226	
	clip_range	0.2	
	entropy_loss	-21.5	
	explained_variance	0.906	
	learning_rate	0.0003	
	loss	130	
	n_updates	1780	

	policy_gradient_loss		-0.0543	
	std		0.86	
	value_loss		292	

---

	rollout/			
	ep_len_mean		77.5	
	ep_rew_mean		370	
	time/			
	fps		456	
	iterations		180	
	time_elapsed		808	
	total_timesteps		368640	
	train/			
	approx_kl		0.041720662	
	clip_fraction		0.256	
	clip_range		0.2	
	entropy_loss		-21.5	
	explained_variance		0.912	
	learning_rate		0.0003	
	loss		69.8	
	n_updates		1790	
	policy_gradient_loss		-0.0546	
	std		0.86	
	value_loss		263	

---



---

	rollout/			
	ep_len_mean		77.7	
	ep_rew_mean		372	
	time/			
	fps		455	
	iterations		181	
	time_elapsed		813	
	total_timesteps		370688	
	train/			
	approx_kl		0.03210812	
	clip_fraction		0.249	
	clip_range		0.2	
	entropy_loss		-21.5	
	explained_variance		0.92	
	learning_rate		0.0003	
	loss		87.7	
	n_updates		1800	
	policy_gradient_loss		-0.0558	
	std		0.86	
	value_loss		280	

---

---

rollout/		
ep_len_mean	78.1	
ep_rew_mean	374	
time/		
fps	455	
iterations	182	
time_elapsed	818	
total_timesteps	372736	
train/		
approx_kl	0.03790813	
clip_fraction	0.226	
clip_range	0.2	
entropy_loss	-21.5	
explained_variance	0.91	
learning_rate	0.0003	
loss	106	
n_updates	1810	
policy_gradient_loss	-0.0513	
std	0.86	
value_loss	336	

---



---

rollout/		
ep_len_mean	79.9	
ep_rew_mean	382	
time/		
fps	455	
iterations	183	
time_elapsed	823	
total_timesteps	374784	
train/		
approx_kl	0.026595809	
clip_fraction	0.223	
clip_range	0.2	
entropy_loss	-21.5	
explained_variance	0.91	
learning_rate	0.0003	
loss	159	
n_updates	1820	
policy_gradient_loss	-0.0553	
std	0.859	
value_loss	329	

---



---

rollout/		
ep_len_mean	78.5	
ep_rew_mean	376	

---

time/		
fps	454	
iterations	184	
time_elapsed	828	
total_timesteps	376832	
train/		
approx_kl	0.033425964	
clip_fraction	0.244	
clip_range	0.2	
entropy_loss	-21.4	
explained_variance	0.921	
learning_rate	0.0003	
loss	117	
n_updates	1830	
policy_gradient_loss	-0.0542	
std	0.857	
value_loss	280	

---

rollout/		
ep_len_mean	78.6	
ep_rew_mean	374	
time/		
fps	454	
iterations	185	
time_elapsed	833	
total_timesteps	378880	
train/		
approx_kl	0.027968632	
clip_fraction	0.222	
clip_range	0.2	
entropy_loss	-21.4	
explained_variance	0.91	
learning_rate	0.0003	
loss	161	
n_updates	1840	
policy_gradient_loss	-0.0568	
std	0.857	
value_loss	312	

---

rollout/		
ep_len_mean	78.1	
ep_rew_mean	372	
time/		
fps	454	
iterations	186	
time_elapsed	838	

	total_timesteps		380928	
	train/			
	approx_kl		0.03528279	
	clip_fraction		0.246	
	clip_range		0.2	
	entropy_loss		-21.4	
	explained_variance		0.911	
	learning_rate		0.0003	
	loss		121	
	n_updates		1850	
	policy_gradient_loss		-0.0558	
	std		0.857	
	value_loss		312	

	rollout/			
	ep_len_mean		76.6	
	ep_rew_mean		364	
	time/			
	fps		453	
	iterations		187	
	time_elapsed		843	
	total_timesteps		382976	
	train/			
	approx_kl		0.029648945	
	clip_fraction		0.235	
	clip_range		0.2	
	entropy_loss		-21.4	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		114	
	n_updates		1860	
	policy_gradient_loss		-0.0551	
	std		0.856	
	value_loss		292	

	rollout/			
	ep_len_mean		79	
	ep_rew_mean		379	
	time/			
	fps		453	
	iterations		188	
	time_elapsed		849	
	total_timesteps		385024	
	train/			
	approx_kl		0.031163404	
	clip_fraction		0.231	



clip_range	0.2	
entropy_loss	-21.4	
explained_variance	0.908	
learning_rate	0.0003	
loss	89.4	
n_updates	1870	
policy_gradient_loss	-0.0542	
std	0.856	
value_loss	283	

rollout/		
ep_len_mean	78.7	
ep_rew_mean	380	
time/		
fps	453	
iterations	189	
time_elapsed	854	
total_timesteps	387072	
train/		
approx_kl	0.034457047	
clip_fraction	0.25	
clip_range	0.2	
entropy_loss	-21.4	
explained_variance	0.931	
learning_rate	0.0003	
loss	81.7	
n_updates	1880	
policy_gradient_loss	-0.0568	
std	0.854	
value_loss	231	

rollout/		
ep_len_mean	79.4	
ep_rew_mean	383	
time/		
fps	452	
iterations	190	
time_elapsed	859	
total_timesteps	389120	
train/		
approx_kl	0.02900227	
clip_fraction	0.229	
clip_range	0.2	
entropy_loss	-21.4	
explained_variance	0.902	
learning_rate	0.0003	

	loss		114	
	n_updates		1890	
	policy_gradient_loss		-0.0527	
	std		0.853	
	value_loss		318	

	rollout/			
	ep_len_mean		78.7	
	ep_rew_mean		380	
	time/			
	fps		452	
	iterations		191	
	time_elapsed		864	
	total_timesteps		391168	
	train/			
	approx_kl		0.031739857	
	clip_fraction		0.243	
	clip_range		0.2	
	entropy_loss		-21.3	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		109	
	n_updates		1900	
	policy_gradient_loss		-0.0546	
	std		0.852	
	value_loss		260	

	rollout/			
	ep_len_mean		74.7	
	ep_rew_mean		360	
	time/			
	fps		452	
	iterations		192	
	time_elapsed		869	
	total_timesteps		393216	
	train/			
	approx_kl		0.033156045	
	clip_fraction		0.224	
	clip_range		0.2	
	entropy_loss		-21.3	
	explained_variance		0.914	
	learning_rate		0.0003	
	loss		80.7	
	n_updates		1910	
	policy_gradient_loss		-0.0512	
	std		0.852	

	value_loss	311	
-----			
	rollout/		
	ep_len_mean	74.8	
	ep_rew_mean	360	
	time/		
	fps	451	
	iterations	193	
	time_elapsed	874	
	total_timesteps	395264	
	train/		
	approx_kl	0.035968557	
	clip_fraction	0.243	
	clip_range	0.2	
	entropy_loss	-21.3	
	explained_variance	0.918	
	learning_rate	0.0003	
	loss	115	
	n_updates	1920	
	policy_gradient_loss	-0.0546	
	std	0.851	
	value_loss	300	
-----			
	rollout/		
	ep_len_mean	74	
	ep_rew_mean	358	
	time/		
	fps	451	
	iterations	194	
	time_elapsed	879	
	total_timesteps	397312	
	train/		
	approx_kl	0.03860105	
	clip_fraction	0.291	
	clip_range	0.2	
	entropy_loss	-21.3	
	explained_variance	0.925	
	learning_rate	0.0003	
	loss	91.4	
	n_updates	1930	
	policy_gradient_loss	-0.0626	
	std	0.85	
	value_loss	244	
-----			
	rollout/		

	ep_len_mean		75.6	
	ep_rew_mean		363	
	time/			
	fps		451	
	iterations		195	
	time_elapsed		885	
	total_timesteps		399360	
	train/			
	approx_kl		0.040775407	
	clip_fraction		0.279	
	clip_range		0.2	
	entropy_loss		-21.3	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		86.9	
	n_updates		1940	
	policy_gradient_loss		-0.0602	
	std		0.849	
	value_loss		223	

	rollout/			
	ep_len_mean		77.5	
	ep_rew_mean		371	
	time/			
	fps		450	
	iterations		196	
	time_elapsed		891	
	total_timesteps		401408	
	train/			
	approx_kl		0.034253523	
	clip_fraction		0.253	
	clip_range		0.2	
	entropy_loss		-21.3	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		90.2	
	n_updates		1950	
	policy_gradient_loss		-0.0558	
	std		0.849	
	value_loss		271	

	rollout/			
	ep_len_mean		77.7	
	ep_rew_mean		373	
	time/			
	fps		450	

iterations	197
time_elapsed	896
total_timesteps	403456
train/	
approx_kl	0.031229533
clip_fraction	0.218
clip_range	0.2
entropy_loss	-21.2
explained_variance	0.906
learning_rate	0.0003
loss	97
n_updates	1960
policy_gradient_loss	-0.0528
std	0.848
value_loss	295

rollout/	
ep_len_mean	78.6
ep_rew_mean	375
time/	
fps	450
iterations	198
time_elapsed	901
total_timesteps	405504
train/	
approx_kl	0.03749676
clip_fraction	0.248
clip_range	0.2
entropy_loss	-21.2
explained_variance	0.923
learning_rate	0.0003
loss	64.3
n_updates	1970
policy_gradient_loss	-0.0549
std	0.847
value_loss	233

rollout/	
ep_len_mean	78.8
ep_rew_mean	378
time/	
fps	450
iterations	199
time_elapsed	905
total_timesteps	407552
train/	

approx_kl	0.03764484	
clip_fraction	0.248	
clip_range	0.2	
entropy_loss	-21.2	
explained_variance	0.902	
learning_rate	0.0003	
loss	100	
n_updates	1980	
policy_gradient_loss	-0.0523	
std	0.846	
value_loss	318	

rollout/		
ep_len_mean	81.3	
ep_rew_mean	391	
time/		
fps	450	
iterations	200	
time_elapsed	909	
total_timesteps	409600	
train/		
approx_kl	0.035831038	
clip_fraction	0.251	
clip_range	0.2	
entropy_loss	-21.2	
explained_variance	0.898	
learning_rate	0.0003	
loss	140	
n_updates	1990	
policy_gradient_loss	-0.0557	
std	0.845	
value_loss	329	

rollout/		
ep_len_mean	81.9	
ep_rew_mean	393	
time/		
fps	450	
iterations	201	
time_elapsed	914	
total_timesteps	411648	
train/		
approx_kl	0.03343161	
clip_fraction	0.241	
clip_range	0.2	
entropy_loss	-21.2	

	explained_variance	0.914	
	learning_rate	0.0003	
	loss	115	
	n_updates	2000	
	policy_gradient_loss	-0.0554	
	std	0.844	
	value_loss	321	

	rollout/		
	ep_len_mean	80.2	
	ep_rew_mean	388	
	time/		
	fps	449	
	iterations	202	
	time_elapsed	919	
	total_timesteps	413696	
	train/		
	approx_kl	0.030496046	
	clip_fraction	0.233	
	clip_range	0.2	
	entropy_loss	-21.2	
	explained_variance	0.911	
	learning_rate	0.0003	
	loss	115	
	n_updates	2010	
	policy_gradient_loss	-0.0534	
	std	0.844	
	value_loss	312	

	rollout/		
	ep_len_mean	79.5	
	ep_rew_mean	386	
	time/		
	fps	449	
	iterations	203	
	time_elapsed	924	
	total_timesteps	415744	
	train/		
	approx_kl	0.036470413	
	clip_fraction	0.257	
	clip_range	0.2	
	entropy_loss	-21.1	
	explained_variance	0.928	
	learning_rate	0.0003	
	loss	113	
	n_updates	2020	

	policy_gradient_loss		-0.0603	
	std		0.842	
	value_loss		265	

	rollout/			
	ep_len_mean		78.5	
	ep_rew_mean		383	
	time/			
	fps		449	
	iterations		204	
	time_elapsed		929	
	total_timesteps		417792	
	train/			
	approx_kl		0.042701513	
	clip_fraction		0.279	
	clip_range		0.2	
	entropy_loss		-21.1	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		68	
	n_updates		2030	
	policy_gradient_loss		-0.0577	
	std		0.841	
	value_loss		223	

	rollout/			
	ep_len_mean		77.7	
	ep_rew_mean		380	
	time/			
	fps		449	
	iterations		205	
	time_elapsed		933	
	total_timesteps		419840	
	train/			
	approx_kl		0.040131405	
	clip_fraction		0.278	
	clip_range		0.2	
	entropy_loss		-21.1	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		66.9	
	n_updates		2040	
	policy_gradient_loss		-0.0586	
	std		0.84	
	value_loss		250	



---

rollout/		
ep_len_mean	78.7	
ep_rew_mean	383	
time/		
fps	449	
iterations	206	
time_elapsed	938	
total_timesteps	421888	
train/		
approx_kl	0.035361134	
clip_fraction	0.255	
clip_range	0.2	
entropy_loss	-21.1	
explained_variance	0.912	
learning_rate	0.0003	
loss	105	
n_updates	2050	
policy_gradient_loss	-0.0553	
std	0.839	
value_loss	336	

---



---

rollout/		
ep_len_mean	80.5	
ep_rew_mean	391	
time/		
fps	449	
iterations	207	
time_elapsed	942	
total_timesteps	423936	
train/		
approx_kl	0.04103043	
clip_fraction	0.254	
clip_range	0.2	
entropy_loss	-21	
explained_variance	0.912	
learning_rate	0.0003	
loss	122	
n_updates	2060	
policy_gradient_loss	-0.0562	
std	0.839	
value_loss	293	

---



---

rollout/		
ep_len_mean	80.2	
ep_rew_mean	387	

---

time/		
fps	449	
iterations	208	
time_elapsed	947	
total_timesteps	425984	
train/		
approx_kl	0.036934413	
clip_fraction	0.247	
clip_range	0.2	
entropy_loss	-21	
explained_variance	0.887	
learning_rate	0.0003	
loss	102	
n_updates	2070	
policy_gradient_loss	-0.0555	
std	0.838	
value_loss	329	

---

rollout/		
ep_len_mean	80.5	
ep_rew_mean	387	
time/		
fps	449	
iterations	209	
time_elapsed	952	
total_timesteps	428032	
train/		
approx_kl	0.03651303	
clip_fraction	0.26	
clip_range	0.2	
entropy_loss	-21	
explained_variance	0.915	
learning_rate	0.0003	
loss	99.5	
n_updates	2080	
policy_gradient_loss	-0.0583	
std	0.837	
value_loss	298	

---

rollout/		
ep_len_mean	80.4	
ep_rew_mean	385	
time/		
fps	449	
iterations	210	
time_elapsed	956	

	total_timesteps		430080	
	train/			
	approx_kl		0.039375164	
	clip_fraction		0.267	
	clip_range		0.2	
	entropy_loss		-21	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		103	
	n_updates		2090	
	policy_gradient_loss		-0.0568	
	std		0.836	
	value_loss		317	

	rollout/			
	ep_len_mean		78.5	
	ep_rew_mean		374	
	time/			
	fps		449	
	iterations		211	
	time_elapsed		961	
	total_timesteps		432128	
	train/			
	approx_kl		0.043419942	
	clip_fraction		0.268	
	clip_range		0.2	
	entropy_loss		-21	
	explained_variance		0.922	
	learning_rate		0.0003	
	loss		98.1	
	n_updates		2100	
	policy_gradient_loss		-0.056	
	std		0.836	
	value_loss		245	

	rollout/			
	ep_len_mean		79.3	
	ep_rew_mean		378	
	time/			
	fps		449	
	iterations		212	
	time_elapsed		966	
	total_timesteps		434176	
	train/			
	approx_kl		0.04389548	
	clip_fraction		0.286	

clip_range	0.2	
entropy_loss	-21	
explained_variance	0.922	
learning_rate	0.0003	
loss	69.1	
n_updates	2110	
policy_gradient_loss	-0.0604	
std	0.837	
value_loss	230	

rollout/		
ep_len_mean	77.8	
ep_rew_mean	374	
time/		
fps	449	
iterations	213	
time_elapsed	971	
total_timesteps	436224	
train/		
approx_kl	0.036449656	
clip_fraction	0.254	
clip_range	0.2	
entropy_loss	-21	
explained_variance	0.89	
learning_rate	0.0003	
loss	104	
n_updates	2120	
policy_gradient_loss	-0.0585	
std	0.836	
value_loss	320	

rollout/		
ep_len_mean	78.2	
ep_rew_mean	380	
time/		
fps	449	
iterations	214	
time_elapsed	975	
total_timesteps	438272	
train/		
approx_kl	0.044456724	
clip_fraction	0.312	
clip_range	0.2	
entropy_loss	-21	
explained_variance	0.919	
learning_rate	0.0003	

	loss		47.5	
	n_updates		2130	
	policy_gradient_loss		-0.0643	
	std		0.836	
	value_loss		252	

	rollout/			
	ep_len_mean		78.6	
	ep_rew_mean		384	
	time/			
	fps		448	
	iterations		215	
	time_elapsed		980	
	total_timesteps		440320	
	train/			
	approx_kl		0.039162196	
	clip_fraction		0.273	
	clip_range		0.2	
	entropy_loss		-21	
	explained_variance		0.92	
	learning_rate		0.0003	
	loss		84.6	
	n_updates		2140	
	policy_gradient_loss		-0.0582	
	std		0.836	
	value_loss		251	

	rollout/			
	ep_len_mean		76.4	
	ep_rew_mean		373	
	time/			
	fps		448	
	iterations		216	
	time_elapsed		985	
	total_timesteps		442368	
	train/			
	approx_kl		0.03722413	
	clip_fraction		0.255	
	clip_range		0.2	
	entropy_loss		-21	
	explained_variance		0.906	
	learning_rate		0.0003	
	loss		125	
	n_updates		2150	
	policy_gradient_loss		-0.0577	
	std		0.836	

	value_loss	346	
-----			
	rollout/		
	ep_len_mean	75.6	
	ep_rew_mean	370	
	time/		
	fps	449	
	iterations	217	
	time_elapsed	989	
	total_timesteps	444416	
	train/		
	approx_kl	0.039952926	
	clip_fraction	0.266	
	clip_range	0.2	
	entropy_loss	-21	
	explained_variance	0.914	
	learning_rate	0.0003	
	loss	89.2	
	n_updates	2160	
	policy_gradient_loss	-0.0607	
	std	0.836	
	value_loss	327	
-----			
	rollout/		
	ep_len_mean	76.7	
	ep_rew_mean	375	
	time/		
	fps	449	
	iterations	218	
	time_elapsed	993	
	total_timesteps	446464	
	train/		
	approx_kl	0.04004877	
	clip_fraction	0.26	
	clip_range	0.2	
	entropy_loss	-21	
	explained_variance	0.918	
	learning_rate	0.0003	
	loss	71.3	
	n_updates	2170	
	policy_gradient_loss	-0.0573	
	std	0.836	
	value_loss	277	
-----			
	rollout/		

	ep_len_mean		78.9	
	ep_rew_mean		383	
	time/			
	fps		449	
	iterations		219	
	time_elapsed		998	
	total_timesteps		448512	
	train/			
	approx_kl		0.043807954	
	clip_fraction		0.258	
	clip_range		0.2	
	entropy_loss		-21	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		104	
	n_updates		2180	
	policy_gradient_loss		-0.0534	
	std		0.836	
	value_loss		287	

	rollout/			
	ep_len_mean		80.7	
	ep_rew_mean		395	
	time/			
	fps		449	
	iterations		220	
	time_elapsed		1003	
	total_timesteps		450560	
	train/			
	approx_kl		0.035356358	
	clip_fraction		0.233	
	clip_range		0.2	
	entropy_loss		-21	
	explained_variance		0.918	
	learning_rate		0.0003	
	loss		149	
	n_updates		2190	
	policy_gradient_loss		-0.0558	
	std		0.835	
	value_loss		299	

	rollout/			
	ep_len_mean		81.2	
	ep_rew_mean		396	
	time/			
	fps		448	

iterations	221
time_elapsed	1008
total_timesteps	452608
train/	
approx_kl	0.036986336
clip_fraction	0.242
clip_range	0.2
entropy_loss	-20.9
explained_variance	0.915
learning_rate	0.0003
loss	84.2
n_updates	2200
policy_gradient_loss	-0.0504
std	0.833
value_loss	293

rollout/	
ep_len_mean	81.1
ep_rew_mean	392
time/	
fps	448
iterations	222
time_elapsed	1014
total_timesteps	454656
train/	
approx_kl	0.041824576
clip_fraction	0.256
clip_range	0.2
entropy_loss	-20.9
explained_variance	0.923
learning_rate	0.0003
loss	137
n_updates	2210
policy_gradient_loss	-0.0575
std	0.832
value_loss	271

rollout/	
ep_len_mean	79
ep_rew_mean	380
time/	
fps	448
iterations	223
time_elapsed	1018
total_timesteps	456704
train/	



approx_kl	0.03782315	
clip_fraction	0.252	
clip_range	0.2	
entropy_loss	-20.9	
explained_variance	0.916	
learning_rate	0.0003	
loss	103	
n_updates	2220	
policy_gradient_loss	-0.0579	
std	0.831	
value_loss	322	

rollout/		
ep_len_mean	79.2	
ep_rew_mean	379	
time/		
fps	448	
iterations	224	
time_elapsed	1023	
total_timesteps	458752	
train/		
approx_kl	0.037564635	
clip_fraction	0.241	
clip_range	0.2	
entropy_loss	-20.9	
explained_variance	0.906	
learning_rate	0.0003	
loss	180	
n_updates	2230	
policy_gradient_loss	-0.0538	
std	0.831	
value_loss	352	

rollout/		
ep_len_mean	79.9	
ep_rew_mean	384	
time/		
fps	448	
iterations	225	
time_elapsed	1027	
total_timesteps	460800	
train/		
approx_kl	0.039805338	
clip_fraction	0.246	
clip_range	0.2	
entropy_loss	-20.9	

	explained_variance		0.895	
	learning_rate		0.0003	
	loss		90.4	
	n_updates		2240	
	policy_gradient_loss		-0.054	
	std		0.83	
	value_loss		313	

	rollout/			
	ep_len_mean		79.3	
	ep_rew_mean		381	
	time/			
	fps		448	
	iterations		226	
	time_elapsed		1032	
	total_timesteps		462848	
	train/			
	approx_kl		0.032377426	
	clip_fraction		0.251	
	clip_range		0.2	
	entropy_loss		-20.8	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		125	
	n_updates		2250	
	policy_gradient_loss		-0.0582	
	std		0.829	
	value_loss		304	

	rollout/			
	ep_len_mean		79.3	
	ep_rew_mean		384	
	time/			
	fps		448	
	iterations		227	
	time_elapsed		1037	
	total_timesteps		464896	
	train/			
	approx_kl		0.03644625	
	clip_fraction		0.276	
	clip_range		0.2	
	entropy_loss		-20.8	
	explained_variance		0.911	
	learning_rate		0.0003	
	loss		84.8	
	n_updates		2260	

	policy_gradient_loss		-0.0587	
	std		0.828	
	value_loss		301	

	rollout/			
	ep_len_mean		80.1	
	ep_rew_mean		389	
	time/			
	fps		447	
	iterations		228	
	time_elapsed		1043	
	total_timesteps		466944	
	train/			
	approx_kl		0.042365782	
	clip_fraction		0.268	
	clip_range		0.2	
	entropy_loss		-20.8	
	explained_variance		0.915	
	learning_rate		0.0003	
	loss		122	
	n_updates		2270	
	policy_gradient_loss		-0.0539	
	std		0.828	
	value_loss		314	

	rollout/			
	ep_len_mean		78.6	
	ep_rew_mean		383	
	time/			
	fps		447	
	iterations		229	
	time_elapsed		1048	
	total_timesteps		468992	
	train/			
	approx_kl		0.038717337	
	clip_fraction		0.258	
	clip_range		0.2	
	entropy_loss		-20.8	
	explained_variance		0.911	
	learning_rate		0.0003	
	loss		116	
	n_updates		2280	
	policy_gradient_loss		-0.0537	
	std		0.829	
	value_loss		326	

---

rollout/		
ep_len_mean	77.5	
ep_rew_mean	378	
time/		
fps	447	
iterations	230	
time_elapsed	1053	
total_timesteps	471040	
train/		
approx_kl	0.039766356	
clip_fraction	0.242	
clip_range	0.2	
entropy_loss	-20.8	
explained_variance	0.922	
learning_rate	0.0003	
loss	89.2	
n_updates	2290	
policy_gradient_loss	-0.0521	
std	0.827	
value_loss	259	

---



---

rollout/		
ep_len_mean	76.9	
ep_rew_mean	374	
time/		
fps	446	
iterations	231	
time_elapsed	1058	
total_timesteps	473088	
train/		
approx_kl	0.036680274	
clip_fraction	0.265	
clip_range	0.2	
entropy_loss	-20.8	
explained_variance	0.906	
learning_rate	0.0003	
loss	137	
n_updates	2300	
policy_gradient_loss	-0.057	
std	0.826	
value_loss	314	

---



---

rollout/		
ep_len_mean	75.8	
ep_rew_mean	367	

---

time/		
fps	446	
iterations	232	
time_elapsed	1063	
total_timesteps	475136	
train/		
approx_kl	0.036717903	
clip_fraction	0.257	
clip_range	0.2	
entropy_loss	-20.8	
explained_variance	0.9	
learning_rate	0.0003	
loss	116	
n_updates	2310	
policy_gradient_loss	-0.054	
std	0.826	
value_loss	325	

---

rollout/		
ep_len_mean	76.5	
ep_rew_mean	371	
time/		
fps	446	
iterations	233	
time_elapsed	1068	
total_timesteps	477184	
train/		
approx_kl	0.045576	
clip_fraction	0.312	
clip_range	0.2	
entropy_loss	-20.7	
explained_variance	0.914	
learning_rate	0.0003	
loss	110	
n_updates	2320	
policy_gradient_loss	-0.0645	
std	0.824	
value_loss	256	

---

rollout/		
ep_len_mean	77	
ep_rew_mean	375	
time/		
fps	446	
iterations	234	
time_elapsed	1073	

	total_timesteps		479232	
	train/			
	approx_kl		0.045172043	
	clip_fraction		0.309	
	clip_range		0.2	
	entropy_loss		-20.7	
	explained_variance		0.927	
	learning_rate		0.0003	
	loss		101	
	n_updates		2330	
	policy_gradient_loss		-0.0624	
	std		0.823	
	value_loss		230	

	rollout/			
	ep_len_mean		77.2	
	ep_rew_mean		374	
	time/			
	fps		446	
	iterations		235	
	time_elapsed		1077	
	total_timesteps		481280	
	train/			
	approx_kl		0.04221381	
	clip_fraction		0.265	
	clip_range		0.2	
	entropy_loss		-20.7	
	explained_variance		0.919	
	learning_rate		0.0003	
	loss		115	
	n_updates		2340	
	policy_gradient_loss		-0.053	
	std		0.823	
	value_loss		261	

	rollout/			
	ep_len_mean		79.5	
	ep_rew_mean		384	
	time/			
	fps		446	
	iterations		236	
	time_elapsed		1082	
	total_timesteps		483328	
	train/			
	approx_kl		0.035580035	
	clip_fraction		0.244	

clip_range	0.2	
entropy_loss	-20.7	
explained_variance	0.904	
learning_rate	0.0003	
loss	101	
n_updates	2350	
policy_gradient_loss	-0.0524	
std	0.822	
value_loss	352	

rollout/		
ep_len_mean	78	
ep_rew_mean	374	
time/		
fps	446	
iterations	237	
time_elapsed	1087	
total_timesteps	485376	
train/		
approx_kl	0.040306874	
clip_fraction	0.258	
clip_range	0.2	
entropy_loss	-20.7	
explained_variance	0.905	
learning_rate	0.0003	
loss	104	
n_updates	2360	
policy_gradient_loss	-0.0547	
std	0.822	
value_loss	286	

rollout/		
ep_len_mean	80	
ep_rew_mean	383	
time/		
fps	446	
iterations	238	
time_elapsed	1091	
total_timesteps	487424	
train/		
approx_kl	0.040272985	
clip_fraction	0.272	
clip_range	0.2	
entropy_loss	-20.7	
explained_variance	0.898	
learning_rate	0.0003	

	loss		93.1	
	n_updates		2370	
	policy_gradient_loss		-0.059	
	std		0.821	
	value_loss		295	

	rollout/			
	ep_len_mean		80.4	
	ep_rew_mean		386	
	time/			
	fps		446	
	iterations		239	
	time_elapsed		1096	
	total_timesteps		489472	
	train/			
	approx_kl		0.047345385	
	clip_fraction		0.265	
	clip_range		0.2	
	entropy_loss		-20.7	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		77.5	
	n_updates		2380	
	policy_gradient_loss		-0.0575	
	std		0.82	
	value_loss		254	

	rollout/			
	ep_len_mean		80.5	
	ep_rew_mean		388	
	time/			
	fps		446	
	iterations		240	
	time_elapsed		1101	
	total_timesteps		491520	
	train/			
	approx_kl		0.04589551	
	clip_fraction		0.266	
	clip_range		0.2	
	entropy_loss		-20.6	
	explained_variance		0.918	
	learning_rate		0.0003	
	loss		96.5	
	n_updates		2390	
	policy_gradient_loss		-0.0514	
	std		0.819	



	value_loss		269	
-----				
	rollout/			
	ep_len_mean		82.6	
	ep_rew_mean		399	
	time/			
	fps		446	
	iterations		241	
	time_elapsed		1106	
	total_timesteps		493568	
	train/			
	approx_kl		0.043148782	
	clip_fraction		0.285	
	clip_range		0.2	
	entropy_loss		-20.6	
	explained_variance		0.914	
	learning_rate		0.0003	
	loss		72	
	n_updates		2400	
	policy_gradient_loss		-0.0584	
	std		0.82	
	value_loss		282	
-----				
	rollout/			
	ep_len_mean		81.8	
	ep_rew_mean		395	
	time/			
	fps		445	
	iterations		242	
	time_elapsed		1111	
	total_timesteps		495616	
	train/			
	approx_kl		0.03665711	
	clip_fraction		0.259	
	clip_range		0.2	
	entropy_loss		-20.6	
	explained_variance		0.912	
	learning_rate		0.0003	
	loss		111	
	n_updates		2410	
	policy_gradient_loss		-0.0553	
	std		0.82	
	value_loss		304	
-----				
	rollout/			

	ep_len_mean		81.9	
	ep_rew_mean		397	
	time/			
	fps		445	
	iterations		243	
	time_elapsed		1116	
	total_timesteps		497664	
	train/			
	approx_kl		0.047176614	
	clip_fraction		0.267	
	clip_range		0.2	
	entropy_loss		-20.6	
	explained_variance		0.91	
	learning_rate		0.0003	
	loss		102	
	n_updates		2420	
	policy_gradient_loss		-0.0558	
	std		0.819	
	value_loss		328	

	rollout/			
	ep_len_mean		80.8	
	ep_rew_mean		391	
	time/			
	fps		445	
	iterations		244	
	time_elapsed		1121	
	total_timesteps		499712	
	train/			
	approx_kl		0.035340432	
	clip_fraction		0.248	
	clip_range		0.2	
	entropy_loss		-20.6	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		76.2	
	n_updates		2430	
	policy_gradient_loss		-0.0566	
	std		0.818	
	value_loss		279	

	rollout/			
	ep_len_mean		80.2	
	ep_rew_mean		391	
	time/			
	fps		445	

iterations	245
time_elapsed	1126
total_timesteps	501760
train/	
approx_kl	0.03206832
clip_fraction	0.22
clip_range	0.2
entropy_loss	-20.6
explained_variance	0.907
learning_rate	0.0003
loss	147
n_updates	2440
policy_gradient_loss	-0.0525
std	0.818
value_loss	354

rollout/	
ep_len_mean	79.7
ep_rew_mean	388
time/	
fps	445
iterations	246
time_elapsed	1131
total_timesteps	503808
train/	
approx_kl	0.039365485
clip_fraction	0.264
clip_range	0.2
entropy_loss	-20.6
explained_variance	0.907
learning_rate	0.0003
loss	106
n_updates	2450
policy_gradient_loss	-0.0548
std	0.818
value_loss	303

rollout/	
ep_len_mean	82.1
ep_rew_mean	398
time/	
fps	445
iterations	247
time_elapsed	1136
total_timesteps	505856
train/	

approx_kl	0.040284425	
clip_fraction	0.261	
clip_range	0.2	
entropy_loss	-20.6	
explained_variance	0.919	
learning_rate	0.0003	
loss	78.2	
n_updates	2460	
policy_gradient_loss	-0.0586	
std	0.817	
value_loss	268	

rollout/		
ep_len_mean	84.4	
ep_rew_mean	410	
time/		
fps	445	
iterations	248	
time_elapsed	1141	
total_timesteps	507904	
train/		
approx_kl	0.039393913	
clip_fraction	0.269	
clip_range	0.2	
entropy_loss	-20.6	
explained_variance	0.917	
learning_rate	0.0003	
loss	113	
n_updates	2470	
policy_gradient_loss	-0.0585	
std	0.817	
value_loss	323	

rollout/		
ep_len_mean	85.5	
ep_rew_mean	414	
time/		
fps	445	
iterations	249	
time_elapsed	1145	
total_timesteps	509952	
train/		
approx_kl	0.039502326	
clip_fraction	0.271	
clip_range	0.2	
entropy_loss	-20.6	

	explained_variance		0.911	
	learning_rate		0.0003	
	loss		124	
	n_updates		2480	
	policy_gradient_loss		-0.0572	
	std		0.815	
	value_loss		317	

	rollout/			
	ep_len_mean		85.8	
	ep_rew_mean		416	
	time/			
	fps		445	
	iterations		250	
	time_elapsed		1150	
	total_timesteps		512000	
	train/			
	approx_kl		0.050802983	
	clip_fraction		0.284	
	clip_range		0.2	
	entropy_loss		-20.5	
	explained_variance		0.919	
	learning_rate		0.0003	
	loss		75.7	
	n_updates		2490	
	policy_gradient_loss		-0.0556	
	std		0.814	
	value_loss		275	

	rollout/			
	ep_len_mean		84.5	
	ep_rew_mean		409	
	time/			
	fps		445	
	iterations		251	
	time_elapsed		1154	
	total_timesteps		514048	
	train/			
	approx_kl		0.040700644	
	clip_fraction		0.244	
	clip_range		0.2	
	entropy_loss		-20.5	
	explained_variance		0.921	
	learning_rate		0.0003	
	loss		77.8	
	n_updates		2500	

	policy_gradient_loss		-0.0556	
	std		0.814	
	value_loss		339	

---

	rollout/			
	ep_len_mean		83.5	
	ep_rew_mean		405	
	time/			
	fps		445	
	iterations		252	
	time_elapsed		1159	
	total_timesteps		516096	
	train/			
	approx_kl		0.04178782	
	clip_fraction		0.26	
	clip_range		0.2	
	entropy_loss		-20.5	
	explained_variance		0.919	
	learning_rate		0.0003	
	loss		150	
	n_updates		2510	
	policy_gradient_loss		-0.0527	
	std		0.814	
	value_loss		326	

---



---

	rollout/			
	ep_len_mean		82.7	
	ep_rew_mean		400	
	time/			
	fps		445	
	iterations		253	
	time_elapsed		1163	
	total_timesteps		518144	
	train/			
	approx_kl		0.035444934	
	clip_fraction		0.241	
	clip_range		0.2	
	entropy_loss		-20.5	
	explained_variance		0.896	
	learning_rate		0.0003	
	loss		80.2	
	n_updates		2520	
	policy_gradient_loss		-0.0552	
	std		0.813	
	value_loss		357	

---

---

rollout/		
ep_len_mean	82	
ep_rew_mean	396	
time/		
fps	445	
iterations	254	
time_elapsed	1168	
total_timesteps	520192	
train/		
approx_kl	0.042498272	
clip_fraction	0.264	
clip_range	0.2	
entropy_loss	-20.5	
explained_variance	0.927	
learning_rate	0.0003	
loss	106	
n_updates	2530	
policy_gradient_loss	-0.0577	
std	0.813	
value_loss	282	

---



---

rollout/		
ep_len_mean	83	
ep_rew_mean	401	
time/		
fps	445	
iterations	255	
time_elapsed	1172	
total_timesteps	522240	
train/		
approx_kl	0.041766584	
clip_fraction	0.269	
clip_range	0.2	
entropy_loss	-20.5	
explained_variance	0.919	
learning_rate	0.0003	
loss	147	
n_updates	2540	
policy_gradient_loss	-0.0535	
std	0.812	
value_loss	290	

---



---

rollout/		
ep_len_mean	82.2	
ep_rew_mean	398	

---

time/		
fps	445	
iterations	256	
time_elapsed	1177	
total_timesteps	524288	
train/		
approx_kl	0.043848693	
clip_fraction	0.279	
clip_range	0.2	
entropy_loss	-20.5	
explained_variance	0.92	
learning_rate	0.0003	
loss	115	
n_updates	2550	
policy_gradient_loss	-0.0576	
std	0.812	
value_loss	294	

rollout/		
ep_len_mean	80.1	
ep_rew_mean	390	
time/		
fps	445	
iterations	257	
time_elapsed	1182	
total_timesteps	526336	
train/		
approx_kl	0.038723737	
clip_fraction	0.252	
clip_range	0.2	
entropy_loss	-20.5	
explained_variance	0.898	
learning_rate	0.0003	
loss	84.1	
n_updates	2560	
policy_gradient_loss	-0.0557	
std	0.811	
value_loss	363	

rollout/		
ep_len_mean	78.8	
ep_rew_mean	383	
time/		
fps	445	
iterations	258	
time_elapsed	1187	



	total_timesteps		528384	
	train/			
	approx_kl		0.050209146	
	clip_fraction		0.298	
	clip_range		0.2	
	entropy_loss		-20.5	
	explained_variance		0.908	
	learning_rate		0.0003	
	loss		160	
	n_updates		2570	
	policy_gradient_loss		-0.0552	
	std		0.81	
	value_loss		337	

	rollout/			
	ep_len_mean		77.7	
	ep_rew_mean		377	
	time/			
	fps		445	
	iterations		259	
	time_elapsed		1191	
	total_timesteps		530432	
	train/			
	approx_kl		0.05248661	
	clip_fraction		0.303	
	clip_range		0.2	
	entropy_loss		-20.4	
	explained_variance		0.923	
	learning_rate		0.0003	
	loss		69	
	n_updates		2580	
	policy_gradient_loss		-0.059	
	std		0.81	
	value_loss		261	

	rollout/			
	ep_len_mean		76.3	
	ep_rew_mean		372	
	time/			
	fps		445	
	iterations		260	
	time_elapsed		1196	
	total_timesteps		532480	
	train/			
	approx_kl		0.040545113	
	clip_fraction		0.268	

clip_range	0.2	
entropy_loss	-20.4	
explained_variance	0.911	
learning_rate	0.0003	
loss	101	
n_updates	2590	
policy_gradient_loss	-0.0565	
std	0.81	
value_loss	275	

rollout/		
ep_len_mean	78.1	
ep_rew_mean	382	
time/		
fps	444	
iterations	261	
time_elapsed	1201	
total_timesteps	534528	
train/		
approx_kl	0.061318714	
clip_fraction	0.312	
clip_range	0.2	
entropy_loss	-20.4	
explained_variance	0.927	
learning_rate	0.0003	
loss	78.3	
n_updates	2600	
policy_gradient_loss	-0.0622	
std	0.809	
value_loss	236	

rollout/		
ep_len_mean	78.7	
ep_rew_mean	386	
time/		
fps	444	
iterations	262	
time_elapsed	1206	
total_timesteps	536576	
train/		
approx_kl	0.053059697	
clip_fraction	0.282	
clip_range	0.2	
entropy_loss	-20.4	
explained_variance	0.908	
learning_rate	0.0003	

	loss		144	
	n_updates		2610	
	policy_gradient_loss		-0.0573	
	std		0.809	
	value_loss		311	

	rollout/			
	ep_len_mean		78.6	
	ep_rew_mean		386	
	time/			
	fps		444	
	iterations		263	
	time_elapsed		1210	
	total_timesteps		538624	
	train/			
	approx_kl		0.042539995	
	clip_fraction		0.285	
	clip_range		0.2	
	entropy_loss		-20.4	
	explained_variance		0.919	
	learning_rate		0.0003	
	loss		109	
	n_updates		2620	
	policy_gradient_loss		-0.0598	
	std		0.808	
	value_loss		286	

	rollout/			
	ep_len_mean		79.7	
	ep_rew_mean		387	
	time/			
	fps		445	
	iterations		264	
	time_elapsed		1214	
	total_timesteps		540672	
	train/			
	approx_kl		0.039041094	
	clip_fraction		0.275	
	clip_range		0.2	
	entropy_loss		-20.4	
	explained_variance		0.91	
	learning_rate		0.0003	
	loss		105	
	n_updates		2630	
	policy_gradient_loss		-0.0547	
	std		0.807	

	value_loss	283	
-----			
	rollout/		
	ep_len_mean	79.4	
	ep_rew_mean	385	
	time/		
	fps	445	
	iterations	265	
	time_elapsed	1219	
	total_timesteps	542720	
	train/		
	approx_kl	0.042524703	
	clip_fraction	0.285	
	clip_range	0.2	
	entropy_loss	-20.4	
	explained_variance	0.91	
	learning_rate	0.0003	
	loss	70.6	
	n_updates	2640	
	policy_gradient_loss	-0.0589	
	std	0.806	
	value_loss	279	
-----			
	rollout/		
	ep_len_mean	81.5	
	ep_rew_mean	393	
	time/		
	fps	445	
	iterations	266	
	time_elapsed	1223	
	total_timesteps	544768	
	train/		
	approx_kl	0.04056418	
	clip_fraction	0.285	
	clip_range	0.2	
	entropy_loss	-20.3	
	explained_variance	0.915	
	learning_rate	0.0003	
	loss	84.2	
	n_updates	2650	
	policy_gradient_loss	-0.0577	
	std	0.805	
	value_loss	294	
-----			
	rollout/		

	ep_len_mean		82	
	ep_rew_mean		395	
	time/			
	fps		445	
	iterations		267	
	time_elapsed		1228	
	total_timesteps		546816	
	train/			
	approx_kl		0.035000592	
	clip_fraction		0.254	
	clip_range		0.2	
	entropy_loss		-20.3	
	explained_variance		0.893	
	learning_rate		0.0003	
	loss		85.1	
	n_updates		2660	
	policy_gradient_loss		-0.0538	
	std		0.803	
	value_loss		363	

	rollout/			
	ep_len_mean		80.5	
	ep_rew_mean		390	
	time/			
	fps		445	
	iterations		268	
	time_elapsed		1232	
	total_timesteps		548864	
	train/			
	approx_kl		0.048445947	
	clip_fraction		0.294	
	clip_range		0.2	
	entropy_loss		-20.3	
	explained_variance		0.922	
	learning_rate		0.0003	
	loss		96.9	
	n_updates		2670	
	policy_gradient_loss		-0.0601	
	std		0.802	
	value_loss		276	

	rollout/			
	ep_len_mean		81.2	
	ep_rew_mean		395	
	time/			
	fps		445	

iterations	269
time_elapsed	1237
total_timesteps	550912
train/	
approx_kl	0.04771092
clip_fraction	0.294
clip_range	0.2
entropy_loss	-20.3
explained_variance	0.919
learning_rate	0.0003
loss	57.6
n_updates	2680
policy_gradient_loss	-0.056
std	0.802
value_loss	259

rollout/	
ep_len_mean	80.5
ep_rew_mean	394
time/	
fps	445
iterations	270
time_elapsed	1242
total_timesteps	552960
train/	
approx_kl	0.03797114
clip_fraction	0.266
clip_range	0.2
entropy_loss	-20.3
explained_variance	0.922
learning_rate	0.0003
loss	91.8
n_updates	2690
policy_gradient_loss	-0.0563
std	0.802
value_loss	261

rollout/	
ep_len_mean	81.4
ep_rew_mean	398
time/	
fps	444
iterations	271
time_elapsed	1247
total_timesteps	555008
train/	

approx_kl	0.048810564	
clip_fraction	0.296	
clip_range	0.2	
entropy_loss	-20.3	
explained_variance	0.926	
learning_rate	0.0003	
loss	100	
n_updates	2700	
policy_gradient_loss	-0.0564	
std	0.802	
value_loss	248	

rollout/		
ep_len_mean	83.1	
ep_rew_mean	408	
time/		
fps	444	
iterations	272	
time_elapsed	1252	
total_timesteps	557056	
train/		
approx_kl	0.041476738	
clip_fraction	0.283	
clip_range	0.2	
entropy_loss	-20.2	
explained_variance	0.928	
learning_rate	0.0003	
loss	114	
n_updates	2710	
policy_gradient_loss	-0.0582	
std	0.801	
value_loss	279	

rollout/		
ep_len_mean	83.3	
ep_rew_mean	409	
time/		
fps	444	
iterations	273	
time_elapsed	1256	
total_timesteps	559104	
train/		
approx_kl	0.05464532	
clip_fraction	0.277	
clip_range	0.2	
entropy_loss	-20.2	

	explained_variance	0.917	
	learning_rate	0.0003	
	loss	98.7	
	n_updates	2720	
	policy_gradient_loss	-0.0568	
	std	0.8	
	value_loss	304	

	rollout/		
	ep_len_mean	82.8	
	ep_rew_mean	402	
	time/		
	fps	444	
	iterations	274	
	time_elapsed	1261	
	total_timesteps	561152	
	train/		
	approx_kl	0.04468993	
	clip_fraction	0.275	
	clip_range	0.2	
	entropy_loss	-20.2	
	explained_variance	0.911	
	learning_rate	0.0003	
	loss	93.4	
	n_updates	2730	
	policy_gradient_loss	-0.0585	
	std	0.8	
	value_loss	317	

	rollout/		
	ep_len_mean	82.5	
	ep_rew_mean	403	
	time/		
	fps	444	
	iterations	275	
	time_elapsed	1266	
	total_timesteps	563200	
	train/		
	approx_kl	0.044219438	
	clip_fraction	0.274	
	clip_range	0.2	
	entropy_loss	-20.2	
	explained_variance	0.909	
	learning_rate	0.0003	
	loss	129	
	n_updates	2740	



	policy_gradient_loss		-0.0564	
	std		0.8	
	value_loss		314	

	rollout/			
	ep_len_mean		82.2	
	ep_rew_mean		398	
	time/			
	fps		444	
	iterations		276	
	time_elapsed		1271	
	total_timesteps		565248	
	train/			
	approx_kl		0.048757993	
	clip_fraction		0.303	
	clip_range		0.2	
	entropy_loss		-20.2	
	explained_variance		0.92	
	learning_rate		0.0003	
	loss		115	
	n_updates		2750	
	policy_gradient_loss		-0.0604	
	std		0.8	
	value_loss		284	

	rollout/			
	ep_len_mean		81.6	
	ep_rew_mean		393	
	time/			
	fps		444	
	iterations		277	
	time_elapsed		1276	
	total_timesteps		567296	
	train/			
	approx_kl		0.040865634	
	clip_fraction		0.271	
	clip_range		0.2	
	entropy_loss		-20.2	
	explained_variance		0.922	
	learning_rate		0.0003	
	loss		83.9	
	n_updates		2760	
	policy_gradient_loss		-0.0585	
	std		0.799	
	value_loss		287	

---

rollout/		
ep_len_mean	82.1	
ep_rew_mean	399	
time/		
fps	444	
iterations	278	
time_elapsed	1280	
total_timesteps	569344	
train/		
approx_kl	0.04948973	
clip_fraction	0.306	
clip_range	0.2	
entropy_loss	-20.2	
explained_variance	0.924	
learning_rate	0.0003	
loss	92.9	
n_updates	2770	
policy_gradient_loss	-0.0569	
std	0.799	
value_loss	259	

---



---

rollout/		
ep_len_mean	81	
ep_rew_mean	391	
time/		
fps	444	
iterations	279	
time_elapsed	1285	
total_timesteps	571392	
train/		
approx_kl	0.042761404	
clip_fraction	0.255	
clip_range	0.2	
entropy_loss	-20.2	
explained_variance	0.915	
learning_rate	0.0003	
loss	84.7	
n_updates	2780	
policy_gradient_loss	-0.0536	
std	0.8	
value_loss	330	

---



---

rollout/		
ep_len_mean	81.1	
ep_rew_mean	393	

---

time/		
fps	444	
iterations	280	
time_elapsed	1289	
total_timesteps	573440	
train/		
approx_kl	0.040996682	
clip_fraction	0.267	
clip_range	0.2	
entropy_loss	-20.2	
explained_variance	0.922	
learning_rate	0.0003	
loss	112	
n_updates	2790	
policy_gradient_loss	-0.0478	
std	0.799	
value_loss	307	

---

rollout/		
ep_len_mean	82.6	
ep_rew_mean	401	
time/		
fps	444	
iterations	281	
time_elapsed	1294	
total_timesteps	575488	
train/		
approx_kl	0.04446674	
clip_fraction	0.28	
clip_range	0.2	
entropy_loss	-20.2	
explained_variance	0.92	
learning_rate	0.0003	
loss	180	
n_updates	2800	
policy_gradient_loss	-0.0575	
std	0.799	
value_loss	323	

---

rollout/		
ep_len_mean	80.9	
ep_rew_mean	392	
time/		
fps	444	
iterations	282	
time_elapsed	1299	

	total_timesteps		577536	
	train/			
	approx_kl		0.04453574	
	clip_fraction		0.272	
	clip_range		0.2	
	entropy_loss		-20.2	
	explained_variance		0.907	
	learning_rate		0.0003	
	loss		113	
	n_updates		2810	
	policy_gradient_loss		-0.0568	
	std		0.799	
	value_loss		325	

	rollout/			
	ep_len_mean		81.7	
	ep_rew_mean		398	
	time/			
	fps		444	
	iterations		283	
	time_elapsed		1304	
	total_timesteps		579584	
	train/			
	approx_kl		0.050051652	
	clip_fraction		0.287	
	clip_range		0.2	
	entropy_loss		-20.2	
	explained_variance		0.926	
	learning_rate		0.0003	
	loss		103	
	n_updates		2820	
	policy_gradient_loss		-0.0575	
	std		0.798	
	value_loss		250	

	rollout/			
	ep_len_mean		81.2	
	ep_rew_mean		396	
	time/			
	fps		444	
	iterations		284	
	time_elapsed		1308	
	total_timesteps		581632	
	train/			
	approx_kl		0.049693346	
	clip_fraction		0.326	

clip_range	0.2	
entropy_loss	-20.2	
explained_variance	0.911	
learning_rate	0.0003	
loss	80.2	
n_updates	2830	
policy_gradient_loss	-0.0612	
std	0.798	
value_loss	269	

rollout/		
ep_len_mean	81.4	
ep_rew_mean	398	
time/		
fps	444	
iterations	285	
time_elapsed	1313	
total_timesteps	583680	
train/		
approx_kl	0.048082463	
clip_fraction	0.307	
clip_range	0.2	
entropy_loss	-20.1	
explained_variance	0.929	
learning_rate	0.0003	
loss	95.7	
n_updates	2840	
policy_gradient_loss	-0.0556	
std	0.797	
value_loss	244	

rollout/		
ep_len_mean	81.9	
ep_rew_mean	404	
time/		
fps	444	
iterations	286	
time_elapsed	1318	
total_timesteps	585728	
train/		
approx_kl	0.042706348	
clip_fraction	0.273	
clip_range	0.2	
entropy_loss	-20.1	
explained_variance	0.894	
learning_rate	0.0003	

	loss		193	
	n_updates		2850	
	policy_gradient_loss		-0.0542	
	std		0.795	
	value_loss		361	

	rollout/			
	ep_len_mean		81.1	
	ep_rew_mean		400	
	time/			
	fps		444	
	iterations		287	
	time_elapsed		1323	
	total_timesteps		587776	
	train/			
	approx_kl		0.050529197	
	clip_fraction		0.316	
	clip_range		0.2	
	entropy_loss		-20.1	
	explained_variance		0.925	
	learning_rate		0.0003	
	loss		125	
	n_updates		2860	
	policy_gradient_loss		-0.0616	
	std		0.794	
	value_loss		262	

	rollout/			
	ep_len_mean		80	
	ep_rew_mean		396	
	time/			
	fps		443	
	iterations		288	
	time_elapsed		1328	
	total_timesteps		589824	
	train/			
	approx_kl		0.04099112	
	clip_fraction		0.277	
	clip_range		0.2	
	entropy_loss		-20.1	
	explained_variance		0.904	
	learning_rate		0.0003	
	loss		150	
	n_updates		2870	
	policy_gradient_loss		-0.0545	
	std		0.794	

	value_loss		375	
-----				
	rollout/			
	ep_len_mean		81.6	
	ep_rew_mean		402	
	time/			
	fps		443	
	iterations		289	
	time_elapsed		1334	
	total_timesteps		591872	
	train/			
	approx_kl		0.05161968	
	clip_fraction		0.313	
	clip_range		0.2	
	entropy_loss		-20.1	
	explained_variance		0.918	
	learning_rate		0.0003	
	loss		73.1	
	n_updates		2880	
	policy_gradient_loss		-0.0599	
	std		0.793	
	value_loss		248	
-----				
	rollout/			
	ep_len_mean		82.5	
	ep_rew_mean		403	
	time/			
	fps		443	
	iterations		290	
	time_elapsed		1338	
	total_timesteps		593920	
	train/			
	approx_kl		0.047628447	
	clip_fraction		0.277	
	clip_range		0.2	
	entropy_loss		-20	
	explained_variance		0.89	
	learning_rate		0.0003	
	loss		193	
	n_updates		2890	
	policy_gradient_loss		-0.0577	
	std		0.791	
	value_loss		393	
-----				
	rollout/			

	ep_len_mean		86	
	ep_rew_mean		419	
	time/			
	fps		443	
	iterations		291	
	time_elapsed		1343	
	total_timesteps		595968	
	train/			
	approx_kl		0.036687147	
	clip_fraction		0.255	
	clip_range		0.2	
	entropy_loss		-20	
	explained_variance		0.881	
	learning_rate		0.0003	
	loss		169	
	n_updates		2900	
	policy_gradient_loss		-0.0524	
	std		0.791	
	value_loss		423	

	rollout/			
	ep_len_mean		85.7	
	ep_rew_mean		413	
	time/			
	fps		443	
	iterations		292	
	time_elapsed		1348	
	total_timesteps		598016	
	train/			
	approx_kl		0.046347342	
	clip_fraction		0.277	
	clip_range		0.2	
	entropy_loss		-20	
	explained_variance		0.901	
	learning_rate		0.0003	
	loss		132	
	n_updates		2910	
	policy_gradient_loss		-0.0565	
	std		0.791	
	value_loss		325	

	rollout/			
	ep_len_mean		84.9	
	ep_rew_mean		410	
	time/			
	fps		443	



iterations	293
time_elapsed	1353
total_timesteps	600064
train/	
approx_kl	0.051867083
clip_fraction	0.292
clip_range	0.2
entropy_loss	-20
explained_variance	0.92
learning_rate	0.0003
loss	98.2
n_updates	2920
policy_gradient_loss	-0.0574
std	0.791
value_loss	302

rollout/	
ep_len_mean	87.6
ep_rew_mean	423
time/	
fps	443
iterations	294
time_elapsed	1358
total_timesteps	602112
train/	
approx_kl	0.041285664
clip_fraction	0.27
clip_range	0.2
entropy_loss	-20
explained_variance	0.886
learning_rate	0.0003
loss	106
n_updates	2930
policy_gradient_loss	-0.0532
std	0.79
value_loss	417

rollout/	
ep_len_mean	85
ep_rew_mean	409
time/	
fps	443
iterations	295
time_elapsed	1362
total_timesteps	604160
train/	

approx_kl	0.053128675	
clip_fraction	0.298	
clip_range	0.2	
entropy_loss	-20	
explained_variance	0.892	
learning_rate	0.0003	
loss	104	
n_updates	2940	
policy_gradient_loss	-0.0575	
std	0.789	
value_loss	358	

rollout/		
ep_len_mean	84.8	
ep_rew_mean	413	
time/		
fps	443	
iterations	296	
time_elapsed	1367	
total_timesteps	606208	
train/		
approx_kl	0.043215174	
clip_fraction	0.279	
clip_range	0.2	
entropy_loss	-19.9	
explained_variance	0.883	
learning_rate	0.0003	
loss	143	
n_updates	2950	
policy_gradient_loss	-0.0575	
std	0.787	
value_loss	402	

rollout/		
ep_len_mean	83.5	
ep_rew_mean	405	
time/		
fps	443	
iterations	297	
time_elapsed	1372	
total_timesteps	608256	
train/		
approx_kl	0.04986246	
clip_fraction	0.292	
clip_range	0.2	
entropy_loss	-19.9	

	explained_variance		0.906	
	learning_rate		0.0003	
	loss		79.1	
	n_updates		2960	
	policy_gradient_loss		-0.0574	
	std		0.786	
	value_loss		352	

	rollout/			
	ep_len_mean		82.5	
	ep_rew_mean		405	
	time/			
	fps		443	
	iterations		298	
	time_elapsed		1377	
	total_timesteps		610304	
	train/			
	approx_kl		0.044506602	
	clip_fraction		0.279	
	clip_range		0.2	
	entropy_loss		-19.9	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		64.8	
	n_updates		2970	
	policy_gradient_loss		-0.0569	
	std		0.785	
	value_loss		324	

	rollout/			
	ep_len_mean		80.8	
	ep_rew_mean		397	
	time/			
	fps		443	
	iterations		299	
	time_elapsed		1381	
	total_timesteps		612352	
	train/			
	approx_kl		0.047854	
	clip_fraction		0.307	
	clip_range		0.2	
	entropy_loss		-19.9	
	explained_variance		0.909	
	learning_rate		0.0003	
	loss		120	
	n_updates		2980	

	policy_gradient_loss		-0.0599	
	std		0.785	
	value_loss		307	

	rollout/			
	ep_len_mean		82.4	
	ep_rew_mean		402	
	time/			
	fps		443	
	iterations		300	
	time_elapsed		1386	
	total_timesteps		614400	
	train/			
	approx_kl		0.04673706	
	clip_fraction		0.294	
	clip_range		0.2	
	entropy_loss		-19.9	
	explained_variance		0.904	
	learning_rate		0.0003	
	loss		71.1	
	n_updates		2990	
	policy_gradient_loss		-0.0564	
	std		0.785	
	value_loss		324	

	rollout/			
	ep_len_mean		81.8	
	ep_rew_mean		398	
	time/			
	fps		443	
	iterations		301	
	time_elapsed		1390	
	total_timesteps		616448	
	train/			
	approx_kl		0.05304559	
	clip_fraction		0.295	
	clip_range		0.2	
	entropy_loss		-19.9	
	explained_variance		0.912	
	learning_rate		0.0003	
	loss		120	
	n_updates		3000	
	policy_gradient_loss		-0.0545	
	std		0.785	
	value_loss		306	

---

rollout/		
ep_len_mean	83	
ep_rew_mean	404	
time/		
fps	442	
iterations	302	
time_elapsed	1396	
total_timesteps	618496	
train/		
approx_kl	0.050448764	
clip_fraction	0.296	
clip_range	0.2	
entropy_loss	-19.9	
explained_variance	0.913	
learning_rate	0.0003	
loss	80.2	
n_updates	3010	
policy_gradient_loss	-0.0569	
std	0.784	
value_loss	347	

---



---

rollout/		
ep_len_mean	83.9	
ep_rew_mean	409	
time/		
fps	442	
iterations	303	
time_elapsed	1401	
total_timesteps	620544	
train/		
approx_kl	0.0516478	
clip_fraction	0.306	
clip_range	0.2	
entropy_loss	-19.8	
explained_variance	0.905	
learning_rate	0.0003	
loss	115	
n_updates	3020	
policy_gradient_loss	-0.0576	
std	0.784	
value_loss	309	

---



---

rollout/		
ep_len_mean	82.9	
ep_rew_mean	404	

---

time/		
fps	442	
iterations	304	
time_elapsed	1405	
total_timesteps	622592	
train/		
approx_kl	0.052967496	
clip_fraction	0.311	
clip_range	0.2	
entropy_loss	-19.8	
explained_variance	0.927	
learning_rate	0.0003	
loss	131	
n_updates	3030	
policy_gradient_loss	-0.0607	
std	0.784	
value_loss	306	

---

rollout/		
ep_len_mean	82.7	
ep_rew_mean	405	
time/		
fps	442	
iterations	305	
time_elapsed	1410	
total_timesteps	624640	
train/		
approx_kl	0.04339763	
clip_fraction	0.286	
clip_range	0.2	
entropy_loss	-19.8	
explained_variance	0.917	
learning_rate	0.0003	
loss	130	
n_updates	3040	
policy_gradient_loss	-0.0606	
std	0.783	
value_loss	348	

---

rollout/		
ep_len_mean	82	
ep_rew_mean	398	
time/		
fps	442	
iterations	306	
time_elapsed	1414	

	total_timesteps		626688	
	train/			
	approx_kl		0.055925384	
	clip_fraction		0.309	
	clip_range		0.2	
	entropy_loss		-19.8	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		109	
	n_updates		3050	
	policy_gradient_loss		-0.0588	
	std		0.782	
	value_loss		314	

	rollout/			
	ep_len_mean		82.1	
	ep_rew_mean		398	
	time/			
	fps		442	
	iterations		307	
	time_elapsed		1419	
	total_timesteps		628736	
	train/			
	approx_kl		0.04505144	
	clip_fraction		0.283	
	clip_range		0.2	
	entropy_loss		-19.8	
	explained_variance		0.909	
	learning_rate		0.0003	
	loss		140	
	n_updates		3060	
	policy_gradient_loss		-0.0577	
	std		0.781	
	value_loss		346	

	rollout/			
	ep_len_mean		81.8	
	ep_rew_mean		399	
	time/			
	fps		442	
	iterations		308	
	time_elapsed		1423	
	total_timesteps		630784	
	train/			
	approx_kl		0.050501477	
	clip_fraction		0.293	

clip_range	0.2	
entropy_loss	-19.8	
explained_variance	0.884	
learning_rate	0.0003	
loss	105	
n_updates	3070	
policy_gradient_loss	-0.0549	
std	0.781	
value_loss	338	

rollout/		
ep_len_mean	82.4	
ep_rew_mean	401	
time/		
fps	442	
iterations	309	
time_elapsed	1428	
total_timesteps	632832	
train/		
approx_kl	0.061516523	
clip_fraction	0.33	
clip_range	0.2	
entropy_loss	-19.8	
explained_variance	0.932	
learning_rate	0.0003	
loss	64.1	
n_updates	3080	
policy_gradient_loss	-0.0556	
std	0.78	
value_loss	238	

rollout/		
ep_len_mean	79.9	
ep_rew_mean	388	
time/		
fps	442	
iterations	310	
time_elapsed	1433	
total_timesteps	634880	
train/		
approx_kl	0.046365023	
clip_fraction	0.299	
clip_range	0.2	
entropy_loss	-19.8	
explained_variance	0.923	
learning_rate	0.0003	



	loss		117	
	n_updates		3090	
	policy_gradient_loss		-0.0583	
	std		0.78	
	value_loss		298	

	rollout/			
	ep_len_mean		78.8	
	ep_rew_mean		381	
	time/			
	fps		442	
	iterations		311	
	time_elapsed		1438	
	total_timesteps		636928	
	train/			
	approx_kl		0.051737256	
	clip_fraction		0.295	
	clip_range		0.2	
	entropy_loss		-19.8	
	explained_variance		0.922	
	learning_rate		0.0003	
	loss		94.3	
	n_updates		3100	
	policy_gradient_loss		-0.0575	
	std		0.779	
	value_loss		289	

	rollout/			
	ep_len_mean		77	
	ep_rew_mean		375	
	time/			
	fps		442	
	iterations		312	
	time_elapsed		1443	
	total_timesteps		638976	
	train/			
	approx_kl		0.05587723	
	clip_fraction		0.306	
	clip_range		0.2	
	entropy_loss		-19.7	
	explained_variance		0.918	
	learning_rate		0.0003	
	loss		81.5	
	n_updates		3110	
	policy_gradient_loss		-0.053	
	std		0.778	

	value_loss		280	
-----				
	rollout/			
	ep_len_mean		79.3	
	ep_rew_mean		389	
	time/			
	fps		442	
	iterations		313	
	time_elapsed		1449	
	total_timesteps		641024	
	train/			
	approx_kl		0.05594056	
	clip_fraction		0.334	
	clip_range		0.2	
	entropy_loss		-19.7	
	explained_variance		0.925	
	learning_rate		0.0003	
	loss		81.7	
	n_updates		3120	
	policy_gradient_loss		-0.0605	
	std		0.776	
	value_loss		238	
-----				
	rollout/			
	ep_len_mean		81.5	
	ep_rew_mean		399	
	time/			
	fps		442	
	iterations		314	
	time_elapsed		1454	
	total_timesteps		643072	
	train/			
	approx_kl		0.046400867	
	clip_fraction		0.294	
	clip_range		0.2	
	entropy_loss		-19.7	
	explained_variance		0.911	
	learning_rate		0.0003	
	loss		141	
	n_updates		3130	
	policy_gradient_loss		-0.0592	
	std		0.776	
	value_loss		331	
-----				
	rollout/			

	ep_len_mean		81.9	
	ep_rew_mean		404	
	time/			
	fps		442	
	iterations		315	
	time_elapsed		1459	
	total_timesteps		645120	
	train/			
	approx_kl		0.052340727	
	clip_fraction		0.302	
	clip_range		0.2	
	entropy_loss		-19.7	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		121	
	n_updates		3140	
	policy_gradient_loss		-0.0607	
	std		0.775	
	value_loss		297	

	rollout/			
	ep_len_mean		82.5	
	ep_rew_mean		405	
	time/			
	fps		442	
	iterations		316	
	time_elapsed		1463	
	total_timesteps		647168	
	train/			
	approx_kl		0.050901216	
	clip_fraction		0.311	
	clip_range		0.2	
	entropy_loss		-19.6	
	explained_variance		0.918	
	learning_rate		0.0003	
	loss		84.8	
	n_updates		3150	
	policy_gradient_loss		-0.0574	
	std		0.774	
	value_loss		270	

	rollout/			
	ep_len_mean		80.8	
	ep_rew_mean		396	
	time/			
	fps		441	

iterations	317
time_elapsed	1469
total_timesteps	649216
train/	
approx_kl	0.056847863
clip_fraction	0.323
clip_range	0.2
entropy_loss	-19.6
explained_variance	0.92
learning_rate	0.0003
loss	144
n_updates	3160
policy_gradient_loss	-0.0619
std	0.772
value_loss	252

rollout/	
ep_len_mean	80.3
ep_rew_mean	397
time/	
fps	441
iterations	318
time_elapsed	1473
total_timesteps	651264
train/	
approx_kl	0.057725646
clip_fraction	0.325
clip_range	0.2
entropy_loss	-19.6
explained_variance	0.927
learning_rate	0.0003
loss	85.1
n_updates	3170
policy_gradient_loss	-0.0623
std	0.771
value_loss	258

rollout/	
ep_len_mean	81.2
ep_rew_mean	400
time/	
fps	441
iterations	319
time_elapsed	1478
total_timesteps	653312
train/	

approx_kl	0.05589725	
clip_fraction	0.309	
clip_range	0.2	
entropy_loss	-19.6	
explained_variance	0.918	
learning_rate	0.0003	
loss	134	
n_updates	3180	
policy_gradient_loss	-0.0579	
std	0.77	
value_loss	293	

rollout/		
ep_len_mean	82.2	
ep_rew_mean	403	
time/		
fps	441	
iterations	320	
time_elapsed	1483	
total_timesteps	655360	
train/		
approx_kl	0.04914855	
clip_fraction	0.284	
clip_range	0.2	
entropy_loss	-19.5	
explained_variance	0.899	
learning_rate	0.0003	
loss	187	
n_updates	3190	
policy_gradient_loss	-0.0537	
std	0.77	
value_loss	367	

rollout/		
ep_len_mean	83.6	
ep_rew_mean	409	
time/		
fps	441	
iterations	321	
time_elapsed	1488	
total_timesteps	657408	
train/		
approx_kl	0.050477326	
clip_fraction	0.293	
clip_range	0.2	
entropy_loss	-19.5	

	explained_variance		0.908	
	learning_rate		0.0003	
	loss		110	
	n_updates		3200	
	policy_gradient_loss		-0.0595	
	std		0.769	
	value_loss		358	

	rollout/			
	ep_len_mean		82.8	
	ep_rew_mean		402	
	time/			
	fps		441	
	iterations		322	
	time_elapsed		1492	
	total_timesteps		659456	
	train/			
	approx_kl		0.05654657	
	clip_fraction		0.3	
	clip_range		0.2	
	entropy_loss		-19.5	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		157	
	n_updates		3210	
	policy_gradient_loss		-0.0588	
	std		0.767	
	value_loss		304	

	rollout/			
	ep_len_mean		84.9	
	ep_rew_mean		412	
	time/			
	fps		441	
	iterations		323	
	time_elapsed		1497	
	total_timesteps		661504	
	train/			
	approx_kl		0.06508553	
	clip_fraction		0.332	
	clip_range		0.2	
	entropy_loss		-19.5	
	explained_variance		0.908	
	learning_rate		0.0003	
	loss		66.5	
	n_updates		3220	

	policy_gradient_loss		-0.0604	
	std		0.767	
	value_loss		289	

	rollout/			
	ep_len_mean		84.9	
	ep_rew_mean		415	
	time/			
	fps		441	
	iterations		324	
	time_elapsed		1502	
	total_timesteps		663552	
	train/			
	approx_kl		0.05246591	
	clip_fraction		0.317	
	clip_range		0.2	
	entropy_loss		-19.5	
	explained_variance		0.903	
	learning_rate		0.0003	
	loss		119	
	n_updates		3230	
	policy_gradient_loss		-0.0577	
	std		0.767	
	value_loss		313	

	rollout/			
	ep_len_mean		85.2	
	ep_rew_mean		419	
	time/			
	fps		441	
	iterations		325	
	time_elapsed		1506	
	total_timesteps		665600	
	train/			
	approx_kl		0.05142635	
	clip_fraction		0.309	
	clip_range		0.2	
	entropy_loss		-19.5	
	explained_variance		0.908	
	learning_rate		0.0003	
	loss		145	
	n_updates		3240	
	policy_gradient_loss		-0.0576	
	std		0.766	
	value_loss		312	

-----		
rollout/		
ep_len_mean	85.4	
ep_rew_mean	420	
time/		
fps	441	
iterations	326	
time_elapsed	1511	
total_timesteps	667648	
train/		
approx_kl	0.051016696	
clip_fraction	0.299	
clip_range	0.2	
entropy_loss	-19.5	
explained_variance	0.916	
learning_rate	0.0003	
loss	108	
n_updates	3250	
policy_gradient_loss	-0.0567	
std	0.766	
value_loss	323	
-----		

-----		
rollout/		
ep_len_mean	82.7	
ep_rew_mean	404	
time/		
fps	441	
iterations	327	
time_elapsed	1516	
total_timesteps	669696	
train/		
approx_kl	0.045645185	
clip_fraction	0.299	
clip_range	0.2	
entropy_loss	-19.5	
explained_variance	0.909	
learning_rate	0.0003	
loss	110	
n_updates	3260	
policy_gradient_loss	-0.0542	
std	0.766	
value_loss	320	
-----		

-----		
rollout/		
ep_len_mean	83.3	
ep_rew_mean	407	



time/		
fps	441	
iterations	328	
time_elapsed	1521	
total_timesteps	671744	
train/		
approx_kl	0.060217503	
clip_fraction	0.333	
clip_range	0.2	
entropy_loss	-19.4	
explained_variance	0.928	
learning_rate	0.0003	
loss	116	
n_updates	3270	
policy_gradient_loss	-0.0635	
std	0.765	
value_loss	280	

---

rollout/		
ep_len_mean	83.2	
ep_rew_mean	410	
time/		
fps	441	
iterations	329	
time_elapsed	1525	
total_timesteps	673792	
train/		
approx_kl	0.046235237	
clip_fraction	0.302	
clip_range	0.2	
entropy_loss	-19.4	
explained_variance	0.907	
learning_rate	0.0003	
loss	133	
n_updates	3280	
policy_gradient_loss	-0.0566	
std	0.765	
value_loss	355	

---

rollout/		
ep_len_mean	85	
ep_rew_mean	421	
time/		
fps	441	
iterations	330	
time_elapsed	1530	

	total_timesteps		675840	
	train/			
	approx_kl		0.05757811	
	clip_fraction		0.299	
	clip_range		0.2	
	entropy_loss		-19.4	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		118	
	n_updates		3290	
	policy_gradient_loss		-0.0562	
	std		0.764	
	value_loss		335	

	rollout/			
	ep_len_mean		85.5	
	ep_rew_mean		424	
	time/			
	fps		441	
	iterations		331	
	time_elapsed		1535	
	total_timesteps		677888	
	train/			
	approx_kl		0.040320173	
	clip_fraction		0.272	
	clip_range		0.2	
	entropy_loss		-19.4	
	explained_variance		0.903	
	learning_rate		0.0003	
	loss		123	
	n_updates		3300	
	policy_gradient_loss		-0.0556	
	std		0.764	
	value_loss		407	

	rollout/			
	ep_len_mean		85.7	
	ep_rew_mean		426	
	time/			
	fps		441	
	iterations		332	
	time_elapsed		1539	
	total_timesteps		679936	
	train/			
	approx_kl		0.042214822	
	clip_fraction		0.279	

clip_range	0.2	
entropy_loss	-19.4	
explained_variance	0.906	
learning_rate	0.0003	
loss	118	
n_updates	3310	
policy_gradient_loss	-0.0536	
std	0.765	
value_loss	395	

rollout/		
ep_len_mean	85.1	
ep_rew_mean	420	
time/		
fps	441	
iterations	333	
time_elapsed	1544	
total_timesteps	681984	
train/		
approx_kl	0.047563687	
clip_fraction	0.279	
clip_range	0.2	
entropy_loss	-19.4	
explained_variance	0.912	
learning_rate	0.0003	
loss	197	
n_updates	3320	
policy_gradient_loss	-0.0547	
std	0.766	
value_loss	376	

rollout/		
ep_len_mean	81.9	
ep_rew_mean	401	
time/		
fps	441	
iterations	334	
time_elapsed	1549	
total_timesteps	684032	
train/		
approx_kl	0.051516123	
clip_fraction	0.29	
clip_range	0.2	
entropy_loss	-19.4	
explained_variance	0.896	
learning_rate	0.0003	

	loss		60.7	
	n_updates		3330	
	policy_gradient_loss		-0.0516	
	std		0.764	
	value_loss		351	

---

	rollout/			
	ep_len_mean		81.5	
	ep_rew_mean		401	
	time/			
	fps		441	
	iterations		335	
	time_elapsed		1554	
	total_timesteps		686080	
	train/			
	approx_kl		0.05019036	
	clip_fraction		0.287	
	clip_range		0.2	
	entropy_loss		-19.4	
	explained_variance		0.9	
	learning_rate		0.0003	
	loss		121	
	n_updates		3340	
	policy_gradient_loss		-0.055	
	std		0.764	
	value_loss		371	

---



---

	rollout/			
	ep_len_mean		81	
	ep_rew_mean		400	
	time/			
	fps		441	
	iterations		336	
	time_elapsed		1558	
	total_timesteps		688128	
	train/			
	approx_kl		0.060526036	
	clip_fraction		0.335	
	clip_range		0.2	
	entropy_loss		-19.4	
	explained_variance		0.907	
	learning_rate		0.0003	
	loss		90.2	
	n_updates		3350	
	policy_gradient_loss		-0.0594	
	std		0.763	

	value_loss		293	
-----				
	rollout/			
	ep_len_mean		83.7	
	ep_rew_mean		413	
	time/			
	fps		441	
	iterations		337	
	time_elapsed		1564	
	total_timesteps		690176	
	train/			
	approx_kl		0.055101484	
	clip_fraction		0.303	
	clip_range		0.2	
	entropy_loss		-19.4	
	explained_variance		0.901	
	learning_rate		0.0003	
	loss		86.9	
	n_updates		3360	
	policy_gradient_loss		-0.0568	
	std		0.762	
	value_loss		345	
-----				
	rollout/			
	ep_len_mean		82.7	
	ep_rew_mean		410	
	time/			
	fps		441	
	iterations		338	
	time_elapsed		1569	
	total_timesteps		692224	
	train/			
	approx_kl		0.04883127	
	clip_fraction		0.297	
	clip_range		0.2	
	entropy_loss		-19.3	
	explained_variance		0.886	
	learning_rate		0.0003	
	loss		112	
	n_updates		3370	
	policy_gradient_loss		-0.0581	
	std		0.761	
	value_loss		390	
-----				
	rollout/			

	ep_len_mean		82.7	
	ep_rew_mean		410	
	time/			
	fps		441	
	iterations		339	
	time_elapsed		1573	
	total_timesteps		694272	
	train/			
	approx_kl		0.053204037	
	clip_fraction		0.307	
	clip_range		0.2	
	entropy_loss		-19.3	
	explained_variance		0.921	
	learning_rate		0.0003	
	loss		113	
	n_updates		3380	
	policy_gradient_loss		-0.0592	
	std		0.761	
	value_loss		311	

	rollout/			
	ep_len_mean		83.7	
	ep_rew_mean		415	
	time/			
	fps		440	
	iterations		340	
	time_elapsed		1579	
	total_timesteps		696320	
	train/			
	approx_kl		0.060068496	
	clip_fraction		0.303	
	clip_range		0.2	
	entropy_loss		-19.3	
	explained_variance		0.903	
	learning_rate		0.0003	
	loss		69.3	
	n_updates		3390	
	policy_gradient_loss		-0.056	
	std		0.759	
	value_loss		324	

	rollout/			
	ep_len_mean		82.1	
	ep_rew_mean		407	
	time/			
	fps		441	

iterations	341
time_elapsed	1583
total_timesteps	698368
train/	
approx_kl	0.06380924
clip_fraction	0.338
clip_range	0.2
entropy_loss	-19.3
explained_variance	0.905
learning_rate	0.0003
loss	86.8
n_updates	3400
policy_gradient_loss	-0.0558
std	0.759
value_loss	317

rollout/	
ep_len_mean	83.4
ep_rew_mean	415
time/	
fps	440
iterations	342
time_elapsed	1588
total_timesteps	700416
train/	
approx_kl	0.056965366
clip_fraction	0.323
clip_range	0.2
entropy_loss	-19.3
explained_variance	0.893
learning_rate	0.0003
loss	85.1
n_updates	3410
policy_gradient_loss	-0.0649
std	0.757
value_loss	348

rollout/	
ep_len_mean	83.8
ep_rew_mean	418
time/	
fps	440
iterations	343
time_elapsed	1593
total_timesteps	702464
train/	

approx_kl	0.06808295	
clip_fraction	0.345	
clip_range	0.2	
entropy_loss	-19.3	
explained_variance	0.917	
learning_rate	0.0003	
loss	82.3	
n_updates	3420	
policy_gradient_loss	-0.0572	
std	0.758	
value_loss	297	

rollout/		
ep_len_mean	83.5	
ep_rew_mean	416	
time/		
fps	440	
iterations	344	
time_elapsed	1598	
total_timesteps	704512	
train/		
approx_kl	0.05505716	
clip_fraction	0.3	
clip_range	0.2	
entropy_loss	-19.3	
explained_variance	0.908	
learning_rate	0.0003	
loss	73.8	
n_updates	3430	
policy_gradient_loss	-0.0611	
std	0.758	
value_loss	360	

rollout/		
ep_len_mean	82.6	
ep_rew_mean	412	
time/		
fps	440	
iterations	345	
time_elapsed	1603	
total_timesteps	706560	
train/		
approx_kl	0.05866832	
clip_fraction	0.325	
clip_range	0.2	
entropy_loss	-19.3	



	explained_variance		0.906	
	learning_rate		0.0003	
	loss		153	
	n_updates		3440	
	policy_gradient_loss		-0.0553	
	std		0.758	
	value_loss		330	

	rollout/			
	ep_len_mean		81.5	
	ep_rew_mean		405	
	time/			
	fps		440	
	iterations		346	
	time_elapsed		1608	
	total_timesteps		708608	
	train/			
	approx_kl		0.048537582	
	clip_fraction		0.288	
	clip_range		0.2	
	entropy_loss		-19.3	
	explained_variance		0.91	
	learning_rate		0.0003	
	loss		130	
	n_updates		3450	
	policy_gradient_loss		-0.0569	
	std		0.758	
	value_loss		347	

	rollout/			
	ep_len_mean		81.3	
	ep_rew_mean		401	
	time/			
	fps		440	
	iterations		347	
	time_elapsed		1613	
	total_timesteps		710656	
	train/			
	approx_kl		0.061983563	
	clip_fraction		0.337	
	clip_range		0.2	
	entropy_loss		-19.3	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		95.4	
	n_updates		3460	

	policy_gradient_loss		-0.0589	
	std		0.758	
	value_loss		294	

---

	rollout/			
	ep_len_mean		83.1	
	ep_rew_mean		408	
	time/			
	fps		440	
	iterations		348	
	time_elapsed		1618	
	total_timesteps		712704	
	train/			
	approx_kl		0.054480635	
	clip_fraction		0.309	
	clip_range		0.2	
	entropy_loss		-19.3	
	explained_variance		0.908	
	learning_rate		0.0003	
	loss		150	
	n_updates		3470	
	policy_gradient_loss		-0.0581	
	std		0.757	
	value_loss		345	

---



---

	rollout/			
	ep_len_mean		84.1	
	ep_rew_mean		412	
	time/			
	fps		440	
	iterations		349	
	time_elapsed		1623	
	total_timesteps		714752	
	train/			
	approx_kl		0.054882854	
	clip_fraction		0.327	
	clip_range		0.2	
	entropy_loss		-19.2	
	explained_variance		0.898	
	learning_rate		0.0003	
	loss		119	
	n_updates		3480	
	policy_gradient_loss		-0.0616	
	std		0.757	
	value_loss		367	

---

-----		
rollout/		
ep_len_mean	85.1	
ep_rew_mean	419	
time/		
fps	440	
iterations	350	
time_elapsed	1628	
total_timesteps	716800	
train/		
approx_kl	0.050189726	
clip_fraction	0.314	
clip_range	0.2	
entropy_loss	-19.2	
explained_variance	0.918	
learning_rate	0.0003	
loss	126	
n_updates	3490	
policy_gradient_loss	-0.0638	
std	0.757	
value_loss	350	
-----		

-----		
rollout/		
ep_len_mean	86.2	
ep_rew_mean	426	
time/		
fps	440	
iterations	351	
time_elapsed	1633	
total_timesteps	718848	
train/		
approx_kl	0.043683745	
clip_fraction	0.282	
clip_range	0.2	
entropy_loss	-19.2	
explained_variance	0.892	
learning_rate	0.0003	
loss	204	
n_updates	3500	
policy_gradient_loss	-0.0578	
std	0.756	
value_loss	444	
-----		

-----		
rollout/		
ep_len_mean	83.2	
ep_rew_mean	413	

time/		
fps	440	
iterations	352	
time_elapsed	1637	
total_timesteps	720896	
train/		
approx_kl	0.05403021	
clip_fraction	0.288	
clip_range	0.2	
entropy_loss	-19.2	
explained_variance	0.906	
learning_rate	0.0003	
loss	113	
n_updates	3510	
policy_gradient_loss	-0.0552	
std	0.755	
value_loss	361	

rollout/		
ep_len_mean	82.9	
ep_rew_mean	414	
time/		
fps	440	
iterations	353	
time_elapsed	1642	
total_timesteps	722944	
train/		
approx_kl	0.057725504	
clip_fraction	0.323	
clip_range	0.2	
entropy_loss	-19.2	
explained_variance	0.898	
learning_rate	0.0003	
loss	152	
n_updates	3520	
policy_gradient_loss	-0.0593	
std	0.754	
value_loss	357	

rollout/		
ep_len_mean	84.3	
ep_rew_mean	418	
time/		
fps	440	
iterations	354	
time_elapsed	1647	

	total_timesteps		724992	
	train/			
	approx_kl		0.060122646	
	clip_fraction		0.314	
	clip_range		0.2	
	entropy_loss		-19.2	
	explained_variance		0.899	
	learning_rate		0.0003	
	loss		75.8	
	n_updates		3530	
	policy_gradient_loss		-0.0534	
	std		0.755	
	value_loss		339	

---

	rollout/			
	ep_len_mean		84.7	
	ep_rew_mean		418	
	time/			
	fps		440	
	iterations		355	
	time_elapsed		1652	
	total_timesteps		727040	
	train/			
	approx_kl		0.0505016	
	clip_fraction		0.293	
	clip_range		0.2	
	entropy_loss		-19.2	
	explained_variance		0.892	
	learning_rate		0.0003	
	loss		181	
	n_updates		3540	
	policy_gradient_loss		-0.0568	
	std		0.755	
	value_loss		436	

---

	rollout/			
	ep_len_mean		83.7	
	ep_rew_mean		413	
	time/			
	fps		440	
	iterations		356	
	time_elapsed		1656	
	total_timesteps		729088	
	train/			
	approx_kl		0.057376258	
	clip_fraction		0.298	

clip_range	0.2	
entropy_loss	-19.2	
explained_variance	0.905	
learning_rate	0.0003	
loss	119	
n_updates	3550	
policy_gradient_loss	-0.0552	
std	0.753	
value_loss	321	

rollout/		
ep_len_mean	83.2	
ep_rew_mean	407	
time/		
fps	440	
iterations	357	
time_elapsed	1661	
total_timesteps	731136	
train/		
approx_kl	0.060498945	
clip_fraction	0.301	
clip_range	0.2	
entropy_loss	-19.1	
explained_variance	0.902	
learning_rate	0.0003	
loss	129	
n_updates	3560	
policy_gradient_loss	-0.0574	
std	0.752	
value_loss	375	

rollout/		
ep_len_mean	80.5	
ep_rew_mean	395	
time/		
fps	440	
iterations	358	
time_elapsed	1666	
total_timesteps	733184	
train/		
approx_kl	0.0615414	
clip_fraction	0.346	
clip_range	0.2	
entropy_loss	-19.1	
explained_variance	0.901	
learning_rate	0.0003	

	loss		130	
	n_updates		3570	
	policy_gradient_loss		-0.0628	
	std		0.751	
	value_loss		361	
-----				
	rollout/			
	ep_len_mean		81.3	
	ep_rew_mean		400	
	time/			
	fps		440	
	iterations		359	
	time_elapsed		1670	
	total_timesteps		735232	
	train/			
	approx_kl		0.057366874	
	clip_fraction		0.328	
	clip_range		0.2	
	entropy_loss		-19.1	
	explained_variance		0.915	
	learning_rate		0.0003	
	loss		111	
	n_updates		3580	
	policy_gradient_loss		-0.062	
	std		0.751	
	value_loss		309	
-----				
	rollout/			
	ep_len_mean		82	
	ep_rew_mean		403	
	time/			
	fps		440	
	iterations		360	
	time_elapsed		1675	
	total_timesteps		737280	
	train/			
	approx_kl		0.060164545	
	clip_fraction		0.314	
	clip_range		0.2	
	entropy_loss		-19.1	
	explained_variance		0.926	
	learning_rate		0.0003	
	loss		76.6	
	n_updates		3590	
	policy_gradient_loss		-0.0568	
	std		0.75	

	value_loss		304	
-----				
	rollout/			
	ep_len_mean		81.4	
	ep_rew_mean		401	
	time/			
	fps		440	
	iterations		361	
	time_elapsed		1680	
	total_timesteps		739328	
	train/			
	approx_kl		0.05647254	
	clip_fraction		0.333	
	clip_range		0.2	
	entropy_loss		-19.1	
	explained_variance		0.928	
	learning_rate		0.0003	
	loss		104	
	n_updates		3600	
	policy_gradient_loss		-0.0573	
	std		0.75	
	value_loss		292	
-----				
	rollout/			
	ep_len_mean		81.5	
	ep_rew_mean		403	
	time/			
	fps		439	
	iterations		362	
	time_elapsed		1685	
	total_timesteps		741376	
	train/			
	approx_kl		0.05830235	
	clip_fraction		0.29	
	clip_range		0.2	
	entropy_loss		-19.1	
	explained_variance		0.912	
	learning_rate		0.0003	
	loss		81	
	n_updates		3610	
	policy_gradient_loss		-0.0593	
	std		0.749	
	value_loss		309	
-----				
	rollout/			



	ep_len_mean		81.9	
	ep_rew_mean		406	
	time/			
	fps		439	
	iterations		363	
	time_elapsed		1689	
	total_timesteps		743424	
	train/			
	approx_kl		0.056153473	
	clip_fraction		0.31	
	clip_range		0.2	
	entropy_loss		-19	
	explained_variance		0.922	
	learning_rate		0.0003	
	loss		137	
	n_updates		3620	
	policy_gradient_loss		-0.056	
	std		0.748	
	value_loss		315	

	rollout/			
	ep_len_mean		81.4	
	ep_rew_mean		404	
	time/			
	fps		439	
	iterations		364	
	time_elapsed		1694	
	total_timesteps		745472	
	train/			
	approx_kl		0.05469892	
	clip_fraction		0.311	
	clip_range		0.2	
	entropy_loss		-19	
	explained_variance		0.908	
	learning_rate		0.0003	
	loss		116	
	n_updates		3630	
	policy_gradient_loss		-0.0591	
	std		0.748	
	value_loss		347	

	rollout/			
	ep_len_mean		80.3	
	ep_rew_mean		399	
	time/			
	fps		439	

iterations	365
time_elapsed	1699
total_timesteps	747520
train/	
approx_kl	0.06090732
clip_fraction	0.313
clip_range	0.2
entropy_loss	-19
explained_variance	0.908
learning_rate	0.0003
loss	103
n_updates	3640
policy_gradient_loss	-0.0591
std	0.748
value_loss	336

rollout/	
ep_len_mean	82.9
ep_rew_mean	411
time/	
fps	439
iterations	366
time_elapsed	1703
total_timesteps	749568
train/	
approx_kl	0.060587663
clip_fraction	0.309
clip_range	0.2
entropy_loss	-19
explained_variance	0.91
learning_rate	0.0003
loss	68.9
n_updates	3650
policy_gradient_loss	-0.0577
std	0.748
value_loss	302

rollout/	
ep_len_mean	81.5
ep_rew_mean	405
time/	
fps	439
iterations	367
time_elapsed	1708
total_timesteps	751616
train/	

approx_kl	0.056497168	
clip_fraction	0.32	
clip_range	0.2	
entropy_loss	-19	
explained_variance	0.899	
learning_rate	0.0003	
loss	146	
n_updates	3660	
policy_gradient_loss	-0.056	
std	0.747	
value_loss	347	

rollout/		
ep_len_mean	82.1	
ep_rew_mean	409	
time/		
fps	439	
iterations	368	
time_elapsed	1713	
total_timesteps	753664	
train/		
approx_kl	0.058048304	
clip_fraction	0.347	
clip_range	0.2	
entropy_loss	-19	
explained_variance	0.915	
learning_rate	0.0003	
loss	113	
n_updates	3670	
policy_gradient_loss	-0.0555	
std	0.747	
value_loss	292	

rollout/		
ep_len_mean	84.5	
ep_rew_mean	420	
time/		
fps	439	
iterations	369	
time_elapsed	1718	
total_timesteps	755712	
train/		
approx_kl	0.06559527	
clip_fraction	0.328	
clip_range	0.2	
entropy_loss	-19	

	explained_variance		0.914	
	learning_rate		0.0003	
	loss		86.6	
	n_updates		3680	
	policy_gradient_loss		-0.0581	
	std		0.747	
	value_loss		317	

	rollout/			
	ep_len_mean		82.6	
	ep_rew_mean		412	
	time/			
	fps		439	
	iterations		370	
	time_elapsed		1723	
	total_timesteps		757760	
	train/			
	approx_kl		0.058114734	
	clip_fraction		0.329	
	clip_range		0.2	
	entropy_loss		-19	
	explained_variance		0.906	
	learning_rate		0.0003	
	loss		68.3	
	n_updates		3690	
	policy_gradient_loss		-0.0581	
	std		0.747	
	value_loss		308	

	rollout/			
	ep_len_mean		82.8	
	ep_rew_mean		411	
	time/			
	fps		439	
	iterations		371	
	time_elapsed		1728	
	total_timesteps		759808	
	train/			
	approx_kl		0.059836984	
	clip_fraction		0.341	
	clip_range		0.2	
	entropy_loss		-19	
	explained_variance		0.919	
	learning_rate		0.0003	
	loss		99.5	
	n_updates		3700	

	policy_gradient_loss		-0.0585	
	std		0.746	
	value_loss		283	

---

	rollout/			
	ep_len_mean		82.5	
	ep_rew_mean		409	
	time/			
	fps		439	
	iterations		372	
	time_elapsed		1733	
	total_timesteps		761856	
	train/			
	approx_kl		0.05289513	
	clip_fraction		0.319	
	clip_range		0.2	
	entropy_loss		-19	
	explained_variance		0.914	
	learning_rate		0.0003	
	loss		126	
	n_updates		3710	
	policy_gradient_loss		-0.0589	
	std		0.745	
	value_loss		309	

---



---

	rollout/			
	ep_len_mean		80.7	
	ep_rew_mean		400	
	time/			
	fps		439	
	iterations		373	
	time_elapsed		1737	
	total_timesteps		763904	
	train/			
	approx_kl		0.059107255	
	clip_fraction		0.312	
	clip_range		0.2	
	entropy_loss		-19	
	explained_variance		0.897	
	learning_rate		0.0003	
	loss		133	
	n_updates		3720	
	policy_gradient_loss		-0.0592	
	std		0.745	
	value_loss		342	

---

---

rollout/		
ep_len_mean	81.8	
ep_rew_mean	404	
time/		
fps	439	
iterations	374	
time_elapsed	1742	
total_timesteps	765952	
train/		
approx_kl	0.06853158	
clip_fraction	0.333	
clip_range	0.2	
entropy_loss	-18.9	
explained_variance	0.908	
learning_rate	0.0003	
loss	107	
n_updates	3730	
policy_gradient_loss	-0.0572	
std	0.744	
value_loss	324	

---



---

rollout/		
ep_len_mean	80.5	
ep_rew_mean	399	
time/		
fps	439	
iterations	375	
time_elapsed	1747	
total_timesteps	768000	
train/		
approx_kl	0.0540872	
clip_fraction	0.335	
clip_range	0.2	
entropy_loss	-18.9	
explained_variance	0.911	
learning_rate	0.0003	
loss	135	
n_updates	3740	
policy_gradient_loss	-0.0579	
std	0.743	
value_loss	332	

---



---

rollout/		
ep_len_mean	80.5	
ep_rew_mean	398	

---

time/		
fps	439	
iterations	376	
time_elapsed	1751	
total_timesteps	770048	
train/		
approx_kl	0.05835658	
clip_fraction	0.324	
clip_range	0.2	
entropy_loss	-18.9	
explained_variance	0.922	
learning_rate	0.0003	
loss	150	
n_updates	3750	
policy_gradient_loss	-0.059	
std	0.743	
value_loss	303	

---

rollout/		
ep_len_mean	82.4	
ep_rew_mean	407	
time/		
fps	439	
iterations	377	
time_elapsed	1756	
total_timesteps	772096	
train/		
approx_kl	0.0486343	
clip_fraction	0.302	
clip_range	0.2	
entropy_loss	-18.9	
explained_variance	0.899	
learning_rate	0.0003	
loss	157	
n_updates	3760	
policy_gradient_loss	-0.0573	
std	0.742	
value_loss	401	

---

rollout/		
ep_len_mean	81.9	
ep_rew_mean	404	
time/		
fps	439	
iterations	378	
time_elapsed	1761	

	total_timesteps		774144	
	train/			
	approx_kl		0.057854936	
	clip_fraction		0.291	
	clip_range		0.2	
	entropy_loss		-18.9	
	explained_variance		0.883	
	learning_rate		0.0003	
	loss		118	
	n_updates		3770	
	policy_gradient_loss		-0.0534	
	std		0.742	
	value_loss		394	

	rollout/			
	ep_len_mean		82.9	
	ep_rew_mean		408	
	time/			
	fps		439	
	iterations		379	
	time_elapsed		1765	
	total_timesteps		776192	
	train/			
	approx_kl		0.063734934	
	clip_fraction		0.341	
	clip_range		0.2	
	entropy_loss		-18.9	
	explained_variance		0.908	
	learning_rate		0.0003	
	loss		111	
	n_updates		3780	
	policy_gradient_loss		-0.0617	
	std		0.742	
	value_loss		340	

	rollout/			
	ep_len_mean		84.4	
	ep_rew_mean		414	
	time/			
	fps		439	
	iterations		380	
	time_elapsed		1770	
	total_timesteps		778240	
	train/			
	approx_kl		0.062593475	
	clip_fraction		0.327	



clip_range	0.2	
entropy_loss	-18.9	
explained_variance	0.913	
learning_rate	0.0003	
loss	128	
n_updates	3790	
policy_gradient_loss	-0.0606	
std	0.741	
value_loss	315	

rollout/		
ep_len_mean	83.3	
ep_rew_mean	410	
time/		
fps	439	
iterations	381	
time_elapsed	1775	
total_timesteps	780288	
train/		
approx_kl	0.051825665	
clip_fraction	0.332	
clip_range	0.2	
entropy_loss	-18.9	
explained_variance	0.916	
learning_rate	0.0003	
loss	152	
n_updates	3800	
policy_gradient_loss	-0.0604	
std	0.74	
value_loss	344	

rollout/		
ep_len_mean	84.1	
ep_rew_mean	417	
time/		
fps	439	
iterations	382	
time_elapsed	1780	
total_timesteps	782336	
train/		
approx_kl	0.052106168	
clip_fraction	0.306	
clip_range	0.2	
entropy_loss	-18.8	
explained_variance	0.916	
learning_rate	0.0003	

	loss		134	
	n_updates		3810	
	policy_gradient_loss		-0.0593	
	std		0.74	
	value_loss		319	

---

	rollout/			
	ep_len_mean		83.5	
	ep_rew_mean		415	
	time/			
	fps		439	
	iterations		383	
	time_elapsed		1785	
	total_timesteps		784384	
	train/			
	approx_kl		0.07114923	
	clip_fraction		0.312	
	clip_range		0.2	
	entropy_loss		-18.8	
	explained_variance		0.907	
	learning_rate		0.0003	
	loss		130	
	n_updates		3820	
	policy_gradient_loss		-0.0571	
	std		0.738	
	value_loss		338	

---

	rollout/			
	ep_len_mean		83.1	
	ep_rew_mean		415	
	time/			
	fps		439	
	iterations		384	
	time_elapsed		1790	
	total_timesteps		786432	
	train/			
	approx_kl		0.061876472	
	clip_fraction		0.317	
	clip_range		0.2	
	entropy_loss		-18.8	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		104	
	n_updates		3830	
	policy_gradient_loss		-0.0564	
	std		0.738	

	value_loss	314	
-----			
	rollout/		
	ep_len_mean	83.5	
	ep_rew_mean	417	
	time/		
	fps	439	
	iterations	385	
	time_elapsed	1795	
	total_timesteps	788480	
	train/		
	approx_kl	0.050785102	
	clip_fraction	0.29	
	clip_range	0.2	
	entropy_loss	-18.8	
	explained_variance	0.913	
	learning_rate	0.0003	
	loss	135	
	n_updates	3840	
	policy_gradient_loss	-0.0558	
	std	0.737	
	value_loss	365	
-----			
	rollout/		
	ep_len_mean	83	
	ep_rew_mean	414	
	time/		
	fps	439	
	iterations	386	
	time_elapsed	1799	
	total_timesteps	790528	
	train/		
	approx_kl	0.06484822	
	clip_fraction	0.303	
	clip_range	0.2	
	entropy_loss	-18.8	
	explained_variance	0.898	
	learning_rate	0.0003	
	loss	115	
	n_updates	3850	
	policy_gradient_loss	-0.0545	
	std	0.738	
	value_loss	372	
-----			
	rollout/		

	ep_len_mean	83.4	
	ep_rew_mean	416	
	time/		
	fps	439	
	iterations	387	
	time_elapsed	1804	
	total_timesteps	792576	
	train/		
	approx_kl	0.05506862	
	clip_fraction	0.321	
	clip_range	0.2	
	entropy_loss	-18.8	
	explained_variance	0.917	
	learning_rate	0.0003	
	loss	124	
	n_updates	3860	
	policy_gradient_loss	-0.0547	
	std	0.737	
	value_loss	325	

	rollout/		
	ep_len_mean	82.7	
	ep_rew_mean	414	
	time/		
	fps	439	
	iterations	388	
	time_elapsed	1809	
	total_timesteps	794624	
	train/		
	approx_kl	0.060606644	
	clip_fraction	0.313	
	clip_range	0.2	
	entropy_loss	-18.8	
	explained_variance	0.905	
	learning_rate	0.0003	
	loss	134	
	n_updates	3870	
	policy_gradient_loss	-0.0566	
	std	0.737	
	value_loss	375	

	rollout/		
	ep_len_mean	82.5	
	ep_rew_mean	414	
	time/		
	fps	439	

iterations	389
time_elapsed	1814
total_timesteps	796672
train/	
approx_kl	0.0661463
clip_fraction	0.31
clip_range	0.2
entropy_loss	-18.8
explained_variance	0.916
learning_rate	0.0003
loss	116
n_updates	3880
policy_gradient_loss	-0.0562
std	0.738
value_loss	360

rollout/	
ep_len_mean	82.1
ep_rew_mean	411
time/	
fps	438
iterations	390
time_elapsed	1819
total_timesteps	798720
train/	
approx_kl	0.06492148
clip_fraction	0.313
clip_range	0.2
entropy_loss	-18.8
explained_variance	0.906
learning_rate	0.0003
loss	97.3
n_updates	3890
policy_gradient_loss	-0.0578
std	0.738
value_loss	316

rollout/	
ep_len_mean	84
ep_rew_mean	417
time/	
fps	438
iterations	391
time_elapsed	1824
total_timesteps	800768
train/	

approx_kl	0.062441565	
clip_fraction	0.327	
clip_range	0.2	
entropy_loss	-18.8	
explained_variance	0.91	
learning_rate	0.0003	
loss	135	
n_updates	3900	
policy_gradient_loss	-0.0604	
std	0.737	
value_loss	343	

rollout/		
ep_len_mean	84.1	
ep_rew_mean	417	
time/		
fps	438	
iterations	392	
time_elapsed	1829	
total_timesteps	802816	
train/		
approx_kl	0.06557526	
clip_fraction	0.311	
clip_range	0.2	
entropy_loss	-18.8	
explained_variance	0.906	
learning_rate	0.0003	
loss	115	
n_updates	3910	
policy_gradient_loss	-0.0565	
std	0.737	
value_loss	332	

rollout/		
ep_len_mean	84.3	
ep_rew_mean	418	
time/		
fps	438	
iterations	393	
time_elapsed	1833	
total_timesteps	804864	
train/		
approx_kl	0.056534767	
clip_fraction	0.306	
clip_range	0.2	
entropy_loss	-18.7	

	explained_variance		0.914	
	learning_rate		0.0003	
	loss		115	
	n_updates		3920	
	policy_gradient_loss		-0.0545	
	std		0.736	
	value_loss		327	

	rollout/			
	ep_len_mean		84.2	
	ep_rew_mean		417	
	time/			
	fps		438	
	iterations		394	
	time_elapsed		1838	
	total_timesteps		806912	
	train/			
	approx_kl		0.04824531	
	clip_fraction		0.293	
	clip_range		0.2	
	entropy_loss		-18.7	
	explained_variance		0.899	
	learning_rate		0.0003	
	loss		120	
	n_updates		3930	
	policy_gradient_loss		-0.0531	
	std		0.736	
	value_loss		388	

	rollout/			
	ep_len_mean		84.3	
	ep_rew_mean		417	
	time/			
	fps		438	
	iterations		395	
	time_elapsed		1843	
	total_timesteps		808960	
	train/			
	approx_kl		0.07535021	
	clip_fraction		0.309	
	clip_range		0.2	
	entropy_loss		-18.7	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		108	
	n_updates		3940	

	policy_gradient_loss		-0.0543	
	std		0.736	
	value_loss		351	

	rollout/			
	ep_len_mean		83.4	
	ep_rew_mean		412	
	time/			
	fps		438	
	iterations		396	
	time_elapsed		1848	
	total_timesteps		811008	
	train/			
	approx_kl		0.052610703	
	clip_fraction		0.303	
	clip_range		0.2	
	entropy_loss		-18.7	
	explained_variance		0.885	
	learning_rate		0.0003	
	loss		98.7	
	n_updates		3950	
	policy_gradient_loss		-0.0568	
	std		0.735	
	value_loss		426	

	rollout/			
	ep_len_mean		84.2	
	ep_rew_mean		416	
	time/			
	fps		438	
	iterations		397	
	time_elapsed		1853	
	total_timesteps		813056	
	train/			
	approx_kl		0.056993686	
	clip_fraction		0.307	
	clip_range		0.2	
	entropy_loss		-18.7	
	explained_variance		0.893	
	learning_rate		0.0003	
	loss		248	
	n_updates		3960	
	policy_gradient_loss		-0.061	
	std		0.735	
	value_loss		430	



---

rollout/		
ep_len_mean	86.2	
ep_rew_mean	426	
time/		
fps	438	
iterations	398	
time_elapsed	1858	
total_timesteps	815104	
train/		
approx_kl	0.07514132	
clip_fraction	0.343	
clip_range	0.2	
entropy_loss	-18.7	
explained_variance	0.92	
learning_rate	0.0003	
loss	88.8	
n_updates	3970	
policy_gradient_loss	-0.0594	
std	0.734	
value_loss	304	

---



---

rollout/		
ep_len_mean	85.8	
ep_rew_mean	425	
time/		
fps	438	
iterations	399	
time_elapsed	1864	
total_timesteps	817152	
train/		
approx_kl	0.07200529	
clip_fraction	0.34	
clip_range	0.2	
entropy_loss	-18.7	
explained_variance	0.919	
learning_rate	0.0003	
loss	90.8	
n_updates	3980	
policy_gradient_loss	-0.0615	
std	0.735	
value_loss	321	

---



---

rollout/		
ep_len_mean	85.5	
ep_rew_mean	424	

---

time/		
fps	438	
iterations	400	
time_elapsed	1869	
total_timesteps	819200	
train/		
approx_kl	0.06531005	
clip_fraction	0.323	
clip_range	0.2	
entropy_loss	-18.7	
explained_variance	0.924	
learning_rate	0.0003	
loss	118	
n_updates	3990	
policy_gradient_loss	-0.0562	
std	0.734	
value_loss	321	

rollout/		
ep_len_mean	85.8	
ep_rew_mean	425	
time/		
fps	438	
iterations	401	
time_elapsed	1874	
total_timesteps	821248	
train/		
approx_kl	0.067447826	
clip_fraction	0.331	
clip_range	0.2	
entropy_loss	-18.7	
explained_variance	0.914	
learning_rate	0.0003	
loss	97.1	
n_updates	4000	
policy_gradient_loss	-0.0586	
std	0.733	
value_loss	320	

rollout/		
ep_len_mean	84.3	
ep_rew_mean	418	
time/		
fps	438	
iterations	402	
time_elapsed	1879	

	total_timesteps		823296	
	train/			
	approx_kl		0.072138056	
	clip_fraction		0.34	
	clip_range		0.2	
	entropy_loss		-18.7	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		128	
	n_updates		4010	
	policy_gradient_loss		-0.0574	
	std		0.732	
	value_loss		279	

---

	rollout/			
	ep_len_mean		83.5	
	ep_rew_mean		415	
	time/			
	fps		437	
	iterations		403	
	time_elapsed		1884	
	total_timesteps		825344	
	train/			
	approx_kl		0.077481925	
	clip_fraction		0.343	
	clip_range		0.2	
	entropy_loss		-18.6	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		91	
	n_updates		4020	
	policy_gradient_loss		-0.0559	
	std		0.732	
	value_loss		328	

---

	rollout/			
	ep_len_mean		85.4	
	ep_rew_mean		426	
	time/			
	fps		437	
	iterations		404	
	time_elapsed		1889	
	total_timesteps		827392	
	train/			
	approx_kl		0.07447772	
	clip_fraction		0.366	

clip_range	0.2	
entropy_loss	-18.6	
explained_variance	0.936	
learning_rate	0.0003	
loss	94.4	
n_updates	4030	
policy_gradient_loss	-0.0613	
std	0.732	
value_loss	233	

rollout/		
ep_len_mean	86	
ep_rew_mean	427	
time/		
fps	437	
iterations	405	
time_elapsed	1893	
total_timesteps	829440	
train/		
approx_kl	0.088220954	
clip_fraction	0.382	
clip_range	0.2	
entropy_loss	-18.6	
explained_variance	0.919	
learning_rate	0.0003	
loss	122	
n_updates	4040	
policy_gradient_loss	-0.0661	
std	0.731	
value_loss	316	

rollout/		
ep_len_mean	85.2	
ep_rew_mean	423	
time/		
fps	437	
iterations	406	
time_elapsed	1898	
total_timesteps	831488	
train/		
approx_kl	0.06623647	
clip_fraction	0.339	
clip_range	0.2	
entropy_loss	-18.6	
explained_variance	0.927	
learning_rate	0.0003	

	loss		127	
	n_updates		4050	
	policy_gradient_loss		-0.0586	
	std		0.73	
	value_loss		309	

	rollout/			
	ep_len_mean		86.9	
	ep_rew_mean		431	
	time/			
	fps		437	
	iterations		407	
	time_elapsed		1903	
	total_timesteps		833536	
	train/			
	approx_kl		0.07471658	
	clip_fraction		0.337	
	clip_range		0.2	
	entropy_loss		-18.6	
	explained_variance		0.919	
	learning_rate		0.0003	
	loss		213	
	n_updates		4060	
	policy_gradient_loss		-0.0534	
	std		0.728	
	value_loss		361	

	rollout/			
	ep_len_mean		83.2	
	ep_rew_mean		410	
	time/			
	fps		437	
	iterations		408	
	time_elapsed		1907	
	total_timesteps		835584	
	train/			
	approx_kl		0.059574764	
	clip_fraction		0.328	
	clip_range		0.2	
	entropy_loss		-18.5	
	explained_variance		0.921	
	learning_rate		0.0003	
	loss		147	
	n_updates		4070	
	policy_gradient_loss		-0.0574	
	std		0.728	

	value_loss	330	
-----			
	rollout/		
	ep_len_mean	81	
	ep_rew_mean	400	
	time/		
	fps	437	
	iterations	409	
	time_elapsed	1912	
	total_timesteps	837632	
	train/		
	approx_kl	0.070989944	
	clip_fraction	0.337	
	clip_range	0.2	
	entropy_loss	-18.5	
	explained_variance	0.927	
	learning_rate	0.0003	
	loss	119	
	n_updates	4080	
	policy_gradient_loss	-0.0572	
	std	0.728	
	value_loss	274	
-----			
	rollout/		
	ep_len_mean	81.6	
	ep_rew_mean	403	
	time/		
	fps	438	
	iterations	410	
	time_elapsed	1916	
	total_timesteps	839680	
	train/		
	approx_kl	0.06835993	
	clip_fraction	0.305	
	clip_range	0.2	
	entropy_loss	-18.5	
	explained_variance	0.905	
	learning_rate	0.0003	
	loss	96.5	
	n_updates	4090	
	policy_gradient_loss	-0.0584	
	std	0.727	
	value_loss	342	
-----			
	rollout/		

	ep_len_mean		81.1	
	ep_rew_mean		399	
	time/			
	fps		438	
	iterations		411	
	time_elapsed		1921	
	total_timesteps		841728	
	train/			
	approx_kl		0.06296618	
	clip_fraction		0.329	
	clip_range		0.2	
	entropy_loss		-18.5	
	explained_variance		0.916	
	learning_rate		0.0003	
	loss		121	
	n_updates		4100	
	policy_gradient_loss		-0.0609	
	std		0.727	
	value_loss		299	

	rollout/			
	ep_len_mean		82.4	
	ep_rew_mean		409	
	time/			
	fps		438	
	iterations		412	
	time_elapsed		1926	
	total_timesteps		843776	
	train/			
	approx_kl		0.06619055	
	clip_fraction		0.357	
	clip_range		0.2	
	entropy_loss		-18.5	
	explained_variance		0.923	
	learning_rate		0.0003	
	loss		91	
	n_updates		4110	
	policy_gradient_loss		-0.0609	
	std		0.725	
	value_loss		279	

	rollout/			
	ep_len_mean		85.3	
	ep_rew_mean		421	
	time/			
	fps		438	

iterations	413
time_elapsed	1930
total_timesteps	845824
train/	
approx_kl	0.06989674
clip_fraction	0.351
clip_range	0.2
entropy_loss	-18.5
explained_variance	0.91
learning_rate	0.0003
loss	115
n_updates	4120
policy_gradient_loss	-0.0567
std	0.724
value_loss	288

rollout/	
ep_len_mean	84.9
ep_rew_mean	419
time/	
fps	438
iterations	414
time_elapsed	1935
total_timesteps	847872
train/	
approx_kl	0.08262791
clip_fraction	0.386
clip_range	0.2
entropy_loss	-18.4
explained_variance	0.921
learning_rate	0.0003
loss	104
n_updates	4130
policy_gradient_loss	-0.0541
std	0.723
value_loss	260

rollout/	
ep_len_mean	84.9
ep_rew_mean	420
time/	
fps	438
iterations	415
time_elapsed	1940
total_timesteps	849920
train/	



approx_kl	0.058272682	
clip_fraction	0.321	
clip_range	0.2	
entropy_loss	-18.4	
explained_variance	0.912	
learning_rate	0.0003	
loss	152	
n_updates	4140	
policy_gradient_loss	-0.0548	
std	0.723	
value_loss	387	

rollout/		
ep_len_mean	86.3	
ep_rew_mean	424	
time/		
fps	438	
iterations	416	
time_elapsed	1944	
total_timesteps	851968	
train/		
approx_kl	0.06249902	
clip_fraction	0.33	
clip_range	0.2	
entropy_loss	-18.4	
explained_variance	0.915	
learning_rate	0.0003	
loss	132	
n_updates	4150	
policy_gradient_loss	-0.057	
std	0.723	
value_loss	303	

rollout/		
ep_len_mean	84.4	
ep_rew_mean	414	
time/		
fps	438	
iterations	417	
time_elapsed	1949	
total_timesteps	854016	
train/		
approx_kl	0.0636737	
clip_fraction	0.319	
clip_range	0.2	
entropy_loss	-18.4	

	explained_variance		0.905	
	learning_rate		0.0003	
	loss		175	
	n_updates		4160	
	policy_gradient_loss		-0.0567	
	std		0.723	
	value_loss		337	

	rollout/			
	ep_len_mean		85.9	
	ep_rew_mean		424	
	time/			
	fps		438	
	iterations		418	
	time_elapsed		1954	
	total_timesteps		856064	
	train/			
	approx_kl		0.06038027	
	clip_fraction		0.343	
	clip_range		0.2	
	entropy_loss		-18.4	
	explained_variance		0.906	
	learning_rate		0.0003	
	loss		86.3	
	n_updates		4170	
	policy_gradient_loss		-0.0598	
	std		0.723	
	value_loss		323	

	rollout/			
	ep_len_mean		86	
	ep_rew_mean		426	
	time/			
	fps		438	
	iterations		419	
	time_elapsed		1959	
	total_timesteps		858112	
	train/			
	approx_kl		0.082924336	
	clip_fraction		0.366	
	clip_range		0.2	
	entropy_loss		-18.4	
	explained_variance		0.927	
	learning_rate		0.0003	
	loss		93.7	
	n_updates		4180	

	policy_gradient_loss		-0.0562	
	std		0.723	
	value_loss		258	

---

	rollout/			
	ep_len_mean		88.5	
	ep_rew_mean		440	
	time/			
	fps		438	
	iterations		420	
	time_elapsed		1963	
	total_timesteps		860160	
	train/			
	approx_kl		0.07298912	
	clip_fraction		0.359	
	clip_range		0.2	
	entropy_loss		-18.4	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		107	
	n_updates		4190	
	policy_gradient_loss		-0.0594	
	std		0.723	
	value_loss		272	

---

	rollout/			
	ep_len_mean		88.1	
	ep_rew_mean		438	
	time/			
	fps		438	
	iterations		421	
	time_elapsed		1967	
	total_timesteps		862208	
	train/			
	approx_kl		0.06400592	
	clip_fraction		0.343	
	clip_range		0.2	
	entropy_loss		-18.4	
	explained_variance		0.923	
	learning_rate		0.0003	
	loss		84.3	
	n_updates		4200	
	policy_gradient_loss		-0.06	
	std		0.722	
	value_loss		302	

---

---

rollout/		
ep_len_mean	88	
ep_rew_mean	437	
time/		
fps	438	
iterations	422	
time_elapsed	1972	
total_timesteps	864256	
train/		
approx_kl	0.06696643	
clip_fraction	0.317	
clip_range	0.2	
entropy_loss	-18.4	
explained_variance	0.925	
learning_rate	0.0003	
loss	122	
n_updates	4210	
policy_gradient_loss	-0.0563	
std	0.722	
value_loss	331	

---



---

rollout/		
ep_len_mean	88.2	
ep_rew_mean	439	
time/		
fps	438	
iterations	423	
time_elapsed	1976	
total_timesteps	866304	
train/		
approx_kl	0.06322785	
clip_fraction	0.326	
clip_range	0.2	
entropy_loss	-18.4	
explained_variance	0.924	
learning_rate	0.0003	
loss	98.5	
n_updates	4220	
policy_gradient_loss	-0.0568	
std	0.721	
value_loss	303	

---



---

rollout/		
ep_len_mean	85.8	
ep_rew_mean	428	

---

time/		
fps	438	
iterations	424	
time_elapsed	1981	
total_timesteps	868352	
train/		
approx_kl	0.075339586	
clip_fraction	0.358	
clip_range	0.2	
entropy_loss	-18.3	
explained_variance	0.92	
learning_rate	0.0003	
loss	136	
n_updates	4230	
policy_gradient_loss	-0.0578	
std	0.719	
value_loss	314	

---

rollout/		
ep_len_mean	85.9	
ep_rew_mean	429	
time/		
fps	438	
iterations	425	
time_elapsed	1985	
total_timesteps	870400	
train/		
approx_kl	0.07037902	
clip_fraction	0.318	
clip_range	0.2	
entropy_loss	-18.3	
explained_variance	0.913	
learning_rate	0.0003	
loss	111	
n_updates	4240	
policy_gradient_loss	-0.0557	
std	0.718	
value_loss	323	

---

rollout/		
ep_len_mean	83.8	
ep_rew_mean	418	
time/		
fps	438	
iterations	426	
time_elapsed	1990	

	total_timesteps		872448	
	train/			
	approx_kl		0.057187397	
	clip_fraction		0.306	
	clip_range		0.2	
	entropy_loss		-18.3	
	explained_variance		0.905	
	learning_rate		0.0003	
	loss		153	
	n_updates		4250	
	policy_gradient_loss		-0.0588	
	std		0.717	
	value_loss		386	

	rollout/			
	ep_len_mean		82.6	
	ep_rew_mean		412	
	time/			
	fps		438	
	iterations		427	
	time_elapsed		1995	
	total_timesteps		874496	
	train/			
	approx_kl		0.06864112	
	clip_fraction		0.329	
	clip_range		0.2	
	entropy_loss		-18.3	
	explained_variance		0.927	
	learning_rate		0.0003	
	loss		91.8	
	n_updates		4260	
	policy_gradient_loss		-0.0546	
	std		0.716	
	value_loss		286	

	rollout/			
	ep_len_mean		83.7	
	ep_rew_mean		417	
	time/			
	fps		438	
	iterations		428	
	time_elapsed		1999	
	total_timesteps		876544	
	train/			
	approx_kl		0.061831407	
	clip_fraction		0.36	

clip_range	0.2	
entropy_loss	-18.3	
explained_variance	0.925	
learning_rate	0.0003	
loss	96.8	
n_updates	4270	
policy_gradient_loss	-0.0621	
std	0.716	
value_loss	290	

rollout/		
ep_len_mean	81.7	
ep_rew_mean	408	
time/		
fps	438	
iterations	429	
time_elapsed	2004	
total_timesteps	878592	
train/		
approx_kl	0.0769848	
clip_fraction	0.359	
clip_range	0.2	
entropy_loss	-18.2	
explained_variance	0.907	
learning_rate	0.0003	
loss	81.3	
n_updates	4280	
policy_gradient_loss	-0.056	
std	0.715	
value_loss	314	

rollout/		
ep_len_mean	81.9	
ep_rew_mean	409	
time/		
fps	438	
iterations	430	
time_elapsed	2008	
total_timesteps	880640	
train/		
approx_kl	0.07216303	
clip_fraction	0.348	
clip_range	0.2	
entropy_loss	-18.2	
explained_variance	0.927	
learning_rate	0.0003	

	loss		121	
	n_updates		4290	
	policy_gradient_loss		-0.0601	
	std		0.715	
	value_loss		310	

	rollout/			
	ep_len_mean		83.7	
	ep_rew_mean		419	
	time/			
	fps		438	
	iterations		431	
	time_elapsed		2012	
	total_timesteps		882688	
	train/			
	approx_kl		0.07125511	
	clip_fraction		0.321	
	clip_range		0.2	
	entropy_loss		-18.2	
	explained_variance		0.922	
	learning_rate		0.0003	
	loss		115	
	n_updates		4300	
	policy_gradient_loss		-0.0546	
	std		0.715	
	value_loss		349	

	rollout/			
	ep_len_mean		82.6	
	ep_rew_mean		410	
	time/			
	fps		438	
	iterations		432	
	time_elapsed		2017	
	total_timesteps		884736	
	train/			
	approx_kl		0.068975285	
	clip_fraction		0.331	
	clip_range		0.2	
	entropy_loss		-18.2	
	explained_variance		0.925	
	learning_rate		0.0003	
	loss		108	
	n_updates		4310	
	policy_gradient_loss		-0.0532	
	std		0.713	



	value_loss		333	
-----				
	rollout/			
	ep_len_mean		84.2	
	ep_rew_mean		417	
	time/			
	fps		438	
	iterations		433	
	time_elapsed		2021	
	total_timesteps		886784	
	train/			
	approx_kl		0.07092878	
	clip_fraction		0.34	
	clip_range		0.2	
	entropy_loss		-18.2	
	explained_variance		0.892	
	learning_rate		0.0003	
	loss		110	
	n_updates		4320	
	policy_gradient_loss		-0.0562	
	std		0.714	
	value_loss		349	
-----				
	rollout/			
	ep_len_mean		87.1	
	ep_rew_mean		432	
	time/			
	fps		438	
	iterations		434	
	time_elapsed		2026	
	total_timesteps		888832	
	train/			
	approx_kl		0.062456477	
	clip_fraction		0.328	
	clip_range		0.2	
	entropy_loss		-18.2	
	explained_variance		0.898	
	learning_rate		0.0003	
	loss		123	
	n_updates		4330	
	policy_gradient_loss		-0.0599	
	std		0.713	
	value_loss		344	
-----				
	rollout/			

	ep_len_mean		87.3	
	ep_rew_mean		431	
	time/			
	fps		438	
	iterations		435	
	time_elapsed		2030	
	total_timesteps		890880	
	train/			
	approx_kl		0.0803453	
	clip_fraction		0.36	
	clip_range		0.2	
	entropy_loss		-18.2	
	explained_variance		0.908	
	learning_rate		0.0003	
	loss		123	
	n_updates		4340	
	policy_gradient_loss		-0.0607	
	std		0.712	
	value_loss		317	

	rollout/			
	ep_len_mean		86.1	
	ep_rew_mean		425	
	time/			
	fps		438	
	iterations		436	
	time_elapsed		2035	
	total_timesteps		892928	
	train/			
	approx_kl		0.07040608	
	clip_fraction		0.351	
	clip_range		0.2	
	entropy_loss		-18.2	
	explained_variance		0.909	
	learning_rate		0.0003	
	loss		157	
	n_updates		4350	
	policy_gradient_loss		-0.0608	
	std		0.712	
	value_loss		335	

	rollout/			
	ep_len_mean		84.8	
	ep_rew_mean		419	
	time/			
	fps		438	

iterations	437
time_elapsed	2039
total_timesteps	894976
train/	
approx_kl	0.058486186
clip_fraction	0.343
clip_range	0.2
entropy_loss	-18.2
explained_variance	0.916
learning_rate	0.0003
loss	104
n_updates	4360
policy_gradient_loss	-0.0537
std	0.712
value_loss	348

rollout/	
ep_len_mean	82.1
ep_rew_mean	407
time/	
fps	438
iterations	438
time_elapsed	2043
total_timesteps	897024
train/	
approx_kl	0.07487102
clip_fraction	0.343
clip_range	0.2
entropy_loss	-18.2
explained_variance	0.92
learning_rate	0.0003
loss	118
n_updates	4370
policy_gradient_loss	-0.0567
std	0.713
value_loss	319

rollout/	
ep_len_mean	83.3
ep_rew_mean	413
time/	
fps	438
iterations	439
time_elapsed	2048
total_timesteps	899072
train/	

approx_kl	0.069438085	
clip_fraction	0.329	
clip_range	0.2	
entropy_loss	-18.1	
explained_variance	0.92	
learning_rate	0.0003	
loss	73.9	
n_updates	4380	
policy_gradient_loss	-0.0587	
std	0.711	
value_loss	289	

rollout/		
ep_len_mean	83.7	
ep_rew_mean	414	
time/		
fps	438	
iterations	440	
time_elapsed	2052	
total_timesteps	901120	
train/		
approx_kl	0.06448495	
clip_fraction	0.337	
clip_range	0.2	
entropy_loss	-18.1	
explained_variance	0.914	
learning_rate	0.0003	
loss	60.3	
n_updates	4390	
policy_gradient_loss	-0.0536	
std	0.711	
value_loss	307	

rollout/		
ep_len_mean	85.5	
ep_rew_mean	423	
time/		
fps	439	
iterations	441	
time_elapsed	2057	
total_timesteps	903168	
train/		
approx_kl	0.05973091	
clip_fraction	0.337	
clip_range	0.2	
entropy_loss	-18.1	

	explained_variance		0.912	
	learning_rate		0.0003	
	loss		113	
	n_updates		4400	
	policy_gradient_loss		-0.057	
	std		0.71	
	value_loss		347	

	rollout/			
	ep_len_mean		83.6	
	ep_rew_mean		414	
	time/			
	fps		439	
	iterations		442	
	time_elapsed		2061	
	total_timesteps		905216	
	train/			
	approx_kl		0.07675816	
	clip_fraction		0.318	
	clip_range		0.2	
	entropy_loss		-18.1	
	explained_variance		0.911	
	learning_rate		0.0003	
	loss		127	
	n_updates		4410	
	policy_gradient_loss		-0.0517	
	std		0.71	
	value_loss		367	

	rollout/			
	ep_len_mean		82.2	
	ep_rew_mean		409	
	time/			
	fps		439	
	iterations		443	
	time_elapsed		2066	
	total_timesteps		907264	
	train/			
	approx_kl		0.077478714	
	clip_fraction		0.39	
	clip_range		0.2	
	entropy_loss		-18.1	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		103	
	n_updates		4420	

	policy_gradient_loss		-0.0598	
	std		0.709	
	value_loss		287	

---

	rollout/			
	ep_len_mean		81.7	
	ep_rew_mean		409	
	time/			
	fps		439	
	iterations		444	
	time_elapsed		2070	
	total_timesteps		909312	
	train/			
	approx_kl		0.057111237	
	clip_fraction		0.329	
	clip_range		0.2	
	entropy_loss		-18.1	
	explained_variance		0.918	
	learning_rate		0.0003	
	loss		103	
	n_updates		4430	
	policy_gradient_loss		-0.058	
	std		0.708	
	value_loss		330	

---



---

	rollout/			
	ep_len_mean		80.1	
	ep_rew_mean		401	
	time/			
	fps		439	
	iterations		445	
	time_elapsed		2074	
	total_timesteps		911360	
	train/			
	approx_kl		0.06967526	
	clip_fraction		0.359	
	clip_range		0.2	
	entropy_loss		-18	
	explained_variance		0.873	
	learning_rate		0.0003	
	loss		151	
	n_updates		4440	
	policy_gradient_loss		-0.054	
	std		0.708	
	value_loss		329	

---

---

rollout/		
ep_len_mean	80.5	
ep_rew_mean	401	
time/		
fps	439	
iterations	446	
time_elapsed	2079	
total_timesteps	913408	
train/		
approx_kl	0.08475846	
clip_fraction	0.381	
clip_range	0.2	
entropy_loss	-18	
explained_variance	0.926	
learning_rate	0.0003	
loss	103	
n_updates	4450	
policy_gradient_loss	-0.0548	
std	0.707	
value_loss	263	

---



---

rollout/		
ep_len_mean	82	
ep_rew_mean	409	
time/		
fps	439	
iterations	447	
time_elapsed	2084	
total_timesteps	915456	
train/		
approx_kl	0.077369176	
clip_fraction	0.343	
clip_range	0.2	
entropy_loss	-18	
explained_variance	0.916	
learning_rate	0.0003	
loss	170	
n_updates	4460	
policy_gradient_loss	-0.0576	
std	0.707	
value_loss	343	

---



---

rollout/		
ep_len_mean	81.7	
ep_rew_mean	404	

---

time/		
fps	439	
iterations	448	
time_elapsed	2088	
total_timesteps	917504	
train/		
approx_kl	0.059179194	
clip_fraction	0.32	
clip_range	0.2	
entropy_loss	-18	
explained_variance	0.894	
learning_rate	0.0003	
loss	107	
n_updates	4470	
policy_gradient_loss	-0.0594	
std	0.708	
value_loss	424	

---

rollout/		
ep_len_mean	83.4	
ep_rew_mean	414	
time/		
fps	439	
iterations	449	
time_elapsed	2093	
total_timesteps	919552	
train/		
approx_kl	0.06622589	
clip_fraction	0.331	
clip_range	0.2	
entropy_loss	-18	
explained_variance	0.922	
learning_rate	0.0003	
loss	88.8	
n_updates	4480	
policy_gradient_loss	-0.0552	
std	0.707	
value_loss	301	

---

rollout/		
ep_len_mean	82.8	
ep_rew_mean	415	
time/		
fps	439	
iterations	450	
time_elapsed	2097	



	total_timesteps		921600	
	train/			
	approx_kl		0.058086224	
	clip_fraction		0.313	
	clip_range		0.2	
	entropy_loss		-18	
	explained_variance		0.902	
	learning_rate		0.0003	
	loss		150	
	n_updates		4490	
	policy_gradient_loss		-0.0561	
	std		0.707	
	value_loss		404	

---

	rollout/			
	ep_len_mean		80.6	
	ep_rew_mean		403	
	time/			
	fps		439	
	iterations		451	
	time_elapsed		2102	
	total_timesteps		923648	
	train/			
	approx_kl		0.086120784	
	clip_fraction		0.365	
	clip_range		0.2	
	entropy_loss		-18	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		159	
	n_updates		4500	
	policy_gradient_loss		-0.059	
	std		0.706	
	value_loss		323	

---

	rollout/			
	ep_len_mean		81.8	
	ep_rew_mean		410	
	time/			
	fps		439	
	iterations		452	
	time_elapsed		2107	
	total_timesteps		925696	
	train/			
	approx_kl		0.073065475	
	clip_fraction		0.356	

clip_range	0.2	
entropy_loss	-18	
explained_variance	0.92	
learning_rate	0.0003	
loss	100	
n_updates	4510	
policy_gradient_loss	-0.057	
std	0.705	
value_loss	278	

rollout/		
ep_len_mean	78.6	
ep_rew_mean	390	
time/		
fps	439	
iterations	453	
time_elapsed	2112	
total_timesteps	927744	
train/		
approx_kl	0.07116476	
clip_fraction	0.347	
clip_range	0.2	
entropy_loss	-18	
explained_variance	0.906	
learning_rate	0.0003	
loss	130	
n_updates	4520	
policy_gradient_loss	-0.0572	
std	0.705	
value_loss	333	

rollout/		
ep_len_mean	80.3	
ep_rew_mean	396	
time/		
fps	439	
iterations	454	
time_elapsed	2116	
total_timesteps	929792	
train/		
approx_kl	0.09735218	
clip_fraction	0.402	
clip_range	0.2	
entropy_loss	-18	
explained_variance	0.938	
learning_rate	0.0003	

	loss		64.4	
	n_updates		4530	
	policy_gradient_loss		-0.0513	
	std		0.704	
	value_loss		243	

	rollout/			
	ep_len_mean		82.7	
	ep_rew_mean		408	
	time/			
	fps		439	
	iterations		455	
	time_elapsed		2121	
	total_timesteps		931840	
	train/			
	approx_kl		0.06580991	
	clip_fraction		0.33	
	clip_range		0.2	
	entropy_loss		-18	
	explained_variance		0.884	
	learning_rate		0.0003	
	loss		131	
	n_updates		4540	
	policy_gradient_loss		-0.0587	
	std		0.704	
	value_loss		369	

	rollout/			
	ep_len_mean		81.9	
	ep_rew_mean		404	
	time/			
	fps		439	
	iterations		456	
	time_elapsed		2126	
	total_timesteps		933888	
	train/			
	approx_kl		0.075444534	
	clip_fraction		0.347	
	clip_range		0.2	
	entropy_loss		-17.9	
	explained_variance		0.92	
	learning_rate		0.0003	
	loss		85.2	
	n_updates		4550	
	policy_gradient_loss		-0.056	
	std		0.703	

	value_loss	288	
-----			
	rollout/		
	ep_len_mean	83.2	
	ep_rew_mean	412	
	time/		
	fps	438	
	iterations	457	
	time_elapsed	2133	
	total_timesteps	935936	
	train/		
	approx_kl	0.059479125	
	clip_fraction	0.344	
	clip_range	0.2	
	entropy_loss	-17.9	
	explained_variance	0.918	
	learning_rate	0.0003	
	loss	88.5	
	n_updates	4560	
	policy_gradient_loss	-0.061	
	std	0.703	
	value_loss	322	
-----			
	rollout/		
	ep_len_mean	83.2	
	ep_rew_mean	413	
	time/		
	fps	438	
	iterations	458	
	time_elapsed	2138	
	total_timesteps	937984	
	train/		
	approx_kl	0.067927815	
	clip_fraction	0.332	
	clip_range	0.2	
	entropy_loss	-17.9	
	explained_variance	0.916	
	learning_rate	0.0003	
	loss	137	
	n_updates	4570	
	policy_gradient_loss	-0.0536	
	std	0.702	
	value_loss	332	
-----			
	rollout/		

	ep_len_mean		83.5	
	ep_rew_mean		413	
	time/			
	fps		438	
	iterations		459	
	time_elapsed		2143	
	total_timesteps		940032	
	train/			
	approx_kl		0.08318827	
	clip_fraction		0.37	
	clip_range		0.2	
	entropy_loss		-17.9	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		105	
	n_updates		4580	
	policy_gradient_loss		-0.0542	
	std		0.702	
	value_loss		295	

	rollout/			
	ep_len_mean		83.6	
	ep_rew_mean		415	
	time/			
	fps		438	
	iterations		460	
	time_elapsed		2148	
	total_timesteps		942080	
	train/			
	approx_kl		0.08546106	
	clip_fraction		0.372	
	clip_range		0.2	
	entropy_loss		-17.9	
	explained_variance		0.931	
	learning_rate		0.0003	
	loss		94.8	
	n_updates		4590	
	policy_gradient_loss		-0.0623	
	std		0.703	
	value_loss		260	

	rollout/			
	ep_len_mean		85.4	
	ep_rew_mean		422	
	time/			
	fps		438	

iterations	461
time_elapsed	2153
total_timesteps	944128
train/	
approx_kl	0.07449889
clip_fraction	0.37
clip_range	0.2
entropy_loss	-17.9
explained_variance	0.932
learning_rate	0.0003
loss	129
n_updates	4600
policy_gradient_loss	-0.0596
std	0.702
value_loss	270

rollout/	
ep_len_mean	84.3
ep_rew_mean	418
time/	
fps	438
iterations	462
time_elapsed	2158
total_timesteps	946176
train/	
approx_kl	0.07129397
clip_fraction	0.338
clip_range	0.2
entropy_loss	-17.9
explained_variance	0.928
learning_rate	0.0003
loss	126
n_updates	4610
policy_gradient_loss	-0.0587
std	0.702
value_loss	308

rollout/	
ep_len_mean	84.2
ep_rew_mean	419
time/	
fps	438
iterations	463
time_elapsed	2163
total_timesteps	948224
train/	

approx_kl	0.08013459	
clip_fraction	0.375	
clip_range	0.2	
entropy_loss	-17.9	
explained_variance	0.913	
learning_rate	0.0003	
loss	93.7	
n_updates	4620	
policy_gradient_loss	-0.0587	
std	0.701	
value_loss	324	

rollout/		
ep_len_mean	84.2	
ep_rew_mean	419	
time/		
fps	438	
iterations	464	
time_elapsed	2168	
total_timesteps	950272	
train/		
approx_kl	0.081791684	
clip_fraction	0.353	
clip_range	0.2	
entropy_loss	-17.9	
explained_variance	0.931	
learning_rate	0.0003	
loss	113	
n_updates	4630	
policy_gradient_loss	-0.0611	
std	0.701	
value_loss	267	

rollout/		
ep_len_mean	87.3	
ep_rew_mean	434	
time/		
fps	438	
iterations	465	
time_elapsed	2173	
total_timesteps	952320	
train/		
approx_kl	0.0767144	
clip_fraction	0.334	
clip_range	0.2	
entropy_loss	-17.9	

	explained_variance		0.905	
	learning_rate		0.0003	
	loss		117	
	n_updates		4640	
	policy_gradient_loss		-0.0546	
	std		0.7	
	value_loss		339	

	rollout/			
	ep_len_mean		89.6	
	ep_rew_mean		447	
	time/			
	fps		438	
	iterations		466	
	time_elapsed		2178	
	total_timesteps		954368	
	train/			
	approx_kl		0.085646346	
	clip_fraction		0.362	
	clip_range		0.2	
	entropy_loss		-17.9	
	explained_variance		0.909	
	learning_rate		0.0003	
	loss		126	
	n_updates		4650	
	policy_gradient_loss		-0.0579	
	std		0.7	
	value_loss		369	

	rollout/			
	ep_len_mean		89.2	
	ep_rew_mean		445	
	time/			
	fps		438	
	iterations		467	
	time_elapsed		2183	
	total_timesteps		956416	
	train/			
	approx_kl		0.06890987	
	clip_fraction		0.35	
	clip_range		0.2	
	entropy_loss		-17.8	
	explained_variance		0.935	
	learning_rate		0.0003	
	loss		88.2	
	n_updates		4660	



	policy_gradient_loss		-0.0636	
	std		0.7	
	value_loss		267	

	rollout/			
	ep_len_mean		87.5	
	ep_rew_mean		436	
	time/			
	fps		437	
	iterations		468	
	time_elapsed		2188	
	total_timesteps		958464	
	train/			
	approx_kl		0.066124156	
	clip_fraction		0.325	
	clip_range		0.2	
	entropy_loss		-17.8	
	explained_variance		0.929	
	learning_rate		0.0003	
	loss		166	
	n_updates		4670	
	policy_gradient_loss		-0.0579	
	std		0.7	
	value_loss		292	

	rollout/			
	ep_len_mean		88	
	ep_rew_mean		441	
	time/			
	fps		438	
	iterations		469	
	time_elapsed		2192	
	total_timesteps		960512	
	train/			
	approx_kl		0.061573446	
	clip_fraction		0.32	
	clip_range		0.2	
	entropy_loss		-17.8	
	explained_variance		0.912	
	learning_rate		0.0003	
	loss		127	
	n_updates		4680	
	policy_gradient_loss		-0.0538	
	std		0.699	
	value_loss		376	

---

rollout/		
ep_len_mean	86.2	
ep_rew_mean	428	
time/		
fps	437	
iterations	470	
time_elapsed	2197	
total_timesteps	962560	
train/		
approx_kl	0.08282314	
clip_fraction	0.347	
clip_range	0.2	
entropy_loss	-17.8	
explained_variance	0.91	
learning_rate	0.0003	
loss	131	
n_updates	4690	
policy_gradient_loss	-0.0527	
std	0.699	
value_loss	360	

---



---

rollout/		
ep_len_mean	87.5	
ep_rew_mean	435	
time/		
fps	437	
iterations	471	
time_elapsed	2202	
total_timesteps	964608	
train/		
approx_kl	0.083549365	
clip_fraction	0.371	
clip_range	0.2	
entropy_loss	-17.8	
explained_variance	0.931	
learning_rate	0.0003	
loss	73.3	
n_updates	4700	
policy_gradient_loss	-0.0576	
std	0.699	
value_loss	277	

---



---

rollout/		
ep_len_mean	87.3	
ep_rew_mean	434	

---

time/		
fps	437	
iterations	472	
time_elapsed	2208	
total_timesteps	966656	
train/		
approx_kl	0.068529174	
clip_fraction	0.336	
clip_range	0.2	
entropy_loss	-17.8	
explained_variance	0.921	
learning_rate	0.0003	
loss	107	
n_updates	4710	
policy_gradient_loss	-0.0589	
std	0.699	
value_loss	302	

rollout/		
ep_len_mean	86.2	
ep_rew_mean	430	
time/		
fps	437	
iterations	473	
time_elapsed	2213	
total_timesteps	968704	
train/		
approx_kl	0.078921616	
clip_fraction	0.353	
clip_range	0.2	
entropy_loss	-17.8	
explained_variance	0.928	
learning_rate	0.0003	
loss	69.4	
n_updates	4720	
policy_gradient_loss	-0.0562	
std	0.698	
value_loss	251	

rollout/		
ep_len_mean	85.8	
ep_rew_mean	432	
time/		
fps	437	
iterations	474	
time_elapsed	2218	

	total_timesteps		970752	
	train/			
	approx_kl		0.070360094	
	clip_fraction		0.37	
	clip_range		0.2	
	entropy_loss		-17.8	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		88.5	
	n_updates		4730	
	policy_gradient_loss		-0.0604	
	std		0.699	
	value_loss		275	

	rollout/			
	ep_len_mean		85.5	
	ep_rew_mean		430	
	time/			
	fps		437	
	iterations		475	
	time_elapsed		2223	
	total_timesteps		972800	
	train/			
	approx_kl		0.06667427	
	clip_fraction		0.365	
	clip_range		0.2	
	entropy_loss		-17.8	
	explained_variance		0.92	
	learning_rate		0.0003	
	loss		126	
	n_updates		4740	
	policy_gradient_loss		-0.0597	
	std		0.698	
	value_loss		289	

	rollout/			
	ep_len_mean		86.5	
	ep_rew_mean		433	
	time/			
	fps		437	
	iterations		476	
	time_elapsed		2228	
	total_timesteps		974848	
	train/			
	approx_kl		0.07137889	
	clip_fraction		0.357	

clip_range	0.2	
entropy_loss	-17.8	
explained_variance	0.929	
learning_rate	0.0003	
loss	99.7	
n_updates	4750	
policy_gradient_loss	-0.0579	
std	0.697	
value_loss	262	

rollout/		
ep_len_mean	86.7	
ep_rew_mean	433	
time/		
fps	437	
iterations	477	
time_elapsed	2233	
total_timesteps	976896	
train/		
approx_kl	0.077442914	
clip_fraction	0.35	
clip_range	0.2	
entropy_loss	-17.8	
explained_variance	0.921	
learning_rate	0.0003	
loss	128	
n_updates	4760	
policy_gradient_loss	-0.0577	
std	0.696	
value_loss	324	

rollout/		
ep_len_mean	87.5	
ep_rew_mean	438	
time/		
fps	437	
iterations	478	
time_elapsed	2238	
total_timesteps	978944	
train/		
approx_kl	0.06906516	
clip_fraction	0.334	
clip_range	0.2	
entropy_loss	-17.7	
explained_variance	0.919	
learning_rate	0.0003	

	loss		59.3	
	n_updates		4770	
	policy_gradient_loss		-0.0576	
	std		0.696	
	value_loss		315	

	rollout/			
	ep_len_mean		86.3	
	ep_rew_mean		432	
	time/			
	fps		437	
	iterations		479	
	time_elapsed		2243	
	total_timesteps		980992	
	train/			
	approx_kl		0.08858864	
	clip_fraction		0.362	
	clip_range		0.2	
	entropy_loss		-17.7	
	explained_variance		0.927	
	learning_rate		0.0003	
	loss		97	
	n_updates		4780	
	policy_gradient_loss		-0.0561	
	std		0.696	
	value_loss		295	

	rollout/			
	ep_len_mean		85.8	
	ep_rew_mean		429	
	time/			
	fps		436	
	iterations		480	
	time_elapsed		2249	
	total_timesteps		983040	
	train/			
	approx_kl		0.073309235	
	clip_fraction		0.376	
	clip_range		0.2	
	entropy_loss		-17.7	
	explained_variance		0.925	
	learning_rate		0.0003	
	loss		124	
	n_updates		4790	
	policy_gradient_loss		-0.0565	
	std		0.695	

	value_loss	263	
-----			
	rollout/		
	ep_len_mean	85.7	
	ep_rew_mean	428	
	time/		
	fps	436	
	iterations	481	
	time_elapsed	2255	
	total_timesteps	985088	
	train/		
	approx_kl	0.08279662	
	clip_fraction	0.355	
	clip_range	0.2	
	entropy_loss	-17.7	
	explained_variance	0.917	
	learning_rate	0.0003	
	loss	187	
	n_updates	4800	
	policy_gradient_loss	-0.0567	
	std	0.694	
	value_loss	350	
-----			
	rollout/		
	ep_len_mean	83.1	
	ep_rew_mean	416	
	time/		
	fps	436	
	iterations	482	
	time_elapsed	2262	
	total_timesteps	987136	
	train/		
	approx_kl	0.084743425	
	clip_fraction	0.379	
	clip_range	0.2	
	entropy_loss	-17.7	
	explained_variance	0.938	
	learning_rate	0.0003	
	loss	87.4	
	n_updates	4810	
	policy_gradient_loss	-0.0618	
	std	0.693	
	value_loss	241	
-----			
	rollout/		

	ep_len_mean		84.2	
	ep_rew_mean		422	
	time/			
	fps		435	
	iterations		483	
	time_elapsed		2269	
	total_timesteps		989184	
	train/			
	approx_kl		0.07644686	
	clip_fraction		0.355	
	clip_range		0.2	
	entropy_loss		-17.7	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		92.1	
	n_updates		4820	
	policy_gradient_loss		-0.0575	
	std		0.692	
	value_loss		319	

	rollout/			
	ep_len_mean		86.8	
	ep_rew_mean		436	
	time/			
	fps		435	
	iterations		484	
	time_elapsed		2275	
	total_timesteps		991232	
	train/			
	approx_kl		0.10253276	
	clip_fraction		0.391	
	clip_range		0.2	
	entropy_loss		-17.6	
	explained_variance		0.929	
	learning_rate		0.0003	
	loss		78.4	
	n_updates		4830	
	policy_gradient_loss		-0.0506	
	std		0.691	
	value_loss		261	

	rollout/			
	ep_len_mean		86.4	
	ep_rew_mean		435	
	time/			
	fps		435	



iterations	485
time_elapsed	2279
total_timesteps	993280
train/	
approx_kl	0.09305653
clip_fraction	0.354
clip_range	0.2
entropy_loss	-17.6
explained_variance	0.91
learning_rate	0.0003
loss	134
n_updates	4840
policy_gradient_loss	-0.0518
std	0.69
value_loss	346

rollout/	
ep_len_mean	89
ep_rew_mean	449
time/	
fps	435
iterations	486
time_elapsed	2284
total_timesteps	995328
train/	
approx_kl	0.07498552
clip_fraction	0.343
clip_range	0.2
entropy_loss	-17.6
explained_variance	0.929
learning_rate	0.0003
loss	121
n_updates	4850
policy_gradient_loss	-0.0564
std	0.691
value_loss	283

rollout/	
ep_len_mean	90.6
ep_rew_mean	454
time/	
fps	435
iterations	487
time_elapsed	2289
total_timesteps	997376
train/	

approx_kl	0.078204036	
clip_fraction	0.368	
clip_range	0.2	
entropy_loss	-17.6	
explained_variance	0.937	
learning_rate	0.0003	
loss	97.1	
n_updates	4860	
policy_gradient_loss	-0.063	
std	0.691	
value_loss	272	

rollout/		
ep_len_mean	88.2	
ep_rew_mean	440	
time/		
fps	435	
iterations	488	
time_elapsed	2294	
total_timesteps	999424	
train/		
approx_kl	0.07978214	
clip_fraction	0.32	
clip_range	0.2	
entropy_loss	-17.6	
explained_variance	0.909	
learning_rate	0.0003	
loss	117	
n_updates	4870	
policy_gradient_loss	-0.0535	
std	0.69	
value_loss	353	

rollout/		
ep_len_mean	87.8	
ep_rew_mean	441	
time/		
fps	435	
iterations	489	
time_elapsed	2299	
total_timesteps	1001472	
train/		
approx_kl	0.070625834	
clip_fraction	0.338	
clip_range	0.2	
entropy_loss	-17.6	

	explained_variance		0.911	
	learning_rate		0.0003	
	loss		122	
	n_updates		4880	
	policy_gradient_loss		-0.0583	
	std		0.69	
	value_loss		327	

	rollout/			
	ep_len_mean		86.8	
	ep_rew_mean		436	
	time/			
	fps		435	
	iterations		490	
	time_elapsed		2305	
	total_timesteps		1003520	
	train/			
	approx_kl		0.062243864	
	clip_fraction		0.328	
	clip_range		0.2	
	entropy_loss		-17.6	
	explained_variance		0.922	
	learning_rate		0.0003	
	loss		121	
	n_updates		4890	
	policy_gradient_loss		-0.0598	
	std		0.69	
	value_loss		308	

	rollout/			
	ep_len_mean		85.5	
	ep_rew_mean		431	
	time/			
	fps		435	
	iterations		491	
	time_elapsed		2310	
	total_timesteps		1005568	
	train/			
	approx_kl		0.08686638	
	clip_fraction		0.354	
	clip_range		0.2	
	entropy_loss		-17.6	
	explained_variance		0.927	
	learning_rate		0.0003	
	loss		91.2	
	n_updates		4900	

	policy_gradient_loss		-0.0537	
	std		0.689	
	value_loss		285	

---

	rollout/			
	ep_len_mean		85	
	ep_rew_mean		430	
	time/			
	fps		435	
	iterations		492	
	time_elapsed		2315	
	total_timesteps		1007616	
	train/			
	approx_kl		0.0840717	
	clip_fraction		0.385	
	clip_range		0.2	
	entropy_loss		-17.6	
	explained_variance		0.929	
	learning_rate		0.0003	
	loss		93.1	
	n_updates		4910	
	policy_gradient_loss		-0.0594	
	std		0.689	
	value_loss		295	

---



---

	rollout/			
	ep_len_mean		87.5	
	ep_rew_mean		442	
	time/			
	fps		434	
	iterations		493	
	time_elapsed		2324	
	total_timesteps		1009664	
	train/			
	approx_kl		0.08653122	
	clip_fraction		0.353	
	clip_range		0.2	
	entropy_loss		-17.5	
	explained_variance		0.918	
	learning_rate		0.0003	
	loss		124	
	n_updates		4920	
	policy_gradient_loss		-0.0588	
	std		0.689	
	value_loss		280	

---

---

rollout/		
ep_len_mean	86.8	
ep_rew_mean	437	
time/		
fps	432	
iterations	494	
time_elapsed	2338	
total_timesteps	1011712	
train/		
approx_kl	0.074489646	
clip_fraction	0.365	
clip_range	0.2	
entropy_loss	-17.5	
explained_variance	0.935	
learning_rate	0.0003	
loss	82.7	
n_updates	4930	
policy_gradient_loss	-0.0572	
std	0.688	
value_loss	250	

---



---

rollout/		
ep_len_mean	84.8	
ep_rew_mean	426	
time/		
fps	432	
iterations	495	
time_elapsed	2343	
total_timesteps	1013760	
train/		
approx_kl	0.08429592	
clip_fraction	0.372	
clip_range	0.2	
entropy_loss	-17.5	
explained_variance	0.941	
learning_rate	0.0003	
loss	108	
n_updates	4940	
policy_gradient_loss	-0.0593	
std	0.687	
value_loss	258	

---



---

rollout/		
ep_len_mean	87.8	
ep_rew_mean	440	

---

time/		
fps	432	
iterations	496	
time_elapsed	2348	
total_timesteps	1015808	
train/		
approx_kl	0.08859577	
clip_fraction	0.384	
clip_range	0.2	
entropy_loss	-17.5	
explained_variance	0.939	
learning_rate	0.0003	
loss	71.5	
n_updates	4950	
policy_gradient_loss	-0.0603	
std	0.686	
value_loss	261	

---

rollout/		
ep_len_mean	88.1	
ep_rew_mean	442	
time/		
fps	432	
iterations	497	
time_elapsed	2353	
total_timesteps	1017856	
train/		
approx_kl	0.07100342	
clip_fraction	0.338	
clip_range	0.2	
entropy_loss	-17.5	
explained_variance	0.905	
learning_rate	0.0003	
loss	125	
n_updates	4960	
policy_gradient_loss	-0.0565	
std	0.686	
value_loss	344	

---

rollout/		
ep_len_mean	89.3	
ep_rew_mean	447	
time/		
fps	432	
iterations	498	
time_elapsed	2358	

	total_timesteps		1019904	
	train/			
	approx_kl		0.09599799	
	clip_fraction		0.379	
	clip_range		0.2	
	entropy_loss		-17.5	
	explained_variance		0.923	
	learning_rate		0.0003	
	loss		107	
	n_updates		4970	
	policy_gradient_loss		-0.0521	
	std		0.685	
	value_loss		317	

	rollout/			
	ep_len_mean		89	
	ep_rew_mean		445	
	time/			
	fps		432	
	iterations		499	
	time_elapsed		2363	
	total_timesteps		1021952	
	train/			
	approx_kl		0.08097498	
	clip_fraction		0.339	
	clip_range		0.2	
	entropy_loss		-17.5	
	explained_variance		0.935	
	learning_rate		0.0003	
	loss		94.6	
	n_updates		4980	
	policy_gradient_loss		-0.0594	
	std		0.685	
	value_loss		299	

	rollout/			
	ep_len_mean		86.6	
	ep_rew_mean		434	
	time/			
	fps		432	
	iterations		500	
	time_elapsed		2368	
	total_timesteps		1024000	
	train/			
	approx_kl		0.0794995	
	clip_fraction		0.385	

clip_range	0.2	
entropy_loss	-17.4	
explained_variance	0.934	
learning_rate	0.0003	
loss	135	
n_updates	4990	
policy_gradient_loss	-0.0621	
std	0.684	
value_loss	282	

rollout/		
ep_len_mean	85.1	
ep_rew_mean	425	
time/		
fps	432	
iterations	501	
time_elapsed	2373	
total_timesteps	1026048	
train/		
approx_kl	0.078638434	
clip_fraction	0.365	
clip_range	0.2	
entropy_loss	-17.4	
explained_variance	0.931	
learning_rate	0.0003	
loss	106	
n_updates	5000	
policy_gradient_loss	-0.0586	
std	0.682	
value_loss	284	

rollout/		
ep_len_mean	84.7	
ep_rew_mean	424	
time/		
fps	432	
iterations	502	
time_elapsed	2377	
total_timesteps	1028096	
train/		
approx_kl	0.075044595	
clip_fraction	0.344	
clip_range	0.2	
entropy_loss	-17.4	
explained_variance	0.918	
learning_rate	0.0003	



	loss		157	
	n_updates		5010	
	policy_gradient_loss		-0.0573	
	std		0.681	
	value_loss		321	

---

	rollout/			
	ep_len_mean		84.7	
	ep_rew_mean		422	
	time/			
	fps		432	
	iterations		503	
	time_elapsed		2382	
	total_timesteps		1030144	
	train/			
	approx_kl		0.07187837	
	clip_fraction		0.332	
	clip_range		0.2	
	entropy_loss		-17.3	
	explained_variance		0.913	
	learning_rate		0.0003	
	loss		80.5	
	n_updates		5020	
	policy_gradient_loss		-0.0521	
	std		0.68	
	value_loss		317	

---

	rollout/			
	ep_len_mean		84.8	
	ep_rew_mean		420	
	time/			
	fps		432	
	iterations		504	
	time_elapsed		2387	
	total_timesteps		1032192	
	train/			
	approx_kl		0.09115274	
	clip_fraction		0.382	
	clip_range		0.2	
	entropy_loss		-17.3	
	explained_variance		0.944	
	learning_rate		0.0003	
	loss		65.5	
	n_updates		5030	
	policy_gradient_loss		-0.0601	
	std		0.68	

	value_loss		237	
-----				
	rollout/			
	ep_len_mean		85.8	
	ep_rew_mean		430	
	time/			
	fps		432	
	iterations		505	
	time_elapsed		2392	
	total_timesteps		1034240	
	train/			
	approx_kl		0.07985607	
	clip_fraction		0.38	
	clip_range		0.2	
	entropy_loss		-17.3	
	explained_variance		0.931	
	learning_rate		0.0003	
	loss		72.7	
	n_updates		5040	
	policy_gradient_loss		-0.0545	
	std		0.68	
	value_loss		291	
-----				
	rollout/			
	ep_len_mean		86.5	
	ep_rew_mean		432	
	time/			
	fps		432	
	iterations		506	
	time_elapsed		2397	
	total_timesteps		1036288	
	train/			
	approx_kl		0.07801108	
	clip_fraction		0.361	
	clip_range		0.2	
	entropy_loss		-17.3	
	explained_variance		0.919	
	learning_rate		0.0003	
	loss		85.1	
	n_updates		5050	
	policy_gradient_loss		-0.0558	
	std		0.679	
	value_loss		296	
-----				
	rollout/			

	ep_len_mean		86.7	
	ep_rew_mean		434	
	time/			
	fps		432	
	iterations		507	
	time_elapsed		2402	
	total_timesteps		1038336	
	train/			
	approx_kl		0.059672557	
	clip_fraction		0.333	
	clip_range		0.2	
	entropy_loss		-17.3	
	explained_variance		0.923	
	learning_rate		0.0003	
	loss		93.3	
	n_updates		5060	
	policy_gradient_loss		-0.0565	
	std		0.679	
	value_loss		323	

	rollout/			
	ep_len_mean		86.8	
	ep_rew_mean		438	
	time/			
	fps		432	
	iterations		508	
	time_elapsed		2407	
	total_timesteps		1040384	
	train/			
	approx_kl		0.08698717	
	clip_fraction		0.374	
	clip_range		0.2	
	entropy_loss		-17.3	
	explained_variance		0.926	
	learning_rate		0.0003	
	loss		135	
	n_updates		5070	
	policy_gradient_loss		-0.0535	
	std		0.678	
	value_loss		291	

	rollout/			
	ep_len_mean		86.4	
	ep_rew_mean		433	
	time/			
	fps		432	

iterations	509
time_elapsed	2412
total_timesteps	1042432
train/	
approx_kl	0.08437585
clip_fraction	0.389
clip_range	0.2
entropy_loss	-17.3
explained_variance	0.931
learning_rate	0.0003
loss	123
n_updates	5080
policy_gradient_loss	-0.0559
std	0.678
value_loss	264

rollout/	
ep_len_mean	86.9
ep_rew_mean	436
time/	
fps	432
iterations	510
time_elapsed	2417
total_timesteps	1044480
train/	
approx_kl	0.07923706
clip_fraction	0.355
clip_range	0.2
entropy_loss	-17.3
explained_variance	0.935
learning_rate	0.0003
loss	109
n_updates	5090
policy_gradient_loss	-0.0578
std	0.677
value_loss	281

rollout/	
ep_len_mean	87.4
ep_rew_mean	441
time/	
fps	432
iterations	511
time_elapsed	2422
total_timesteps	1046528
train/	

approx_kl	0.07049539	
clip_fraction	0.348	
clip_range	0.2	
entropy_loss	-17.3	
explained_variance	0.922	
learning_rate	0.0003	
loss	90.7	
n_updates	5100	
policy_gradient_loss	-0.058	
std	0.678	
value_loss	318	

rollout/		
ep_len_mean	89.5	
ep_rew_mean	452	
time/		
fps	431	
iterations	512	
time_elapsed	2427	
total_timesteps	1048576	
train/		
approx_kl	0.08377897	
clip_fraction	0.377	
clip_range	0.2	
entropy_loss	-17.3	
explained_variance	0.932	
learning_rate	0.0003	
loss	109	
n_updates	5110	
policy_gradient_loss	-0.0563	
std	0.677	
value_loss	317	

rollout/		
ep_len_mean	90.5	
ep_rew_mean	458	
time/		
fps	431	
iterations	513	
time_elapsed	2432	
total_timesteps	1050624	
train/		
approx_kl	0.10383383	
clip_fraction	0.388	
clip_range	0.2	
entropy_loss	-17.2	

	explained_variance		0.933	
	learning_rate		0.0003	
	loss		79.2	
	n_updates		5120	
	policy_gradient_loss		-0.0532	
	std		0.676	
	value_loss		277	

	rollout/			
	ep_len_mean		92.3	
	ep_rew_mean		467	
	time/			
	fps		431	
	iterations		514	
	time_elapsed		2437	
	total_timesteps		1052672	
	train/			
	approx_kl		0.08095491	
	clip_fraction		0.349	
	clip_range		0.2	
	entropy_loss		-17.2	
	explained_variance		0.92	
	learning_rate		0.0003	
	loss		136	
	n_updates		5130	
	policy_gradient_loss		-0.0548	
	std		0.675	
	value_loss		297	

	rollout/			
	ep_len_mean		90.3	
	ep_rew_mean		457	
	time/			
	fps		431	
	iterations		515	
	time_elapsed		2442	
	total_timesteps		1054720	
	train/			
	approx_kl		0.07116136	
	clip_fraction		0.34	
	clip_range		0.2	
	entropy_loss		-17.2	
	explained_variance		0.927	
	learning_rate		0.0003	
	loss		117	
	n_updates		5140	

	policy_gradient_loss		-0.0537	
	std		0.675	
	value_loss		299	

---

	rollout/			
	ep_len_mean		88.3	
	ep_rew_mean		446	
	time/			
	fps		431	
	iterations		516	
	time_elapsed		2447	
	total_timesteps		1056768	
	train/			
	approx_kl		0.08773682	
	clip_fraction		0.346	
	clip_range		0.2	
	entropy_loss		-17.2	
	explained_variance		0.925	
	learning_rate		0.0003	
	loss		79.2	
	n_updates		5150	
	policy_gradient_loss		-0.0566	
	std		0.675	
	value_loss		329	

---



---

	rollout/			
	ep_len_mean		88.3	
	ep_rew_mean		445	
	time/			
	fps		431	
	iterations		517	
	time_elapsed		2452	
	total_timesteps		1058816	
	train/			
	approx_kl		0.093651414	
	clip_fraction		0.374	
	clip_range		0.2	
	entropy_loss		-17.2	
	explained_variance		0.939	
	learning_rate		0.0003	
	loss		90	
	n_updates		5160	
	policy_gradient_loss		-0.0528	
	std		0.676	
	value_loss		291	

---

---

rollout/		
ep_len_mean	85.6	
ep_rew_mean	433	
time/		
fps	431	
iterations	518	
time_elapsed	2457	
total_timesteps	1060864	
train/		
approx_kl	0.08649035	
clip_fraction	0.379	
clip_range	0.2	
entropy_loss	-17.2	
explained_variance	0.937	
learning_rate	0.0003	
loss	95.8	
n_updates	5170	
policy_gradient_loss	-0.0562	
std	0.675	
value_loss	264	

---



---

rollout/		
ep_len_mean	86.7	
ep_rew_mean	439	
time/		
fps	431	
iterations	519	
time_elapsed	2462	
total_timesteps	1062912	
train/		
approx_kl	0.076252796	
clip_fraction	0.369	
clip_range	0.2	
entropy_loss	-17.2	
explained_variance	0.937	
learning_rate	0.0003	
loss	91.3	
n_updates	5180	
policy_gradient_loss	-0.0542	
std	0.674	
value_loss	242	

---



---

rollout/		
ep_len_mean	89	
ep_rew_mean	452	

---



time/		
fps	431	
iterations	520	
time_elapsed	2467	
total_timesteps	1064960	
train/		
approx_kl	0.09112887	
clip_fraction	0.386	
clip_range	0.2	
entropy_loss	-17.2	
explained_variance	0.929	
learning_rate	0.0003	
loss	107	
n_updates	5190	
policy_gradient_loss	-0.0557	
std	0.674	
value_loss	272	

rollout/		
ep_len_mean	88.2	
ep_rew_mean	446	
time/		
fps	431	
iterations	521	
time_elapsed	2472	
total_timesteps	1067008	
train/		
approx_kl	0.09074654	
clip_fraction	0.357	
clip_range	0.2	
entropy_loss	-17.2	
explained_variance	0.941	
learning_rate	0.0003	
loss	104	
n_updates	5200	
policy_gradient_loss	-0.0551	
std	0.674	
value_loss	253	

rollout/		
ep_len_mean	88.3	
ep_rew_mean	444	
time/		
fps	431	
iterations	522	
time_elapsed	2477	

	total_timesteps		1069056	
	train/			
	approx_kl		0.09931326	
	clip_fraction		0.407	
	clip_range		0.2	
	entropy_loss		-17.1	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		126	
	n_updates		5210	
	policy_gradient_loss		-0.0591	
	std		0.673	
	value_loss		248	

	rollout/			
	ep_len_mean		87	
	ep_rew_mean		438	
	time/			
	fps		431	
	iterations		523	
	time_elapsed		2482	
	total_timesteps		1071104	
	train/			
	approx_kl		0.079299286	
	clip_fraction		0.366	
	clip_range		0.2	
	entropy_loss		-17.1	
	explained_variance		0.931	
	learning_rate		0.0003	
	loss		120	
	n_updates		5220	
	policy_gradient_loss		-0.0537	
	std		0.674	
	value_loss		282	

	rollout/			
	ep_len_mean		87.7	
	ep_rew_mean		440	
	time/			
	fps		431	
	iterations		524	
	time_elapsed		2487	
	total_timesteps		1073152	
	train/			
	approx_kl		0.07213847	
	clip_fraction		0.329	

clip_range	0.2	
entropy_loss	-17.1	
explained_variance	0.935	
learning_rate	0.0003	
loss	78.3	
n_updates	5230	
policy_gradient_loss	-0.0607	
std	0.672	
value_loss	276	

rollout/		
ep_len_mean	86.2	
ep_rew_mean	433	
time/		
fps	431	
iterations	525	
time_elapsed	2493	
total_timesteps	1075200	
train/		
approx_kl	0.07628214	
clip_fraction	0.341	
clip_range	0.2	
entropy_loss	-17.1	
explained_variance	0.923	
learning_rate	0.0003	
loss	110	
n_updates	5240	
policy_gradient_loss	-0.058	
std	0.672	
value_loss	326	

rollout/		
ep_len_mean	87.3	
ep_rew_mean	438	
time/		
fps	431	
iterations	526	
time_elapsed	2498	
total_timesteps	1077248	
train/		
approx_kl	0.11069459	
clip_fraction	0.414	
clip_range	0.2	
entropy_loss	-17.1	
explained_variance	0.948	
learning_rate	0.0003	

	loss		80.7	
	n_updates		5250	
	policy_gradient_loss		-0.0572	
	std		0.672	
	value_loss		227	

	rollout/			
	ep_len_mean		89.1	
	ep_rew_mean		445	
	time/			
	fps		431	
	iterations		527	
	time_elapsed		2503	
	total_timesteps		1079296	
	train/			
	approx_kl		0.08294566	
	clip_fraction		0.393	
	clip_range		0.2	
	entropy_loss		-17.1	
	explained_variance		0.935	
	learning_rate		0.0003	
	loss		91.4	
	n_updates		5260	
	policy_gradient_loss		-0.052	
	std		0.671	
	value_loss		282	

	rollout/			
	ep_len_mean		88.2	
	ep_rew_mean		443	
	time/			
	fps		431	
	iterations		528	
	time_elapsed		2508	
	total_timesteps		1081344	
	train/			
	approx_kl		0.08628504	
	clip_fraction		0.361	
	clip_range		0.2	
	entropy_loss		-17.1	
	explained_variance		0.931	
	learning_rate		0.0003	
	loss		84.5	
	n_updates		5270	
	policy_gradient_loss		-0.0557	
	std		0.67	

	value_loss	282	
-----			
	rollout/		
	ep_len_mean	91	
	ep_rew_mean	456	
	time/		
	fps	431	
	iterations	529	
	time_elapsed	2512	
	total_timesteps	1083392	
	train/		
	approx_kl	0.073115274	
	clip_fraction	0.337	
	clip_range	0.2	
	entropy_loss	-17	
	explained_variance	0.914	
	learning_rate	0.0003	
	loss	171	
	n_updates	5280	
	policy_gradient_loss	-0.0555	
	std	0.669	
	value_loss	360	
-----			
	rollout/		
	ep_len_mean	93.7	
	ep_rew_mean	469	
	time/		
	fps	431	
	iterations	530	
	time_elapsed	2518	
	total_timesteps	1085440	
	train/		
	approx_kl	0.09974989	
	clip_fraction	0.387	
	clip_range	0.2	
	entropy_loss	-17	
	explained_variance	0.925	
	learning_rate	0.0003	
	loss	116	
	n_updates	5290	
	policy_gradient_loss	-0.0467	
	std	0.669	
	value_loss	303	
-----			
	rollout/		

	ep_len_mean		92.3	
	ep_rew_mean		463	
	time/			
	fps		431	
	iterations		531	
	time_elapsed		2523	
	total_timesteps		1087488	
	train/			
	approx_kl		0.088874415	
	clip_fraction		0.364	
	clip_range		0.2	
	entropy_loss		-17	
	explained_variance		0.932	
	learning_rate		0.0003	
	loss		115	
	n_updates		5300	
	policy_gradient_loss		-0.0534	
	std		0.669	
	value_loss		299	

---

	rollout/			
	ep_len_mean		90.7	
	ep_rew_mean		456	
	time/			
	fps		430	
	iterations		532	
	time_elapsed		2527	
	total_timesteps		1089536	
	train/			
	approx_kl		0.10205275	
	clip_fraction		0.401	
	clip_range		0.2	
	entropy_loss		-17	
	explained_variance		0.935	
	learning_rate		0.0003	
	loss		88.5	
	n_updates		5310	
	policy_gradient_loss		-0.055	
	std		0.669	
	value_loss		272	

---

	rollout/			
	ep_len_mean		88.7	
	ep_rew_mean		446	
	time/			
	fps		430	

iterations	533
time_elapsed	2534
total_timesteps	1091584
train/	
approx_kl	0.07506962
clip_fraction	0.331
clip_range	0.2
entropy_loss	-17
explained_variance	0.92
learning_rate	0.0003
loss	86.3
n_updates	5320
policy_gradient_loss	-0.0548
std	0.669
value_loss	368

rollout/	
ep_len_mean	85.8
ep_rew_mean	432
time/	
fps	430
iterations	534
time_elapsed	2539
total_timesteps	1093632
train/	
approx_kl	0.09652905
clip_fraction	0.367
clip_range	0.2
entropy_loss	-17
explained_variance	0.914
learning_rate	0.0003
loss	90.7
n_updates	5330
policy_gradient_loss	-0.0476
std	0.669
value_loss	274

rollout/	
ep_len_mean	85.1
ep_rew_mean	428
time/	
fps	430
iterations	535
time_elapsed	2543
total_timesteps	1095680
train/	

approx_kl	0.113457106	
clip_fraction	0.397	
clip_range	0.2	
entropy_loss	-17	
explained_variance	0.946	
learning_rate	0.0003	
loss	68.9	
n_updates	5340	
policy_gradient_loss	-0.0521	
std	0.669	
value_loss	218	

rollout/		
ep_len_mean	86.6	
ep_rew_mean	433	
time/		
fps	430	
iterations	536	
time_elapsed	2548	
total_timesteps	1097728	
train/		
approx_kl	0.0961183	
clip_fraction	0.393	
clip_range	0.2	
entropy_loss	-17	
explained_variance	0.919	
learning_rate	0.0003	
loss	86.6	
n_updates	5350	
policy_gradient_loss	-0.0551	
std	0.669	
value_loss	273	

rollout/		
ep_len_mean	87.8	
ep_rew_mean	436	
time/		
fps	430	
iterations	537	
time_elapsed	2554	
total_timesteps	1099776	
train/		
approx_kl	0.08868196	
clip_fraction	0.373	
clip_range	0.2	
entropy_loss	-17	



	explained_variance		0.927	
	learning_rate		0.0003	
	loss		113	
	n_updates		5360	
	policy_gradient_loss		-0.0494	
	std		0.668	
	value_loss		295	

	rollout/			
	ep_len_mean		88.9	
	ep_rew_mean		444	
	time/			
	fps		430	
	iterations		538	
	time_elapsed		2559	
	total_timesteps		1101824	
	train/			
	approx_kl		0.114243746	
	clip_fraction		0.425	
	clip_range		0.2	
	entropy_loss		-17	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		77.9	
	n_updates		5370	
	policy_gradient_loss		-0.0583	
	std		0.668	
	value_loss		211	

	rollout/			
	ep_len_mean		90.7	
	ep_rew_mean		454	
	time/			
	fps		430	
	iterations		539	
	time_elapsed		2564	
	total_timesteps		1103872	
	train/			
	approx_kl		0.09514439	
	clip_fraction		0.368	
	clip_range		0.2	
	entropy_loss		-17	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		107	
	n_updates		5380	

	policy_gradient_loss		-0.0544	
	std		0.668	
	value_loss		286	

	rollout/			
	ep_len_mean		90.1	
	ep_rew_mean		454	
	time/			
	fps		430	
	iterations		540	
	time_elapsed		2569	
	total_timesteps		1105920	
	train/			
	approx_kl		0.096515305	
	clip_fraction		0.384	
	clip_range		0.2	
	entropy_loss		-17	
	explained_variance		0.92	
	learning_rate		0.0003	
	loss		79.6	
	n_updates		5390	
	policy_gradient_loss		-0.0496	
	std		0.666	
	value_loss		322	

	rollout/			
	ep_len_mean		89.2	
	ep_rew_mean		448	
	time/			
	fps		430	
	iterations		541	
	time_elapsed		2574	
	total_timesteps		1107968	
	train/			
	approx_kl		0.08724432	
	clip_fraction		0.355	
	clip_range		0.2	
	entropy_loss		-16.9	
	explained_variance		0.94	
	learning_rate		0.0003	
	loss		106	
	n_updates		5400	
	policy_gradient_loss		-0.0504	
	std		0.666	
	value_loss		269	

---

rollout/		
ep_len_mean	90.1	
ep_rew_mean	451	
time/		
fps	430	
iterations	542	
time_elapsed	2579	
total_timesteps	1110016	
train/		
approx_kl	0.09737831	
clip_fraction	0.361	
clip_range	0.2	
entropy_loss	-16.9	
explained_variance	0.926	
learning_rate	0.0003	
loss	117	
n_updates	5410	
policy_gradient_loss	-0.0564	
std	0.666	
value_loss	327	

---



---

rollout/		
ep_len_mean	93.2	
ep_rew_mean	464	
time/		
fps	430	
iterations	543	
time_elapsed	2583	
total_timesteps	1112064	
train/		
approx_kl	0.09691867	
clip_fraction	0.399	
clip_range	0.2	
entropy_loss	-16.9	
explained_variance	0.92	
learning_rate	0.0003	
loss	127	
n_updates	5420	
policy_gradient_loss	-0.0568	
std	0.666	
value_loss	313	

---



---

rollout/		
ep_len_mean	91.1	
ep_rew_mean	454	

---

time/		
fps	430	
iterations	544	
time_elapsed	2588	
total_timesteps	1114112	
train/		
approx_kl	0.08606841	
clip_fraction	0.392	
clip_range	0.2	
entropy_loss	-16.9	
explained_variance	0.929	
learning_rate	0.0003	
loss	168	
n_updates	5430	
policy_gradient_loss	-0.0599	
std	0.666	
value_loss	291	

rollout/		
ep_len_mean	90.9	
ep_rew_mean	453	
time/		
fps	430	
iterations	545	
time_elapsed	2593	
total_timesteps	1116160	
train/		
approx_kl	0.07244842	
clip_fraction	0.363	
clip_range	0.2	
entropy_loss	-16.9	
explained_variance	0.926	
learning_rate	0.0003	
loss	86.2	
n_updates	5440	
policy_gradient_loss	-0.0607	
std	0.666	
value_loss	290	

rollout/		
ep_len_mean	91.1	
ep_rew_mean	453	
time/		
fps	430	
iterations	546	
time_elapsed	2598	

	total_timesteps		1118208	
	train/			
	approx_kl		0.0951781	
	clip_fraction		0.372	
	clip_range		0.2	
	entropy_loss		-16.9	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		69.4	
	n_updates		5450	
	policy_gradient_loss		-0.0553	
	std		0.666	
	value_loss		265	

	rollout/			
	ep_len_mean		88.6	
	ep_rew_mean		443	
	time/			
	fps		430	
	iterations		547	
	time_elapsed		2603	
	total_timesteps		1120256	
	train/			
	approx_kl		0.0923885	
	clip_fraction		0.389	
	clip_range		0.2	
	entropy_loss		-16.9	
	explained_variance		0.932	
	learning_rate		0.0003	
	loss		84	
	n_updates		5460	
	policy_gradient_loss		-0.057	
	std		0.665	
	value_loss		250	

	rollout/			
	ep_len_mean		88.4	
	ep_rew_mean		442	
	time/			
	fps		430	
	iterations		548	
	time_elapsed		2607	
	total_timesteps		1122304	
	train/			
	approx_kl		0.077675164	
	clip_fraction		0.366	

clip_range	0.2	
entropy_loss	-16.9	
explained_variance	0.933	
learning_rate	0.0003	
loss	87.7	
n_updates	5470	
policy_gradient_loss	-0.0597	
std	0.665	
value_loss	272	

rollout/		
ep_len_mean	91	
ep_rew_mean	453	
time/		
fps	430	
iterations	549	
time_elapsed	2612	
total_timesteps	1124352	
train/		
approx_kl	0.08131513	
clip_fraction	0.382	
clip_range	0.2	
entropy_loss	-16.9	
explained_variance	0.927	
learning_rate	0.0003	
loss	93.2	
n_updates	5480	
policy_gradient_loss	-0.0544	
std	0.665	
value_loss	290	

rollout/		
ep_len_mean	92.2	
ep_rew_mean	459	
time/		
fps	430	
iterations	550	
time_elapsed	2617	
total_timesteps	1126400	
train/		
approx_kl	0.08350122	
clip_fraction	0.382	
clip_range	0.2	
entropy_loss	-16.9	
explained_variance	0.921	
learning_rate	0.0003	

	loss		137	
	n_updates		5490	
	policy_gradient_loss		-0.0592	
	std		0.664	
	value_loss		300	

	rollout/			
	ep_len_mean		90.8	
	ep_rew_mean		454	
	time/			
	fps		430	
	iterations		551	
	time_elapsed		2622	
	total_timesteps		1128448	
	train/			
	approx_kl		0.09384123	
	clip_fraction		0.365	
	clip_range		0.2	
	entropy_loss		-16.9	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		94	
	n_updates		5500	
	policy_gradient_loss		-0.0571	
	std		0.664	
	value_loss		304	

	rollout/			
	ep_len_mean		90.8	
	ep_rew_mean		453	
	time/			
	fps		430	
	iterations		552	
	time_elapsed		2627	
	total_timesteps		1130496	
	train/			
	approx_kl		0.08892276	
	clip_fraction		0.401	
	clip_range		0.2	
	entropy_loss		-16.9	
	explained_variance		0.938	
	learning_rate		0.0003	
	loss		90.8	
	n_updates		5510	
	policy_gradient_loss		-0.0554	
	std		0.664	

	value_loss		263	
-----				
	rollout/			
	ep_len_mean		90	
	ep_rew_mean		450	
	time/			
	fps		430	
	iterations		553	
	time_elapsed		2632	
	total_timesteps		1132544	
	train/			
	approx_kl		0.09438517	
	clip_fraction		0.374	
	clip_range		0.2	
	entropy_loss		-16.9	
	explained_variance		0.92	
	learning_rate		0.0003	
	loss		80.7	
	n_updates		5520	
	policy_gradient_loss		-0.0591	
	std		0.664	
	value_loss		333	
-----				
	rollout/			
	ep_len_mean		86.5	
	ep_rew_mean		432	
	time/			
	fps		430	
	iterations		554	
	time_elapsed		2637	
	total_timesteps		1134592	
	train/			
	approx_kl		0.0940791	
	clip_fraction		0.395	
	clip_range		0.2	
	entropy_loss		-16.8	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		117	
	n_updates		5530	
	policy_gradient_loss		-0.0569	
	std		0.662	
	value_loss		307	
-----				
	rollout/			



	ep_len_mean		88.8	
	ep_rew_mean		442	
	time/			
	fps		430	
	iterations		555	
	time_elapsed		2642	
	total_timesteps		1136640	
	train/			
	approx_kl		0.08366933	
	clip_fraction		0.353	
	clip_range		0.2	
	entropy_loss		-16.8	
	explained_variance		0.929	
	learning_rate		0.0003	
	loss		114	
	n_updates		5540	
	policy_gradient_loss		-0.055	
	std		0.662	
	value_loss		330	

	rollout/			
	ep_len_mean		88.6	
	ep_rew_mean		442	
	time/			
	fps		430	
	iterations		556	
	time_elapsed		2647	
	total_timesteps		1138688	
	train/			
	approx_kl		0.08720486	
	clip_fraction		0.357	
	clip_range		0.2	
	entropy_loss		-16.8	
	explained_variance		0.91	
	learning_rate		0.0003	
	loss		121	
	n_updates		5550	
	policy_gradient_loss		-0.0539	
	std		0.661	
	value_loss		403	

	rollout/			
	ep_len_mean		88.5	
	ep_rew_mean		440	
	time/			
	fps		430	

iterations	557
time_elapsed	2652
total_timesteps	1140736
train/	
approx_kl	0.08996601
clip_fraction	0.363
clip_range	0.2
entropy_loss	-16.8
explained_variance	0.936
learning_rate	0.0003
loss	84.6
n_updates	5560
policy_gradient_loss	-0.0533
std	0.66
value_loss	253

rollout/	
ep_len_mean	88.9
ep_rew_mean	444
time/	
fps	430
iterations	558
time_elapsed	2657
total_timesteps	1142784
train/	
approx_kl	0.0710542
clip_fraction	0.344
clip_range	0.2
entropy_loss	-16.8
explained_variance	0.925
learning_rate	0.0003
loss	110
n_updates	5570
policy_gradient_loss	-0.054
std	0.66
value_loss	345

rollout/	
ep_len_mean	90.2
ep_rew_mean	453
time/	
fps	429
iterations	559
time_elapsed	2662
total_timesteps	1144832
train/	

approx_kl	0.10262425	
clip_fraction	0.393	
clip_range	0.2	
entropy_loss	-16.8	
explained_variance	0.949	
learning_rate	0.0003	
loss	77.9	
n_updates	5580	
policy_gradient_loss	-0.0606	
std	0.66	
value_loss	235	

rollout/		
ep_len_mean	90.2	
ep_rew_mean	454	
time/		
fps	429	
iterations	560	
time_elapsed	2667	
total_timesteps	1146880	
train/		
approx_kl	0.12096631	
clip_fraction	0.412	
clip_range	0.2	
entropy_loss	-16.8	
explained_variance	0.938	
learning_rate	0.0003	
loss	87	
n_updates	5590	
policy_gradient_loss	-0.0544	
std	0.659	
value_loss	239	

rollout/		
ep_len_mean	88.3	
ep_rew_mean	445	
time/		
fps	429	
iterations	561	
time_elapsed	2672	
total_timesteps	1148928	
train/		
approx_kl	0.078916654	
clip_fraction	0.377	
clip_range	0.2	
entropy_loss	-16.8	

	explained_variance		0.931	
	learning_rate		0.0003	
	loss		133	
	n_updates		5600	
	policy_gradient_loss		-0.0571	
	std		0.659	
	value_loss		306	

	rollout/			
	ep_len_mean		90.5	
	ep_rew_mean		458	
	time/			
	fps		429	
	iterations		562	
	time_elapsed		2677	
	total_timesteps		1150976	
	train/			
	approx_kl		0.09281088	
	clip_fraction		0.393	
	clip_range		0.2	
	entropy_loss		-16.7	
	explained_variance		0.926	
	learning_rate		0.0003	
	loss		71.6	
	n_updates		5610	
	policy_gradient_loss		-0.0573	
	std		0.659	
	value_loss		301	

	rollout/			
	ep_len_mean		92.1	
	ep_rew_mean		464	
	time/			
	fps		429	
	iterations		563	
	time_elapsed		2683	
	total_timesteps		1153024	
	train/			
	approx_kl		0.07345067	
	clip_fraction		0.369	
	clip_range		0.2	
	entropy_loss		-16.7	
	explained_variance		0.921	
	learning_rate		0.0003	
	loss		117	
	n_updates		5620	

	policy_gradient_loss		-0.0512	
	std		0.659	
	value_loss		307	

	rollout/			
	ep_len_mean		89.2	
	ep_rew_mean		448	
	time/			
	fps		429	
	iterations		564	
	time_elapsed		2688	
	total_timesteps		1155072	
	train/			
	approx_kl		0.08356528	
	clip_fraction		0.349	
	clip_range		0.2	
	entropy_loss		-16.7	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		102	
	n_updates		5630	
	policy_gradient_loss		-0.0423	
	std		0.658	
	value_loss		345	

	rollout/			
	ep_len_mean		89.2	
	ep_rew_mean		449	
	time/			
	fps		429	
	iterations		565	
	time_elapsed		2693	
	total_timesteps		1157120	
	train/			
	approx_kl		0.093104795	
	clip_fraction		0.374	
	clip_range		0.2	
	entropy_loss		-16.7	
	explained_variance		0.935	
	learning_rate		0.0003	
	loss		110	
	n_updates		5640	
	policy_gradient_loss		-0.0543	
	std		0.659	
	value_loss		305	

---

rollout/		
ep_len_mean	87.8	
ep_rew_mean	440	
time/		
fps	429	
iterations	566	
time_elapsed	2698	
total_timesteps	1159168	
train/		
approx_kl	0.08614381	
clip_fraction	0.363	
clip_range	0.2	
entropy_loss	-16.7	
explained_variance	0.933	
learning_rate	0.0003	
loss	98.6	
n_updates	5650	
policy_gradient_loss	-0.051	
std	0.658	
value_loss	322	

---



---

rollout/		
ep_len_mean	85.7	
ep_rew_mean	429	
time/		
fps	429	
iterations	567	
time_elapsed	2703	
total_timesteps	1161216	
train/		
approx_kl	0.075065345	
clip_fraction	0.372	
clip_range	0.2	
entropy_loss	-16.7	
explained_variance	0.922	
learning_rate	0.0003	
loss	109	
n_updates	5660	
policy_gradient_loss	-0.0583	
std	0.658	
value_loss	310	

---



---

rollout/		
ep_len_mean	87.2	
ep_rew_mean	438	

---

time/		
fps	429	
iterations	568	
time_elapsed	2708	
total_timesteps	1163264	
train/		
approx_kl	0.09792314	
clip_fraction	0.398	
clip_range	0.2	
entropy_loss	-16.7	
explained_variance	0.938	
learning_rate	0.0003	
loss	83.4	
n_updates	5670	
policy_gradient_loss	-0.0518	
std	0.657	
value_loss	239	

rollout/		
ep_len_mean	88.6	
ep_rew_mean	443	
time/		
fps	429	
iterations	569	
time_elapsed	2713	
total_timesteps	1165312	
train/		
approx_kl	0.07265391	
clip_fraction	0.357	
clip_range	0.2	
entropy_loss	-16.7	
explained_variance	0.918	
learning_rate	0.0003	
loss	140	
n_updates	5680	
policy_gradient_loss	-0.049	
std	0.656	
value_loss	352	

rollout/		
ep_len_mean	86.3	
ep_rew_mean	432	
time/		
fps	429	
iterations	570	
time_elapsed	2718	

	total_timesteps		1167360	
	train/			
	approx_kl		0.111679554	
	clip_fraction		0.384	
	clip_range		0.2	
	entropy_loss		-16.7	
	explained_variance		0.932	
	learning_rate		0.0003	
	loss		80.1	
	n_updates		5690	
	policy_gradient_loss		-0.0433	
	std		0.656	
	value_loss		243	

	rollout/			
	ep_len_mean		88.5	
	ep_rew_mean		443	
	time/			
	fps		429	
	iterations		571	
	time_elapsed		2723	
	total_timesteps		1169408	
	train/			
	approx_kl		0.08503549	
	clip_fraction		0.367	
	clip_range		0.2	
	entropy_loss		-16.7	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		96	
	n_updates		5700	
	policy_gradient_loss		-0.0588	
	std		0.656	
	value_loss		289	

	rollout/			
	ep_len_mean		89	
	ep_rew_mean		445	
	time/			
	fps		429	
	iterations		572	
	time_elapsed		2728	
	total_timesteps		1171456	
	train/			
	approx_kl		0.10892215	
	clip_fraction		0.395	



clip_range	0.2	
entropy_loss	-16.7	
explained_variance	0.935	
learning_rate	0.0003	
loss	69.8	
n_updates	5710	
policy_gradient_loss	-0.0531	
std	0.656	
value_loss	257	

rollout/		
ep_len_mean	87.7	
ep_rew_mean	440	
time/		
fps	429	
iterations	573	
time_elapsed	2733	
total_timesteps	1173504	
train/		
approx_kl	0.10355513	
clip_fraction	0.407	
clip_range	0.2	
entropy_loss	-16.7	
explained_variance	0.927	
learning_rate	0.0003	
loss	75.5	
n_updates	5720	
policy_gradient_loss	-0.0547	
std	0.656	
value_loss	265	

rollout/		
ep_len_mean	90.4	
ep_rew_mean	453	
time/		
fps	429	
iterations	574	
time_elapsed	2739	
total_timesteps	1175552	
train/		
approx_kl	0.09331849	
clip_fraction	0.428	
clip_range	0.2	
entropy_loss	-16.7	
explained_variance	0.939	
learning_rate	0.0003	

	loss		102	
	n_updates		5730	
	policy_gradient_loss		-0.0603	
	std		0.656	
	value_loss		254	

	rollout/			
	ep_len_mean		90.5	
	ep_rew_mean		456	
	time/			
	fps		429	
	iterations		575	
	time_elapsed		2744	
	total_timesteps		1177600	
	train/			
	approx_kl		0.09567415	
	clip_fraction		0.375	
	clip_range		0.2	
	entropy_loss		-16.7	
	explained_variance		0.927	
	learning_rate		0.0003	
	loss		76.4	
	n_updates		5740	
	policy_gradient_loss		-0.0516	
	std		0.656	
	value_loss		293	

	rollout/			
	ep_len_mean		90.5	
	ep_rew_mean		456	
	time/			
	fps		428	
	iterations		576	
	time_elapsed		2750	
	total_timesteps		1179648	
	train/			
	approx_kl		0.0956952	
	clip_fraction		0.384	
	clip_range		0.2	
	entropy_loss		-16.6	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		68.4	
	n_updates		5750	
	policy_gradient_loss		-0.0602	
	std		0.656	

	value_loss	274	
-----			
	rollout/		
	ep_len_mean	90.2	
	ep_rew_mean	454	
	time/		
	fps	428	
	iterations	577	
	time_elapsed	2754	
	total_timesteps	1181696	
	train/		
	approx_kl	0.07491121	
	clip_fraction	0.357	
	clip_range	0.2	
	entropy_loss	-16.6	
	explained_variance	0.914	
	learning_rate	0.0003	
	loss	119	
	n_updates	5760	
	policy_gradient_loss	-0.053	
	std	0.655	
	value_loss	311	
-----			
	rollout/		
	ep_len_mean	90.2	
	ep_rew_mean	455	
	time/		
	fps	428	
	iterations	578	
	time_elapsed	2759	
	total_timesteps	1183744	
	train/		
	approx_kl	0.09433505	
	clip_fraction	0.361	
	clip_range	0.2	
	entropy_loss	-16.6	
	explained_variance	0.921	
	learning_rate	0.0003	
	loss	151	
	n_updates	5770	
	policy_gradient_loss	-0.056	
	std	0.654	
	value_loss	350	
-----			
	rollout/		

	ep_len_mean		89.2	
	ep_rew_mean		446	
	time/			
	fps		428	
	iterations		579	
	time_elapsed		2764	
	total_timesteps		1185792	
	train/			
	approx_kl		0.0972748	
	clip_fraction		0.373	
	clip_range		0.2	
	entropy_loss		-16.6	
	explained_variance		0.933	
	learning_rate		0.0003	
	loss		93.5	
	n_updates		5780	
	policy_gradient_loss		-0.0546	
	std		0.654	
	value_loss		301	

	rollout/			
	ep_len_mean		89.1	
	ep_rew_mean		445	
	time/			
	fps		428	
	iterations		580	
	time_elapsed		2769	
	total_timesteps		1187840	
	train/			
	approx_kl		0.10865825	
	clip_fraction		0.393	
	clip_range		0.2	
	entropy_loss		-16.6	
	explained_variance		0.935	
	learning_rate		0.0003	
	loss		77.4	
	n_updates		5790	
	policy_gradient_loss		-0.0475	
	std		0.654	
	value_loss		285	

	rollout/			
	ep_len_mean		90	
	ep_rew_mean		450	
	time/			
	fps		428	

iterations	581
time_elapsed	2774
total_timesteps	1189888
train/	
approx_kl	0.07744074
clip_fraction	0.333
clip_range	0.2
entropy_loss	-16.6
explained_variance	0.916
learning_rate	0.0003
loss	76.8
n_updates	5800
policy_gradient_loss	-0.0526
std	0.654
value_loss	370

rollout/	
ep_len_mean	90
ep_rew_mean	450
time/	
fps	428
iterations	582
time_elapsed	2778
total_timesteps	1191936
train/	
approx_kl	0.085066766
clip_fraction	0.382
clip_range	0.2
entropy_loss	-16.6
explained_variance	0.923
learning_rate	0.0003
loss	114
n_updates	5810
policy_gradient_loss	-0.0525
std	0.654
value_loss	303

rollout/	
ep_len_mean	88
ep_rew_mean	443
time/	
fps	429
iterations	583
time_elapsed	2783
total_timesteps	1193984
train/	

approx_kl	0.08828515	
clip_fraction	0.379	
clip_range	0.2	
entropy_loss	-16.6	
explained_variance	0.936	
learning_rate	0.0003	
loss	132	
n_updates	5820	
policy_gradient_loss	-0.0506	
std	0.654	
value_loss	307	

rollout/		
ep_len_mean	88.5	
ep_rew_mean	446	
time/		
fps	429	
iterations	584	
time_elapsed	2787	
total_timesteps	1196032	
train/		
approx_kl	0.13192126	
clip_fraction	0.41	
clip_range	0.2	
entropy_loss	-16.6	
explained_variance	0.94	
learning_rate	0.0003	
loss	117	
n_updates	5830	
policy_gradient_loss	-0.0475	
std	0.653	
value_loss	262	

rollout/		
ep_len_mean	87.9	
ep_rew_mean	446	
time/		
fps	429	
iterations	585	
time_elapsed	2791	
total_timesteps	1198080	
train/		
approx_kl	0.11705263	
clip_fraction	0.442	
clip_range	0.2	
entropy_loss	-16.6	

	explained_variance		0.935	
	learning_rate		0.0003	
	loss		102	
	n_updates		5840	
	policy_gradient_loss		-0.0554	
	std		0.653	
	value_loss		261	

	rollout/			
	ep_len_mean		87.5	
	ep_rew_mean		442	
	time/			
	fps		429	
	iterations		586	
	time_elapsed		2796	
	total_timesteps		1200128	
	train/			
	approx_kl		0.08364184	
	clip_fraction		0.359	
	clip_range		0.2	
	entropy_loss		-16.6	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		88.5	
	n_updates		5850	
	policy_gradient_loss		-0.0533	
	std		0.652	
	value_loss		325	

	rollout/			
	ep_len_mean		90.2	
	ep_rew_mean		454	
	time/			
	fps		429	
	iterations		587	
	time_elapsed		2800	
	total_timesteps		1202176	
	train/			
	approx_kl		0.074589826	
	clip_fraction		0.365	
	clip_range		0.2	
	entropy_loss		-16.6	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		205	
	n_updates		5860	

	policy_gradient_loss		-0.0502	
	std		0.653	
	value_loss		336	

---

	rollout/			
	ep_len_mean		89.9	
	ep_rew_mean		453	
	time/			
	fps		429	
	iterations		588	
	time_elapsed		2805	
	total_timesteps		1204224	
	train/			
	approx_kl		0.10219	
	clip_fraction		0.381	
	clip_range		0.2	
	entropy_loss		-16.5	
	explained_variance		0.918	
	learning_rate		0.0003	
	loss		124	
	n_updates		5870	
	policy_gradient_loss		-0.0559	
	std		0.652	
	value_loss		301	

---



---

	rollout/			
	ep_len_mean		92.9	
	ep_rew_mean		468	
	time/			
	fps		429	
	iterations		589	
	time_elapsed		2809	
	total_timesteps		1206272	
	train/			
	approx_kl		0.08858838	
	clip_fraction		0.383	
	clip_range		0.2	
	entropy_loss		-16.5	
	explained_variance		0.934	
	learning_rate		0.0003	
	loss		88.9	
	n_updates		5880	
	policy_gradient_loss		-0.0565	
	std		0.652	
	value_loss		289	

---



---

rollout/		
ep_len_mean	90.4	
ep_rew_mean	456	
time/		
fps	429	
iterations	590	
time_elapsed	2813	
total_timesteps	1208320	
train/		
approx_kl	0.08479087	
clip_fraction	0.391	
clip_range	0.2	
entropy_loss	-16.5	
explained_variance	0.933	
learning_rate	0.0003	
loss	102	
n_updates	5890	
policy_gradient_loss	-0.0501	
std	0.652	
value_loss	260	

---



---

rollout/		
ep_len_mean	94.1	
ep_rew_mean	474	
time/		
fps	429	
iterations	591	
time_elapsed	2818	
total_timesteps	1210368	
train/		
approx_kl	0.09907493	
clip_fraction	0.389	
clip_range	0.2	
entropy_loss	-16.5	
explained_variance	0.93	
learning_rate	0.0003	
loss	80.5	
n_updates	5900	
policy_gradient_loss	-0.0533	
std	0.651	
value_loss	277	

---



---

rollout/		
ep_len_mean	93.5	
ep_rew_mean	472	

---

time/		
fps	429	
iterations	592	
time_elapsed	2822	
total_timesteps	1212416	
train/		
approx_kl	0.08779505	
clip_fraction	0.355	
clip_range	0.2	
entropy_loss	-16.5	
explained_variance	0.897	
learning_rate	0.0003	
loss	110	
n_updates	5910	
policy_gradient_loss	-0.0571	
std	0.651	
value_loss	437	

rollout/		
ep_len_mean	91.7	
ep_rew_mean	464	
time/		
fps	429	
iterations	593	
time_elapsed	2827	
total_timesteps	1214464	
train/		
approx_kl	0.116464145	
clip_fraction	0.425	
clip_range	0.2	
entropy_loss	-16.5	
explained_variance	0.94	
learning_rate	0.0003	
loss	74.7	
n_updates	5920	
policy_gradient_loss	-0.0563	
std	0.65	
value_loss	232	

rollout/		
ep_len_mean	90.3	
ep_rew_mean	457	
time/		
fps	429	
iterations	594	
time_elapsed	2832	

	total_timesteps		1216512	
	train/			
	approx_kl		0.08167019	
	clip_fraction		0.355	
	clip_range		0.2	
	entropy_loss		-16.5	
	explained_variance		0.938	
	learning_rate		0.0003	
	loss		81.2	
	n_updates		5930	
	policy_gradient_loss		-0.0536	
	std		0.65	
	value_loss		279	

	rollout/			
	ep_len_mean		90.4	
	ep_rew_mean		458	
	time/			
	fps		429	
	iterations		595	
	time_elapsed		2837	
	total_timesteps		1218560	
	train/			
	approx_kl		0.10230133	
	clip_fraction		0.393	
	clip_range		0.2	
	entropy_loss		-16.5	
	explained_variance		0.928	
	learning_rate		0.0003	
	loss		91.6	
	n_updates		5940	
	policy_gradient_loss		-0.0581	
	std		0.651	
	value_loss		275	

	rollout/			
	ep_len_mean		88.8	
	ep_rew_mean		450	
	time/			
	fps		429	
	iterations		596	
	time_elapsed		2842	
	total_timesteps		1220608	
	train/			
	approx_kl		0.10759689	
	clip_fraction		0.399	

clip_range	0.2	
entropy_loss	-16.5	
explained_variance	0.913	
learning_rate	0.0003	
loss	117	
n_updates	5950	
policy_gradient_loss	-0.0563	
std	0.65	
value_loss	331	

rollout/		
ep_len_mean	90	
ep_rew_mean	454	
time/		
fps	429	
iterations	597	
time_elapsed	2846	
total_timesteps	1222656	
train/		
approx_kl	0.1569413	
clip_fraction	0.428	
clip_range	0.2	
entropy_loss	-16.5	
explained_variance	0.943	
learning_rate	0.0003	
loss	64	
n_updates	5960	
policy_gradient_loss	-0.0531	
std	0.65	
value_loss	230	

rollout/		
ep_len_mean	90.9	
ep_rew_mean	458	
time/		
fps	428	
iterations	598	
time_elapsed	2857	
total_timesteps	1224704	
train/		
approx_kl	0.09458484	
clip_fraction	0.395	
clip_range	0.2	
entropy_loss	-16.5	
explained_variance	0.935	
learning_rate	0.0003	

	loss		99.7	
	n_updates		5970	
	policy_gradient_loss		-0.0566	
	std		0.649	
	value_loss		254	

	rollout/			
	ep_len_mean		92.2	
	ep_rew_mean		464	
	time/			
	fps		424	
	iterations		599	
	time_elapsed		2888	
	total_timesteps		1226752	
	train/			
	approx_kl		0.08691824	
	clip_fraction		0.383	
	clip_range		0.2	
	entropy_loss		-16.4	
	explained_variance		0.933	
	learning_rate		0.0003	
	loss		116	
	n_updates		5980	
	policy_gradient_loss		-0.0579	
	std		0.649	
	value_loss		302	

	rollout/			
	ep_len_mean		94.6	
	ep_rew_mean		476	
	time/			
	fps		423	
	iterations		600	
	time_elapsed		2899	
	total_timesteps		1228800	
	train/			
	approx_kl		0.08563466	
	clip_fraction		0.375	
	clip_range		0.2	
	entropy_loss		-16.4	
	explained_variance		0.947	
	learning_rate		0.0003	
	loss		108	
	n_updates		5990	
	policy_gradient_loss		-0.0598	
	std		0.648	

	value_loss		257	
-----				
	rollout/			
	ep_len_mean		95.4	
	ep_rew_mean		484	
	time/			
	fps		422	
	iterations		601	
	time_elapsed		2913	
	total_timesteps		1230848	
	train/			
	approx_kl		0.10550974	
	clip_fraction		0.393	
	clip_range		0.2	
	entropy_loss		-16.4	
	explained_variance		0.933	
	learning_rate		0.0003	
	loss		83.5	
	n_updates		6000	
	policy_gradient_loss		-0.0537	
	std		0.647	
	value_loss		329	
-----				
	rollout/			
	ep_len_mean		95	
	ep_rew_mean		482	
	time/			
	fps		421	
	iterations		602	
	time_elapsed		2925	
	total_timesteps		1232896	
	train/			
	approx_kl		0.11184597	
	clip_fraction		0.407	
	clip_range		0.2	
	entropy_loss		-16.4	
	explained_variance		0.943	
	learning_rate		0.0003	
	loss		120	
	n_updates		6010	
	policy_gradient_loss		-0.052	
	std		0.647	
	value_loss		246	
-----				
	rollout/			

	ep_len_mean		94.7	
	ep_rew_mean		479	
	time/			
	fps		420	
	iterations		603	
	time_elapsed		2937	
	total_timesteps		1234944	
	train/			
	approx_kl		0.097407304	
	clip_fraction		0.383	
	clip_range		0.2	
	entropy_loss		-16.4	
	explained_variance		0.933	
	learning_rate		0.0003	
	loss		98.2	
	n_updates		6020	
	policy_gradient_loss		-0.0509	
	std		0.647	
	value_loss		339	

	rollout/			
	ep_len_mean		93.7	
	ep_rew_mean		472	
	time/			
	fps		420	
	iterations		604	
	time_elapsed		2944	
	total_timesteps		1236992	
	train/			
	approx_kl		0.11328743	
	clip_fraction		0.366	
	clip_range		0.2	
	entropy_loss		-16.4	
	explained_variance		0.925	
	learning_rate		0.0003	
	loss		89.1	
	n_updates		6030	
	policy_gradient_loss		-0.052	
	std		0.647	
	value_loss		327	

	rollout/			
	ep_len_mean		91.9	
	ep_rew_mean		461	
	time/			
	fps		419	

iterations	605
time_elapsed	2950
total_timesteps	1239040
train/	
approx_kl	0.10307905
clip_fraction	0.394
clip_range	0.2
entropy_loss	-16.4
explained_variance	0.938
learning_rate	0.0003
loss	88
n_updates	6040
policy_gradient_loss	-0.0505
std	0.646
value_loss	285

rollout/	
ep_len_mean	90.5
ep_rew_mean	454
time/	
fps	419
iterations	606
time_elapsed	2955
total_timesteps	1241088
train/	
approx_kl	0.13203554
clip_fraction	0.46
clip_range	0.2
entropy_loss	-16.4
explained_variance	0.941
learning_rate	0.0003
loss	112
n_updates	6050
policy_gradient_loss	-0.0505
std	0.646
value_loss	240

rollout/	
ep_len_mean	90.2
ep_rew_mean	452
time/	
fps	419
iterations	607
time_elapsed	2960
total_timesteps	1243136
train/	



approx_kl	0.10987794	
clip_fraction	0.398	
clip_range	0.2	
entropy_loss	-16.4	
explained_variance	0.922	
learning_rate	0.0003	
loss	124	
n_updates	6060	
policy_gradient_loss	-0.0525	
std	0.645	
value_loss	295	

rollout/		
ep_len_mean	92.9	
ep_rew_mean	467	
time/		
fps	419	
iterations	608	
time_elapsed	2965	
total_timesteps	1245184	
train/		
approx_kl	0.107371464	
clip_fraction	0.437	
clip_range	0.2	
entropy_loss	-16.3	
explained_variance	0.947	
learning_rate	0.0003	
loss	71.9	
n_updates	6070	
policy_gradient_loss	-0.056	
std	0.644	
value_loss	208	

rollout/		
ep_len_mean	93.9	
ep_rew_mean	473	
time/		
fps	419	
iterations	609	
time_elapsed	2970	
total_timesteps	1247232	
train/		
approx_kl	0.11264426	
clip_fraction	0.423	
clip_range	0.2	
entropy_loss	-16.3	

	explained_variance		0.938	
	learning_rate		0.0003	
	loss		83	
	n_updates		6080	
	policy_gradient_loss		-0.0449	
	std		0.644	
	value_loss		261	

	rollout/			
	ep_len_mean		93.7	
	ep_rew_mean		473	
	time/			
	fps		419	
	iterations		610	
	time_elapsed		2977	
	total_timesteps		1249280	
	train/			
	approx_kl		0.08998528	
	clip_fraction		0.368	
	clip_range		0.2	
	entropy_loss		-16.3	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		113	
	n_updates		6090	
	policy_gradient_loss		-0.062	
	std		0.644	
	value_loss		326	

	rollout/			
	ep_len_mean		93.2	
	ep_rew_mean		471	
	time/			
	fps		419	
	iterations		611	
	time_elapsed		2983	
	total_timesteps		1251328	
	train/			
	approx_kl		0.11184265	
	clip_fraction		0.386	
	clip_range		0.2	
	entropy_loss		-16.3	
	explained_variance		0.935	
	learning_rate		0.0003	
	loss		73.2	
	n_updates		6100	

	policy_gradient_loss		-0.0531	
	std		0.644	
	value_loss		278	

---

	rollout/			
	ep_len_mean		92.9	
	ep_rew_mean		470	
	time/			
	fps		419	
	iterations		612	
	time_elapsed		2988	
	total_timesteps		1253376	
	train/			
	approx_kl		0.10652393	
	clip_fraction		0.392	
	clip_range		0.2	
	entropy_loss		-16.3	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		85.5	
	n_updates		6110	
	policy_gradient_loss		-0.0604	
	std		0.644	
	value_loss		292	

---



---

	rollout/			
	ep_len_mean		92.5	
	ep_rew_mean		467	
	time/			
	fps		419	
	iterations		613	
	time_elapsed		2992	
	total_timesteps		1255424	
	train/			
	approx_kl		0.083113626	
	clip_fraction		0.403	
	clip_range		0.2	
	entropy_loss		-16.3	
	explained_variance		0.941	
	learning_rate		0.0003	
	loss		110	
	n_updates		6120	
	policy_gradient_loss		-0.0593	
	std		0.644	
	value_loss		286	

---

---

rollout/		
ep_len_mean	91.3	
ep_rew_mean	460	
time/		
fps	419	
iterations	614	
time_elapsed	2997	
total_timesteps	1257472	
train/		
approx_kl	0.12890622	
clip_fraction	0.411	
clip_range	0.2	
entropy_loss	-16.3	
explained_variance	0.942	
learning_rate	0.0003	
loss	107	
n_updates	6130	
policy_gradient_loss	-0.0536	
std	0.643	
value_loss	246	

---



---

rollout/		
ep_len_mean	92.4	
ep_rew_mean	466	
time/		
fps	419	
iterations	615	
time_elapsed	3003	
total_timesteps	1259520	
train/		
approx_kl	0.09710881	
clip_fraction	0.423	
clip_range	0.2	
entropy_loss	-16.3	
explained_variance	0.934	
learning_rate	0.0003	
loss	85.4	
n_updates	6140	
policy_gradient_loss	-0.0536	
std	0.643	
value_loss	267	

---



---

rollout/		
ep_len_mean	90.7	
ep_rew_mean	456	

---

time/		
fps	418	
iterations	616	
time_elapsed	3012	
total_timesteps	1261568	
train/		
approx_kl	0.08372107	
clip_fraction	0.335	
clip_range	0.2	
entropy_loss	-16.3	
explained_variance	0.915	
learning_rate	0.0003	
loss	140	
n_updates	6150	
policy_gradient_loss	-0.0517	
std	0.642	
value_loss	417	

rollout/		
ep_len_mean	90.6	
ep_rew_mean	456	
time/		
fps	418	
iterations	617	
time_elapsed	3020	
total_timesteps	1263616	
train/		
approx_kl	0.08831097	
clip_fraction	0.376	
clip_range	0.2	
entropy_loss	-16.2	
explained_variance	0.935	
learning_rate	0.0003	
loss	88	
n_updates	6160	
policy_gradient_loss	-0.0553	
std	0.642	
value_loss	301	

rollout/		
ep_len_mean	89	
ep_rew_mean	449	
time/		
fps	418	
iterations	618	
time_elapsed	3025	

	total_timesteps		1265664	
	train/			
	approx_kl		0.08354174	
	clip_fraction		0.371	
	clip_range		0.2	
	entropy_loss		-16.2	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		132	
	n_updates		6170	
	policy_gradient_loss		-0.0597	
	std		0.641	
	value_loss		326	

	rollout/			
	ep_len_mean		91	
	ep_rew_mean		457	
	time/			
	fps		418	
	iterations		619	
	time_elapsed		3031	
	total_timesteps		1267712	
	train/			
	approx_kl		0.09194334	
	clip_fraction		0.378	
	clip_range		0.2	
	entropy_loss		-16.2	
	explained_variance		0.939	
	learning_rate		0.0003	
	loss		88	
	n_updates		6180	
	policy_gradient_loss		-0.0565	
	std		0.64	
	value_loss		278	

	rollout/			
	ep_len_mean		92.1	
	ep_rew_mean		462	
	time/			
	fps		418	
	iterations		620	
	time_elapsed		3037	
	total_timesteps		1269760	
	train/			
	approx_kl		0.09142141	
	clip_fraction		0.387	

clip_range	0.2	
entropy_loss	-16.2	
explained_variance	0.922	
learning_rate	0.0003	
loss	141	
n_updates	6190	
policy_gradient_loss	-0.0516	
std	0.64	
value_loss	303	

rollout/		
ep_len_mean	92.2	
ep_rew_mean	463	
time/		
fps	418	
iterations	621	
time_elapsed	3041	
total_timesteps	1271808	
train/		
approx_kl	0.11475712	
clip_fraction	0.405	
clip_range	0.2	
entropy_loss	-16.2	
explained_variance	0.917	
learning_rate	0.0003	
loss	75.9	
n_updates	6200	
policy_gradient_loss	-0.0526	
std	0.64	
value_loss	257	

rollout/		
ep_len_mean	92.8	
ep_rew_mean	465	
time/		
fps	418	
iterations	622	
time_elapsed	3046	
total_timesteps	1273856	
train/		
approx_kl	0.116162986	
clip_fraction	0.39	
clip_range	0.2	
entropy_loss	-16.2	
explained_variance	0.913	
learning_rate	0.0003	

	loss		148	
	n_updates		6210	
	policy_gradient_loss		-0.052	
	std		0.639	
	value_loss		355	

---

	rollout/			
	ep_len_mean		91.5	
	ep_rew_mean		459	
	time/			
	fps		418	
	iterations		623	
	time_elapsed		3050	
	total_timesteps		1275904	
	train/			
	approx_kl		0.11897094	
	clip_fraction		0.412	
	clip_range		0.2	
	entropy_loss		-16.2	
	explained_variance		0.938	
	learning_rate		0.0003	
	loss		74.3	
	n_updates		6220	
	policy_gradient_loss		-0.0418	
	std		0.639	
	value_loss		249	

---

	rollout/			
	ep_len_mean		89.9	
	ep_rew_mean		453	
	time/			
	fps		418	
	iterations		624	
	time_elapsed		3056	
	total_timesteps		1277952	
	train/			
	approx_kl		0.10101673	
	clip_fraction		0.392	
	clip_range		0.2	
	entropy_loss		-16.2	
	explained_variance		0.941	
	learning_rate		0.0003	
	loss		75	
	n_updates		6230	
	policy_gradient_loss		-0.0532	
	std		0.639	



value_loss	252
-----	
rollout/	
ep_len_mean	87.8
ep_rew_mean	443
time/	
fps	418
iterations	625
time_elapsed	3060
total_timesteps	1280000
train/	
approx_kl	0.14648922
clip_fraction	0.495
clip_range	0.2
entropy_loss	-16.2
explained_variance	0.947
learning_rate	0.0003
loss	86.4
n_updates	6240
policy_gradient_loss	-0.0484
std	0.64
value_loss	187
-----	
rollout/	
ep_len_mean	89.7
ep_rew_mean	452
time/	
fps	418
iterations	626
time_elapsed	3064
total_timesteps	1282048
train/	
approx_kl	0.117746755
clip_fraction	0.418
clip_range	0.2
entropy_loss	-16.2
explained_variance	0.927
learning_rate	0.0003
loss	150
n_updates	6250
policy_gradient_loss	-0.0524
std	0.639
value_loss	291
-----	
rollout/	

	ep_len_mean	93.2	
	ep_rew_mean	468	
	time/		
	fps	418	
	iterations	627	
	time_elapsed	3069	
	total_timesteps	1284096	
	train/		
	approx_kl	0.15652288	
	clip_fraction	0.442	
	clip_range	0.2	
	entropy_loss	-16.1	
	explained_variance	0.935	
	learning_rate	0.0003	
	loss	115	
	n_updates	6260	
	policy_gradient_loss	-0.0522	
	std	0.638	
	value_loss	250	

	rollout/		
	ep_len_mean	95.6	
	ep_rew_mean	479	
	time/		
	fps	418	
	iterations	628	
	time_elapsed	3073	
	total_timesteps	1286144	
	train/		
	approx_kl	0.09987262	
	clip_fraction	0.439	
	clip_range	0.2	
	entropy_loss	-16.1	
	explained_variance	0.916	
	learning_rate	0.0003	
	loss	126	
	n_updates	6270	
	policy_gradient_loss	-0.0505	
	std	0.638	
	value_loss	367	

	rollout/		
	ep_len_mean	97.3	
	ep_rew_mean	488	
	time/		
	fps	418	

iterations	629
time_elapsed	3078
total_timesteps	1288192
train/	
approx_kl	0.106919505
clip_fraction	0.371
clip_range	0.2
entropy_loss	-16.1
explained_variance	0.915
learning_rate	0.0003
loss	92.9
n_updates	6280
policy_gradient_loss	-0.0532
std	0.638
value_loss	343

rollout/	
ep_len_mean	99.9
ep_rew_mean	500
time/	
fps	418
iterations	630
time_elapsed	3082
total_timesteps	1290240
train/	
approx_kl	0.12448252
clip_fraction	0.394
clip_range	0.2
entropy_loss	-16.1
explained_variance	0.929
learning_rate	0.0003
loss	64.4
n_updates	6290
policy_gradient_loss	-0.047
std	0.637
value_loss	306

rollout/	
ep_len_mean	97.5
ep_rew_mean	488
time/	
fps	418
iterations	631
time_elapsed	3087
total_timesteps	1292288
train/	

approx_kl	0.08473733	
clip_fraction	0.356	
clip_range	0.2	
entropy_loss	-16.1	
explained_variance	0.922	
learning_rate	0.0003	
loss	126	
n_updates	6300	
policy_gradient_loss	-0.0558	
std	0.637	
value_loss	354	

rollout/		
ep_len_mean	95.2	
ep_rew_mean	477	
time/		
fps	418	
iterations	632	
time_elapsed	3091	
total_timesteps	1294336	
train/		
approx_kl	0.075785734	
clip_fraction	0.373	
clip_range	0.2	
entropy_loss	-16.1	
explained_variance	0.945	
learning_rate	0.0003	
loss	107	
n_updates	6310	
policy_gradient_loss	-0.0569	
std	0.636	
value_loss	293	

rollout/		
ep_len_mean	95	
ep_rew_mean	479	
time/		
fps	418	
iterations	633	
time_elapsed	3096	
total_timesteps	1296384	
train/		
approx_kl	0.09912253	
clip_fraction	0.401	
clip_range	0.2	
entropy_loss	-16.1	

	explained_variance		0.927	
	learning_rate		0.0003	
	loss		165	
	n_updates		6320	
	policy_gradient_loss		-0.057	
	std		0.636	
	value_loss		313	

	rollout/			
	ep_len_mean		94.6	
	ep_rew_mean		475	
	time/			
	fps		418	
	iterations		634	
	time_elapsed		3103	
	total_timesteps		1298432	
	train/			
	approx_kl		0.08801077	
	clip_fraction		0.388	
	clip_range		0.2	
	entropy_loss		-16.1	
	explained_variance		0.939	
	learning_rate		0.0003	
	loss		85.8	
	n_updates		6330	
	policy_gradient_loss		-0.0523	
	std		0.635	
	value_loss		273	

	rollout/			
	ep_len_mean		93.7	
	ep_rew_mean		471	
	time/			
	fps		418	
	iterations		635	
	time_elapsed		3109	
	total_timesteps		1300480	
	train/			
	approx_kl		0.13671522	
	clip_fraction		0.408	
	clip_range		0.2	
	entropy_loss		-16	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		83.3	
	n_updates		6340	

	policy_gradient_loss		-0.0483	
	std		0.635	
	value_loss		295	

---

	rollout/			
	ep_len_mean		92.8	
	ep_rew_mean		469	
	time/			
	fps		418	
	iterations		636	
	time_elapsed		3114	
	total_timesteps		1302528	
	train/			
	approx_kl		0.09269383	
	clip_fraction		0.391	
	clip_range		0.2	
	entropy_loss		-16	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		81.2	
	n_updates		6350	
	policy_gradient_loss		-0.0524	
	std		0.635	
	value_loss		309	

---

	rollout/			
	ep_len_mean		95.3	
	ep_rew_mean		483	
	time/			
	fps		418	
	iterations		637	
	time_elapsed		3118	
	total_timesteps		1304576	
	train/			
	approx_kl		0.09995041	
	clip_fraction		0.368	
	clip_range		0.2	
	entropy_loss		-16	
	explained_variance		0.934	
	learning_rate		0.0003	
	loss		81.8	
	n_updates		6360	
	policy_gradient_loss		-0.0514	
	std		0.635	
	value_loss		336	

---

-----		
rollout/		
ep_len_mean	95.6	
ep_rew_mean	486	
time/		
fps	418	
iterations	638	
time_elapsed	3122	
total_timesteps	1306624	
train/		
approx_kl	0.112704135	
clip_fraction	0.413	
clip_range	0.2	
entropy_loss	-16.1	
explained_variance	0.941	
learning_rate	0.0003	
loss	89.1	
n_updates	6370	
policy_gradient_loss	-0.0502	
std	0.635	
value_loss	282	
-----		

-----		
rollout/		
ep_len_mean	96.4	
ep_rew_mean	491	
time/		
fps	418	
iterations	639	
time_elapsed	3127	
total_timesteps	1308672	
train/		
approx_kl	0.10916017	
clip_fraction	0.41	
clip_range	0.2	
entropy_loss	-16.1	
explained_variance	0.948	
learning_rate	0.0003	
loss	73.7	
n_updates	6380	
policy_gradient_loss	-0.0559	
std	0.635	
value_loss	256	
-----		

-----		
rollout/		
ep_len_mean	98	
ep_rew_mean	499	

time/		
fps	418	
iterations	640	
time_elapsed	3131	
total_timesteps	1310720	
train/		
approx_kl	0.07799986	
clip_fraction	0.366	
clip_range	0.2	
entropy_loss	-16	
explained_variance	0.938	
learning_rate	0.0003	
loss	100	
n_updates	6390	
policy_gradient_loss	-0.0556	
std	0.635	
value_loss	289	

---

rollout/		
ep_len_mean	97.9	
ep_rew_mean	501	
time/		
fps	418	
iterations	641	
time_elapsed	3136	
total_timesteps	1312768	
train/		
approx_kl	0.088360265	
clip_fraction	0.374	
clip_range	0.2	
entropy_loss	-16	
explained_variance	0.933	
learning_rate	0.0003	
loss	83.7	
n_updates	6400	
policy_gradient_loss	-0.0539	
std	0.635	
value_loss	334	

---

rollout/		
ep_len_mean	97.5	
ep_rew_mean	495	
time/		
fps	418	
iterations	642	
time_elapsed	3141	



	total_timesteps		1314816	
	train/			
	approx_kl		0.07195428	
	clip_fraction		0.365	
	clip_range		0.2	
	entropy_loss		-16	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		111	
	n_updates		6410	
	policy_gradient_loss		-0.0543	
	std		0.634	
	value_loss		309	

	rollout/			
	ep_len_mean		95.5	
	ep_rew_mean		488	
	time/			
	fps		418	
	iterations		643	
	time_elapsed		3145	
	total_timesteps		1316864	
	train/			
	approx_kl		0.12340614	
	clip_fraction		0.387	
	clip_range		0.2	
	entropy_loss		-16	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		119	
	n_updates		6420	
	policy_gradient_loss		-0.0536	
	std		0.633	
	value_loss		358	

	rollout/			
	ep_len_mean		94.2	
	ep_rew_mean		480	
	time/			
	fps		418	
	iterations		644	
	time_elapsed		3150	
	total_timesteps		1318912	
	train/			
	approx_kl		0.13205907	
	clip_fraction		0.421	

clip_range	0.2	
entropy_loss	-16	
explained_variance	0.931	
learning_rate	0.0003	
loss	73.8	
n_updates	6430	
policy_gradient_loss	-0.0509	
std	0.633	
value_loss	279	

rollout/		
ep_len_mean	94.8	
ep_rew_mean	482	
time/		
fps	418	
iterations	645	
time_elapsed	3154	
total_timesteps	1320960	
train/		
approx_kl	0.10468187	
clip_fraction	0.399	
clip_range	0.2	
entropy_loss	-16	
explained_variance	0.94	
learning_rate	0.0003	
loss	125	
n_updates	6440	
policy_gradient_loss	-0.0538	
std	0.632	
value_loss	283	

rollout/		
ep_len_mean	92.5	
ep_rew_mean	470	
time/		
fps	418	
iterations	646	
time_elapsed	3159	
total_timesteps	1323008	
train/		
approx_kl	0.100728095	
clip_fraction	0.4	
clip_range	0.2	
entropy_loss	-15.9	
explained_variance	0.945	
learning_rate	0.0003	

	loss		71.6	
	n_updates		6450	
	policy_gradient_loss		-0.0454	
	std		0.632	
	value_loss		261	

	rollout/			
	ep_len_mean		90.8	
	ep_rew_mean		463	
	time/			
	fps		418	
	iterations		647	
	time_elapsed		3163	
	total_timesteps		1325056	
	train/			
	approx_kl		0.119350694	
	clip_fraction		0.405	
	clip_range		0.2	
	entropy_loss		-15.9	
	explained_variance		0.938	
	learning_rate		0.0003	
	loss		114	
	n_updates		6460	
	policy_gradient_loss		-0.0477	
	std		0.631	
	value_loss		292	

	rollout/			
	ep_len_mean		91.6	
	ep_rew_mean		468	
	time/			
	fps		418	
	iterations		648	
	time_elapsed		3167	
	total_timesteps		1327104	
	train/			
	approx_kl		0.12210329	
	clip_fraction		0.415	
	clip_range		0.2	
	entropy_loss		-15.9	
	explained_variance		0.923	
	learning_rate		0.0003	
	loss		82.8	
	n_updates		6470	
	policy_gradient_loss		-0.0546	
	std		0.631	

	value_loss		317	
-----				
	rollout/			
	ep_len_mean		91.2	
	ep_rew_mean		467	
	time/			
	fps		419	
	iterations		649	
	time_elapsed		3172	
	total_timesteps		1329152	
	train/			
	approx_kl		0.1418406	
	clip_fraction		0.43	
	clip_range		0.2	
	entropy_loss		-15.9	
	explained_variance		0.945	
	learning_rate		0.0003	
	loss		121	
	n_updates		6480	
	policy_gradient_loss		-0.0532	
	std		0.63	
	value_loss		252	
-----				
	rollout/			
	ep_len_mean		90.3	
	ep_rew_mean		463	
	time/			
	fps		419	
	iterations		650	
	time_elapsed		3176	
	total_timesteps		1331200	
	train/			
	approx_kl		0.10296649	
	clip_fraction		0.409	
	clip_range		0.2	
	entropy_loss		-15.9	
	explained_variance		0.934	
	learning_rate		0.0003	
	loss		77.8	
	n_updates		6490	
	policy_gradient_loss		-0.0518	
	std		0.63	
	value_loss		299	
-----				
	rollout/			

	ep_len_mean		91.7	
	ep_rew_mean		468	
	time/			
	fps		419	
	iterations		651	
	time_elapsed		3180	
	total_timesteps		1333248	
	train/			
	approx_kl		0.098235056	
	clip_fraction		0.406	
	clip_range		0.2	
	entropy_loss		-15.9	
	explained_variance		0.943	
	learning_rate		0.0003	
	loss		135	
	n_updates		6500	
	policy_gradient_loss		-0.0555	
	std		0.63	
	value_loss		289	

	rollout/			
	ep_len_mean		91.6	
	ep_rew_mean		468	
	time/			
	fps		419	
	iterations		652	
	time_elapsed		3185	
	total_timesteps		1335296	
	train/			
	approx_kl		0.10583433	
	clip_fraction		0.381	
	clip_range		0.2	
	entropy_loss		-15.9	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		91.6	
	n_updates		6510	
	policy_gradient_loss		-0.0514	
	std		0.63	
	value_loss		315	

	rollout/			
	ep_len_mean		89.9	
	ep_rew_mean		458	
	time/			
	fps		419	

iterations	653
time_elapsed	3189
total_timesteps	1337344
train/	
approx_kl	0.080802396
clip_fraction	0.394
clip_range	0.2
entropy_loss	-15.9
explained_variance	0.939
learning_rate	0.0003
loss	110
n_updates	6520
policy_gradient_loss	-0.0552
std	0.63
value_loss	268

rollout/	
ep_len_mean	91.9
ep_rew_mean	465
time/	
fps	419
iterations	654
time_elapsed	3193
total_timesteps	1339392
train/	
approx_kl	0.12993746
clip_fraction	0.447
clip_range	0.2
entropy_loss	-15.9
explained_variance	0.949
learning_rate	0.0003
loss	71.4
n_updates	6530
policy_gradient_loss	-0.0537
std	0.629
value_loss	219

rollout/	
ep_len_mean	92.2
ep_rew_mean	467
time/	
fps	419
iterations	655
time_elapsed	3198
total_timesteps	1341440
train/	

approx_kl	0.09892955	
clip_fraction	0.402	
clip_range	0.2	
entropy_loss	-15.9	
explained_variance	0.921	
learning_rate	0.0003	
loss	73.1	
n_updates	6540	
policy_gradient_loss	-0.0567	
std	0.629	
value_loss	315	

rollout/		
ep_len_mean	92.6	
ep_rew_mean	469	
time/		
fps	419	
iterations	656	
time_elapsed	3202	
total_timesteps	1343488	
train/		
approx_kl	0.103409246	
clip_fraction	0.409	
clip_range	0.2	
entropy_loss	-15.9	
explained_variance	0.942	
learning_rate	0.0003	
loss	72.2	
n_updates	6550	
policy_gradient_loss	-0.0501	
std	0.628	
value_loss	259	

rollout/		
ep_len_mean	95.9	
ep_rew_mean	482	
time/		
fps	419	
iterations	657	
time_elapsed	3207	
total_timesteps	1345536	
train/		
approx_kl	0.13971555	
clip_fraction	0.408	
clip_range	0.2	
entropy_loss	-15.8	

	explained_variance		0.944	
	learning_rate		0.0003	
	loss		72.2	
	n_updates		6560	
	policy_gradient_loss		-0.0508	
	std		0.627	
	value_loss		276	

	rollout/			
	ep_len_mean		96.7	
	ep_rew_mean		487	
	time/			
	fps		419	
	iterations		658	
	time_elapsed		3211	
	total_timesteps		1347584	
	train/			
	approx_kl		0.16371807	
	clip_fraction		0.45	
	clip_range		0.2	
	entropy_loss		-15.8	
	explained_variance		0.934	
	learning_rate		0.0003	
	loss		132	
	n_updates		6570	
	policy_gradient_loss		-0.0499	
	std		0.627	
	value_loss		251	

	rollout/			
	ep_len_mean		94.3	
	ep_rew_mean		473	
	time/			
	fps		419	
	iterations		659	
	time_elapsed		3215	
	total_timesteps		1349632	
	train/			
	approx_kl		0.12069722	
	clip_fraction		0.415	
	clip_range		0.2	
	entropy_loss		-15.8	
	explained_variance		0.942	
	learning_rate		0.0003	
	loss		93	
	n_updates		6580	



	policy_gradient_loss		-0.0566	
	std		0.626	
	value_loss		250	

	rollout/			
	ep_len_mean		93.4	
	ep_rew_mean		468	
	time/			
	fps		419	
	iterations		660	
	time_elapsed		3220	
	total_timesteps		1351680	
	train/			
	approx_kl		0.120775476	
	clip_fraction		0.429	
	clip_range		0.2	
	entropy_loss		-15.8	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		150	
	n_updates		6590	
	policy_gradient_loss		-0.052	
	std		0.625	
	value_loss		282	

	rollout/			
	ep_len_mean		92	
	ep_rew_mean		461	
	time/			
	fps		419	
	iterations		661	
	time_elapsed		3225	
	total_timesteps		1353728	
	train/			
	approx_kl		0.15104917	
	clip_fraction		0.414	
	clip_range		0.2	
	entropy_loss		-15.8	
	explained_variance		0.939	
	learning_rate		0.0003	
	loss		88.9	
	n_updates		6600	
	policy_gradient_loss		-0.0506	
	std		0.626	
	value_loss		278	

---

rollout/		
ep_len_mean	92.8	
ep_rew_mean	467	
time/		
fps	419	
iterations	662	
time_elapsed	3230	
total_timesteps	1355776	
train/		
approx_kl	0.09460601	
clip_fraction	0.367	
clip_range	0.2	
entropy_loss	-15.8	
explained_variance	0.915	
learning_rate	0.0003	
loss	140	
n_updates	6610	
policy_gradient_loss	-0.0496	
std	0.626	
value_loss	349	

---



---

rollout/		
ep_len_mean	95.7	
ep_rew_mean	481	
time/		
fps	419	
iterations	663	
time_elapsed	3235	
total_timesteps	1357824	
train/		
approx_kl	0.11514655	
clip_fraction	0.429	
clip_range	0.2	
entropy_loss	-15.8	
explained_variance	0.945	
learning_rate	0.0003	
loss	127	
n_updates	6620	
policy_gradient_loss	-0.044	
std	0.626	
value_loss	243	

---



---

rollout/		
ep_len_mean	93.3	
ep_rew_mean	471	

---

time/		
fps	419	
iterations	664	
time_elapsed	3239	
total_timesteps	1359872	
train/		
approx_kl	0.09531881	
clip_fraction	0.39	
clip_range	0.2	
entropy_loss	-15.8	
explained_variance	0.92	
learning_rate	0.0003	
loss	170	
n_updates	6630	
policy_gradient_loss	-0.047	
std	0.626	
value_loss	329	

---

rollout/		
ep_len_mean	94	
ep_rew_mean	476	
time/		
fps	419	
iterations	665	
time_elapsed	3243	
total_timesteps	1361920	
train/		
approx_kl	0.10867367	
clip_fraction	0.395	
clip_range	0.2	
entropy_loss	-15.8	
explained_variance	0.94	
learning_rate	0.0003	
loss	83.9	
n_updates	6640	
policy_gradient_loss	-0.0482	
std	0.626	
value_loss	252	

---

rollout/		
ep_len_mean	93.3	
ep_rew_mean	472	
time/		
fps	419	
iterations	666	
time_elapsed	3248	

	total_timesteps		1363968	
	train/			
	approx_kl		0.08637962	
	clip_fraction		0.38	
	clip_range		0.2	
	entropy_loss		-15.8	
	explained_variance		0.934	
	learning_rate		0.0003	
	loss		117	
	n_updates		6650	
	policy_gradient_loss		-0.0493	
	std		0.626	
	value_loss		307	

	rollout/			
	ep_len_mean		91.2	
	ep_rew_mean		464	
	time/			
	fps		420	
	iterations		667	
	time_elapsed		3252	
	total_timesteps		1366016	
	train/			
	approx_kl		0.14032133	
	clip_fraction		0.381	
	clip_range		0.2	
	entropy_loss		-15.8	
	explained_variance		0.934	
	learning_rate		0.0003	
	loss		108	
	n_updates		6660	
	policy_gradient_loss		-0.0483	
	std		0.626	
	value_loss		302	

	rollout/			
	ep_len_mean		91.6	
	ep_rew_mean		464	
	time/			
	fps		420	
	iterations		668	
	time_elapsed		3256	
	total_timesteps		1368064	
	train/			
	approx_kl		0.08928265	
	clip_fraction		0.372	

clip_range	0.2	
entropy_loss	-15.8	
explained_variance	0.929	
learning_rate	0.0003	
loss	110	
n_updates	6670	
policy_gradient_loss	-0.055	
std	0.626	
value_loss	298	

rollout/		
ep_len_mean	93.3	
ep_rew_mean	472	
time/		
fps	420	
iterations	669	
time_elapsed	3260	
total_timesteps	1370112	
train/		
approx_kl	0.10979276	
clip_fraction	0.385	
clip_range	0.2	
entropy_loss	-15.8	
explained_variance	0.918	
learning_rate	0.0003	
loss	102	
n_updates	6680	
policy_gradient_loss	-0.0574	
std	0.625	
value_loss	317	

rollout/		
ep_len_mean	91.7	
ep_rew_mean	465	
time/		
fps	420	
iterations	670	
time_elapsed	3264	
total_timesteps	1372160	
train/		
approx_kl	0.13945822	
clip_fraction	0.458	
clip_range	0.2	
entropy_loss	-15.7	
explained_variance	0.935	
learning_rate	0.0003	

	loss		70.8	
	n_updates		6690	
	policy_gradient_loss		-0.0503	
	std		0.624	
	value_loss		273	

	rollout/			
	ep_len_mean		92.8	
	ep_rew_mean		470	
	time/			
	fps		420	
	iterations		671	
	time_elapsed		3268	
	total_timesteps		1374208	
	train/			
	approx_kl		0.11450811	
	clip_fraction		0.428	
	clip_range		0.2	
	entropy_loss		-15.7	
	explained_variance		0.947	
	learning_rate		0.0003	
	loss		46.8	
	n_updates		6700	
	policy_gradient_loss		-0.0561	
	std		0.623	
	value_loss		235	

	rollout/			
	ep_len_mean		94.4	
	ep_rew_mean		479	
	time/			
	fps		420	
	iterations		672	
	time_elapsed		3272	
	total_timesteps		1376256	
	train/			
	approx_kl		0.13721976	
	clip_fraction		0.441	
	clip_range		0.2	
	entropy_loss		-15.7	
	explained_variance		0.94	
	learning_rate		0.0003	
	loss		136	
	n_updates		6710	
	policy_gradient_loss		-0.0495	
	std		0.623	

	value_loss		262	
-----				
	rollout/			
	ep_len_mean		95.3	
	ep_rew_mean		485	
	time/			
	fps		420	
	iterations		673	
	time_elapsed		3277	
	total_timesteps		1378304	
	train/			
	approx_kl		0.108025916	
	clip_fraction		0.41	
	clip_range		0.2	
	entropy_loss		-15.7	
	explained_variance		0.936	
	learning_rate		0.0003	
	loss		98.9	
	n_updates		6720	
	policy_gradient_loss		-0.0535	
	std		0.622	
	value_loss		307	
-----				
	rollout/			
	ep_len_mean		96.9	
	ep_rew_mean		495	
	time/			
	fps		420	
	iterations		674	
	time_elapsed		3281	
	total_timesteps		1380352	
	train/			
	approx_kl		0.13471457	
	clip_fraction		0.441	
	clip_range		0.2	
	entropy_loss		-15.7	
	explained_variance		0.937	
	learning_rate		0.0003	
	loss		114	
	n_updates		6730	
	policy_gradient_loss		-0.05	
	std		0.622	
	value_loss		305	
-----				
	rollout/			

	ep_len_mean		97.2	
	ep_rew_mean		494	
	time/			
	fps		420	
	iterations		675	
	time_elapsed		3285	
	total_timesteps		1382400	
	train/			
	approx_kl		0.09807783	
	clip_fraction		0.373	
	clip_range		0.2	
	entropy_loss		-15.7	
	explained_variance		0.915	
	learning_rate		0.0003	
	loss		122	
	n_updates		6740	
	policy_gradient_loss		-0.0522	
	std		0.622	
	value_loss		367	

	rollout/			
	ep_len_mean		94.8	
	ep_rew_mean		480	
	time/			
	fps		420	
	iterations		676	
	time_elapsed		3289	
	total_timesteps		1384448	
	train/			
	approx_kl		0.15179682	
	clip_fraction		0.441	
	clip_range		0.2	
	entropy_loss		-15.6	
	explained_variance		0.917	
	learning_rate		0.0003	
	loss		142	
	n_updates		6750	
	policy_gradient_loss		-0.0446	
	std		0.62	
	value_loss		290	

	rollout/			
	ep_len_mean		93.1	
	ep_rew_mean		470	
	time/			
	fps		420	



	iterations		677	
	time_elapsed		3294	
	total_timesteps		1386496	
	train/			
	approx_kl		0.09062878	
	clip_fraction		0.372	
	clip_range		0.2	
	entropy_loss		-15.6	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		106	
	n_updates		6760	
	policy_gradient_loss		-0.0457	
	std		0.621	
	value_loss		276	

	rollout/			
	ep_len_mean		94.5	
	ep_rew_mean		475	
	time/			
	fps		421	
	iterations		678	
	time_elapsed		3298	
	total_timesteps		1388544	
	train/			
	approx_kl		0.17539057	
	clip_fraction		0.472	
	clip_range		0.2	
	entropy_loss		-15.6	
	explained_variance		0.938	
	learning_rate		0.0003	
	loss		79.9	
	n_updates		6770	
	policy_gradient_loss		-0.0517	
	std		0.62	
	value_loss		236	

	rollout/			
	ep_len_mean		95.2	
	ep_rew_mean		478	
	time/			
	fps		421	
	iterations		679	
	time_elapsed		3302	
	total_timesteps		1390592	
	train/			

approx_kl	0.12289764	
clip_fraction	0.435	
clip_range	0.2	
entropy_loss	-15.6	
explained_variance	0.94	
learning_rate	0.0003	
loss	88.2	
n_updates	6780	
policy_gradient_loss	-0.0543	
std	0.62	
value_loss	249	

rollout/		
ep_len_mean	96.7	
ep_rew_mean	485	
time/		
fps	421	
iterations	680	
time_elapsed	3306	
total_timesteps	1392640	
train/		
approx_kl	0.10593432	
clip_fraction	0.398	
clip_range	0.2	
entropy_loss	-15.6	
explained_variance	0.94	
learning_rate	0.0003	
loss	108	
n_updates	6790	
policy_gradient_loss	-0.0535	
std	0.621	
value_loss	294	

rollout/		
ep_len_mean	97.2	
ep_rew_mean	489	
time/		
fps	421	
iterations	681	
time_elapsed	3310	
total_timesteps	1394688	
train/		
approx_kl	0.15259296	
clip_fraction	0.374	
clip_range	0.2	
entropy_loss	-15.6	

	explained_variance	0.928	
	learning_rate	0.0003	
	loss	120	
	n_updates	6800	
	policy_gradient_loss	-0.0524	
	std	0.62	
	value_loss	371	

	rollout/		
	ep_len_mean	97.2	
	ep_rew_mean	489	
	time/		
	fps	421	
	iterations	682	
	time_elapsed	3315	
	total_timesteps	1396736	
	train/		
	approx_kl	0.11089732	
	clip_fraction	0.413	
	clip_range	0.2	
	entropy_loss	-15.6	
	explained_variance	0.945	
	learning_rate	0.0003	
	loss	107	
	n_updates	6810	
	policy_gradient_loss	-0.0585	
	std	0.62	
	value_loss	271	

	rollout/		
	ep_len_mean	95.3	
	ep_rew_mean	480	
	time/		
	fps	421	
	iterations	683	
	time_elapsed	3320	
	total_timesteps	1398784	
	train/		
	approx_kl	0.112854086	
	clip_fraction	0.403	
	clip_range	0.2	
	entropy_loss	-15.6	
	explained_variance	0.941	
	learning_rate	0.0003	
	loss	94.1	
	n_updates	6820	

	policy_gradient_loss		-0.0488	
	std		0.62	
	value_loss		279	

---

	rollout/			
	ep_len_mean		94.7	
	ep_rew_mean		479	
	time/			
	fps		421	
	iterations		684	
	time_elapsed		3324	
	total_timesteps		1400832	
	train/			
	approx_kl		0.09281428	
	clip_fraction		0.399	
	clip_range		0.2	
	entropy_loss		-15.6	
	explained_variance		0.936	
	learning_rate		0.0003	
	loss		95.2	
	n_updates		6830	
	policy_gradient_loss		-0.0517	
	std		0.62	
	value_loss		289	

---



---

	rollout/			
	ep_len_mean		92	
	ep_rew_mean		466	
	time/			
	fps		421	
	iterations		685	
	time_elapsed		3329	
	total_timesteps		1402880	
	train/			
	approx_kl		0.093377456	
	clip_fraction		0.379	
	clip_range		0.2	
	entropy_loss		-15.6	
	explained_variance		0.925	
	learning_rate		0.0003	
	loss		104	
	n_updates		6840	
	policy_gradient_loss		-0.0555	
	std		0.62	
	value_loss		314	

---

---

rollout/		
ep_len_mean	94.1	
ep_rew_mean	479	
time/		
fps	421	
iterations	686	
time_elapsed	3333	
total_timesteps	1404928	
train/		
approx_kl	0.10561605	
clip_fraction	0.424	
clip_range	0.2	
entropy_loss	-15.6	
explained_variance	0.939	
learning_rate	0.0003	
loss	69.9	
n_updates	6850	
policy_gradient_loss	-0.0533	
std	0.619	
value_loss	272	

---



---

rollout/		
ep_len_mean	96.4	
ep_rew_mean	491	
time/		
fps	421	
iterations	687	
time_elapsed	3338	
total_timesteps	1406976	
train/		
approx_kl	0.11148706	
clip_fraction	0.387	
clip_range	0.2	
entropy_loss	-15.6	
explained_variance	0.915	
learning_rate	0.0003	
loss	140	
n_updates	6860	
policy_gradient_loss	-0.051	
std	0.618	
value_loss	423	

---



---

rollout/		
ep_len_mean	98.5	
ep_rew_mean	502	

---

time/		
fps	421	
iterations	688	
time_elapsed	3342	
total_timesteps	1409024	
train/		
approx_kl	0.15837029	
clip_fraction	0.427	
clip_range	0.2	
entropy_loss	-15.6	
explained_variance	0.93	
learning_rate	0.0003	
loss	84.3	
n_updates	6870	
policy_gradient_loss	-0.0443	
std	0.618	
value_loss	296	

---

rollout/		
ep_len_mean	99	
ep_rew_mean	503	
time/		
fps	421	
iterations	689	
time_elapsed	3347	
total_timesteps	1411072	
train/		
approx_kl	0.14313918	
clip_fraction	0.445	
clip_range	0.2	
entropy_loss	-15.5	
explained_variance	0.94	
learning_rate	0.0003	
loss	88.7	
n_updates	6880	
policy_gradient_loss	-0.0381	
std	0.617	
value_loss	303	

---

rollout/		
ep_len_mean	98.8	
ep_rew_mean	502	
time/		
fps	421	
iterations	690	
time_elapsed	3352	

	total_timesteps		1413120	
	train/			
	approx_kl		0.09487152	
	clip_fraction		0.397	
	clip_range		0.2	
	entropy_loss		-15.5	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		129	
	n_updates		6890	
	policy_gradient_loss		-0.0506	
	std		0.617	
	value_loss		359	

	rollout/			
	ep_len_mean		97	
	ep_rew_mean		491	
	time/			
	fps		421	
	iterations		691	
	time_elapsed		3357	
	total_timesteps		1415168	
	train/			
	approx_kl		0.098106906	
	clip_fraction		0.375	
	clip_range		0.2	
	entropy_loss		-15.5	
	explained_variance		0.934	
	learning_rate		0.0003	
	loss		134	
	n_updates		6900	
	policy_gradient_loss		-0.055	
	std		0.616	
	value_loss		284	

	rollout/			
	ep_len_mean		95	
	ep_rew_mean		483	
	time/			
	fps		421	
	iterations		692	
	time_elapsed		3362	
	total_timesteps		1417216	
	train/			
	approx_kl		0.11021967	
	clip_fraction		0.383	

clip_range	0.2	
entropy_loss	-15.5	
explained_variance	0.928	
learning_rate	0.0003	
loss	150	
n_updates	6910	
policy_gradient_loss	-0.0562	
std	0.616	
value_loss	346	

rollout/		
ep_len_mean	92.9	
ep_rew_mean	471	
time/		
fps	421	
iterations	693	
time_elapsed	3366	
total_timesteps	1419264	
train/		
approx_kl	0.09188216	
clip_fraction	0.379	
clip_range	0.2	
entropy_loss	-15.5	
explained_variance	0.933	
learning_rate	0.0003	
loss	125	
n_updates	6920	
policy_gradient_loss	-0.0534	
std	0.616	
value_loss	312	

rollout/		
ep_len_mean	93.3	
ep_rew_mean	472	
time/		
fps	421	
iterations	694	
time_elapsed	3371	
total_timesteps	1421312	
train/		
approx_kl	0.12247704	
clip_fraction	0.415	
clip_range	0.2	
entropy_loss	-15.5	
explained_variance	0.944	
learning_rate	0.0003	



	loss		112	
	n_updates		6930	
	policy_gradient_loss		-0.0503	
	std		0.616	
	value_loss		282	

	rollout/			
	ep_len_mean		94.2	
	ep_rew_mean		477	
	time/			
	fps		421	
	iterations		695	
	time_elapsed		3375	
	total_timesteps		1423360	
	train/			
	approx_kl		0.12064403	
	clip_fraction		0.406	
	clip_range		0.2	
	entropy_loss		-15.5	
	explained_variance		0.938	
	learning_rate		0.0003	
	loss		100	
	n_updates		6940	
	policy_gradient_loss		-0.0535	
	std		0.616	
	value_loss		296	

	rollout/			
	ep_len_mean		94.5	
	ep_rew_mean		479	
	time/			
	fps		421	
	iterations		696	
	time_elapsed		3380	
	total_timesteps		1425408	
	train/			
	approx_kl		0.09833868	
	clip_fraction		0.395	
	clip_range		0.2	
	entropy_loss		-15.5	
	explained_variance		0.934	
	learning_rate		0.0003	
	loss		92.3	
	n_updates		6950	
	policy_gradient_loss		-0.0557	
	std		0.615	

	value_loss		308	
-----				
	rollout/			
	ep_len_mean		97.6	
	ep_rew_mean		490	
	time/			
	fps		421	
	iterations		697	
	time_elapsed		3385	
	total_timesteps		1427456	
	train/			
	approx_kl		0.10932177	
	clip_fraction		0.404	
	clip_range		0.2	
	entropy_loss		-15.5	
	explained_variance		0.94	
	learning_rate		0.0003	
	loss		96.5	
	n_updates		6960	
	policy_gradient_loss		-0.0508	
	std		0.615	
	value_loss		303	
-----				
	rollout/			
	ep_len_mean		97.3	
	ep_rew_mean		490	
	time/			
	fps		421	
	iterations		698	
	time_elapsed		3390	
	total_timesteps		1429504	
	train/			
	approx_kl		0.13352333	
	clip_fraction		0.435	
	clip_range		0.2	
	entropy_loss		-15.5	
	explained_variance		0.94	
	learning_rate		0.0003	
	loss		99.1	
	n_updates		6970	
	policy_gradient_loss		-0.0509	
	std		0.616	
	value_loss		285	
-----				
	rollout/			

	ep_len_mean		96.2	
	ep_rew_mean		485	
	time/			
	fps		421	
	iterations		699	
	time_elapsed		3395	
	total_timesteps		1431552	
	train/			
	approx_kl		0.0980304	
	clip_fraction		0.381	
	clip_range		0.2	
	entropy_loss		-15.5	
	explained_variance		0.939	
	learning_rate		0.0003	
	loss		112	
	n_updates		6980	
	policy_gradient_loss		-0.0549	
	std		0.616	
	value_loss		308	

	rollout/			
	ep_len_mean		95.6	
	ep_rew_mean		483	
	time/			
	fps		420	
	iterations		700	
	time_elapsed		3406	
	total_timesteps		1433600	
	train/			
	approx_kl		0.08349624	
	clip_fraction		0.361	
	clip_range		0.2	
	entropy_loss		-15.5	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		186	
	n_updates		6990	
	policy_gradient_loss		-0.0596	
	std		0.615	
	value_loss		382	

	rollout/			
	ep_len_mean		95.3	
	ep_rew_mean		481	
	time/			
	fps		419	

iterations	701
time_elapsed	3419
total_timesteps	1435648
train/	
approx_kl	0.15034527
clip_fraction	0.42
clip_range	0.2
entropy_loss	-15.5
explained_variance	0.935
learning_rate	0.0003
loss	82.9
n_updates	7000
policy_gradient_loss	-0.0532
std	0.614
value_loss	298

rollout/	
ep_len_mean	94.2
ep_rew_mean	477
time/	
fps	419
iterations	702
time_elapsed	3426
total_timesteps	1437696
train/	
approx_kl	0.12414686
clip_fraction	0.424
clip_range	0.2
entropy_loss	-15.4
explained_variance	0.93
learning_rate	0.0003
loss	118
n_updates	7010
policy_gradient_loss	-0.0533
std	0.614
value_loss	280

rollout/	
ep_len_mean	93.3
ep_rew_mean	474
time/	
fps	419
iterations	703
time_elapsed	3431
total_timesteps	1439744
train/	

approx_kl	0.11354199	
clip_fraction	0.41	
clip_range	0.2	
entropy_loss	-15.4	
explained_variance	0.928	
learning_rate	0.0003	
loss	83	
n_updates	7020	
policy_gradient_loss	-0.0538	
std	0.613	
value_loss	289	

rollout/		
ep_len_mean	95.7	
ep_rew_mean	485	
time/		
fps	419	
iterations	704	
time_elapsed	3436	
total_timesteps	1441792	
train/		
approx_kl	0.109243646	
clip_fraction	0.412	
clip_range	0.2	
entropy_loss	-15.4	
explained_variance	0.932	
learning_rate	0.0003	
loss	139	
n_updates	7030	
policy_gradient_loss	-0.0559	
std	0.612	
value_loss	303	

rollout/		
ep_len_mean	94.5	
ep_rew_mean	480	
time/		
fps	419	
iterations	705	
time_elapsed	3440	
total_timesteps	1443840	
train/		
approx_kl	0.5140816	
clip_fraction	0.436	
clip_range	0.2	
entropy_loss	-15.4	

	explained_variance		0.936	
	learning_rate		0.0003	
	loss		134	
	n_updates		7040	
	policy_gradient_loss		-0.0536	
	std		0.612	
	value_loss		292	

	rollout/			
	ep_len_mean		96.7	
	ep_rew_mean		492	
	time/			
	fps		419	
	iterations		706	
	time_elapsed		3445	
	total_timesteps		1445888	
	train/			
	approx_kl		0.10646391	
	clip_fraction		0.387	
	clip_range		0.2	
	entropy_loss		-15.4	
	explained_variance		0.909	
	learning_rate		0.0003	
	loss		109	
	n_updates		7050	
	policy_gradient_loss		-0.0516	
	std		0.612	
	value_loss		343	

	rollout/			
	ep_len_mean		96	
	ep_rew_mean		488	
	time/			
	fps		419	
	iterations		707	
	time_elapsed		3449	
	total_timesteps		1447936	
	train/			
	approx_kl		0.109043375	
	clip_fraction		0.417	
	clip_range		0.2	
	entropy_loss		-15.4	
	explained_variance		0.903	
	learning_rate		0.0003	
	loss		144	
	n_updates		7060	

	policy_gradient_loss		-0.0458	
	std		0.612	
	value_loss		400	

---

	rollout/			
	ep_len_mean		97.9	
	ep_rew_mean		497	
	time/			
	fps		419	
	iterations		708	
	time_elapsed		3453	
	total_timesteps		1449984	
	train/			
	approx_kl		0.11357801	
	clip_fraction		0.425	
	clip_range		0.2	
	entropy_loss		-15.4	
	explained_variance		0.942	
	learning_rate		0.0003	
	loss		59.6	
	n_updates		7070	
	policy_gradient_loss		-0.055	
	std		0.611	
	value_loss		249	

---



---

	rollout/			
	ep_len_mean		97.7	
	ep_rew_mean		494	
	time/			
	fps		419	
	iterations		709	
	time_elapsed		3458	
	total_timesteps		1452032	
	train/			
	approx_kl		0.11430616	
	clip_fraction		0.388	
	clip_range		0.2	
	entropy_loss		-15.4	
	explained_variance		0.936	
	learning_rate		0.0003	
	loss		87.4	
	n_updates		7080	
	policy_gradient_loss		-0.0527	
	std		0.611	
	value_loss		310	

---

-----		
rollout/		
ep_len_mean	99.5	
ep_rew_mean	503	
time/		
fps	419	
iterations	710	
time_elapsed	3463	
total_timesteps	1454080	
train/		
approx_kl	0.113251776	
clip_fraction	0.359	
clip_range	0.2	
entropy_loss	-15.3	
explained_variance	0.917	
learning_rate	0.0003	
loss	120	
n_updates	7090	
policy_gradient_loss	-0.0474	
std	0.61	
value_loss	390	
-----		

-----		
rollout/		
ep_len_mean	96.8	
ep_rew_mean	487	
time/		
fps	419	
iterations	711	
time_elapsed	3467	
total_timesteps	1456128	
train/		
approx_kl	0.14542091	
clip_fraction	0.454	
clip_range	0.2	
entropy_loss	-15.3	
explained_variance	0.93	
learning_rate	0.0003	
loss	84.3	
n_updates	7100	
policy_gradient_loss	-0.0442	
std	0.61	
value_loss	320	
-----		

-----		
rollout/		
ep_len_mean	98.5	
ep_rew_mean	494	



time/		
fps	419	
iterations	712	
time_elapsed	3472	
total_timesteps	1458176	
train/		
approx_kl	0.14066258	
clip_fraction	0.408	
clip_range	0.2	
entropy_loss	-15.3	
explained_variance	0.93	
learning_rate	0.0003	
loss	133	
n_updates	7110	
policy_gradient_loss	-0.0495	
std	0.609	
value_loss	305	

---

rollout/		
ep_len_mean	98.3	
ep_rew_mean	494	
time/		
fps	419	
iterations	713	
time_elapsed	3477	
total_timesteps	1460224	
train/		
approx_kl	0.10138592	
clip_fraction	0.403	
clip_range	0.2	
entropy_loss	-15.3	
explained_variance	0.934	
learning_rate	0.0003	
loss	111	
n_updates	7120	
policy_gradient_loss	-0.0506	
std	0.609	
value_loss	305	

---

rollout/		
ep_len_mean	100	
ep_rew_mean	506	
time/		
fps	420	
iterations	714	
time_elapsed	3481	

	total_timesteps		1462272	
	train/			
	approx_kl		0.14037758	
	clip_fraction		0.445	
	clip_range		0.2	
	entropy_loss		-15.3	
	explained_variance		0.932	
	learning_rate		0.0003	
	loss		87	
	n_updates		7130	
	policy_gradient_loss		-0.0498	
	std		0.609	
	value_loss		297	

	rollout/			
	ep_len_mean		98.1	
	ep_rew_mean		494	
	time/			
	fps		420	
	iterations		715	
	time_elapsed		3485	
	total_timesteps		1464320	
	train/			
	approx_kl		0.1165221	
	clip_fraction		0.43	
	clip_range		0.2	
	entropy_loss		-15.3	
	explained_variance		0.942	
	learning_rate		0.0003	
	loss		108	
	n_updates		7140	
	policy_gradient_loss		-0.0458	
	std		0.608	
	value_loss		314	

	rollout/			
	ep_len_mean		98.6	
	ep_rew_mean		497	
	time/			
	fps		420	
	iterations		716	
	time_elapsed		3489	
	total_timesteps		1466368	
	train/			
	approx_kl		0.106205195	
	clip_fraction		0.388	

clip_range	0.2	
entropy_loss	-15.3	
explained_variance	0.931	
learning_rate	0.0003	
loss	122	
n_updates	7150	
policy_gradient_loss	-0.0595	
std	0.609	
value_loss	338	

rollout/		
ep_len_mean	96.8	
ep_rew_mean	489	
time/		
fps	420	
iterations	717	
time_elapsed	3494	
total_timesteps	1468416	
train/		
approx_kl	0.096550375	
clip_fraction	0.364	
clip_range	0.2	
entropy_loss	-15.3	
explained_variance	0.929	
learning_rate	0.0003	
loss	230	
n_updates	7160	
policy_gradient_loss	-0.0508	
std	0.608	
value_loss	384	

rollout/		
ep_len_mean	98.3	
ep_rew_mean	497	
time/		
fps	420	
iterations	718	
time_elapsed	3498	
total_timesteps	1470464	
train/		
approx_kl	0.086926766	
clip_fraction	0.357	
clip_range	0.2	
entropy_loss	-15.3	
explained_variance	0.912	
learning_rate	0.0003	

	loss		141	
	n_updates		7170	
	policy_gradient_loss		-0.0536	
	std		0.608	
	value_loss		394	

---

	rollout/			
	ep_len_mean		98.7	
	ep_rew_mean		500	
	time/			
	fps		420	
	iterations		719	
	time_elapsed		3503	
	total_timesteps		1472512	
	train/			
	approx_kl		0.10319322	
	clip_fraction		0.384	
	clip_range		0.2	
	entropy_loss		-15.3	
	explained_variance		0.923	
	learning_rate		0.0003	
	loss		113	
	n_updates		7180	
	policy_gradient_loss		-0.0551	
	std		0.608	
	value_loss		370	

---

	rollout/			
	ep_len_mean		96.8	
	ep_rew_mean		491	
	time/			
	fps		420	
	iterations		720	
	time_elapsed		3507	
	total_timesteps		1474560	
	train/			
	approx_kl		0.11657783	
	clip_fraction		0.39	
	clip_range		0.2	
	entropy_loss		-15.3	
	explained_variance		0.933	
	learning_rate		0.0003	
	loss		117	
	n_updates		7190	
	policy_gradient_loss		-0.052	
	std		0.607	

value_loss	328
-----	
rollout/	
ep_len_mean	97.2
ep_rew_mean	495
time/	
fps	420
iterations	721
time_elapsed	3512
total_timesteps	1476608
train/	
approx_kl	0.1136899
clip_fraction	0.434
clip_range	0.2
entropy_loss	-15.3
explained_variance	0.936
learning_rate	0.0003
loss	99.4
n_updates	7200
policy_gradient_loss	-0.0508
std	0.606
value_loss	298
-----	
rollout/	
ep_len_mean	97.1
ep_rew_mean	493
time/	
fps	420
iterations	722
time_elapsed	3516
total_timesteps	1478656
train/	
approx_kl	0.09685972
clip_fraction	0.385
clip_range	0.2
entropy_loss	-15.2
explained_variance	0.923
learning_rate	0.0003
loss	90.7
n_updates	7210
policy_gradient_loss	-0.051
std	0.606
value_loss	345
-----	
rollout/	

	ep_len_mean		96	
	ep_rew_mean		484	
	time/			
	fps		420	
	iterations		723	
	time_elapsed		3521	
	total_timesteps		1480704	
	train/			
	approx_kl		0.1300199	
	clip_fraction		0.44	
	clip_range		0.2	
	entropy_loss		-15.2	
	explained_variance		0.924	
	learning_rate		0.0003	
	loss		90	
	n_updates		7220	
	policy_gradient_loss		-0.0549	
	std		0.606	
	value_loss		308	

	rollout/			
	ep_len_mean		96.3	
	ep_rew_mean		484	
	time/			
	fps		419	
	iterations		724	
	time_elapsed		3530	
	total_timesteps		1482752	
	train/			
	approx_kl		0.10419019	
	clip_fraction		0.412	
	clip_range		0.2	
	entropy_loss		-15.2	
	explained_variance		0.93	
	learning_rate		0.0003	
	loss		102	
	n_updates		7230	
	policy_gradient_loss		-0.0513	
	std		0.606	
	value_loss		310	

	rollout/			
	ep_len_mean		97.8	
	ep_rew_mean		490	
	time/			
	fps		419	

iterations	725
time_elapsed	3538
total_timesteps	1484800
train/	
approx_kl	0.17748393
clip_fraction	0.448
clip_range	0.2
entropy_loss	-15.2
explained_variance	0.939
learning_rate	0.0003
loss	88.4
n_updates	7240
policy_gradient_loss	-0.0478
std	0.605
value_loss	268

rollout/	
ep_len_mean	96.5
ep_rew_mean	483
time/	
fps	419
iterations	726
time_elapsed	3545
total_timesteps	1486848
train/	
approx_kl	0.118045665
clip_fraction	0.402
clip_range	0.2
entropy_loss	-15.2
explained_variance	0.923
learning_rate	0.0003
loss	105
n_updates	7250
policy_gradient_loss	-0.0553
std	0.604
value_loss	363

rollout/	
ep_len_mean	94.6
ep_rew_mean	474
time/	
fps	419
iterations	727
time_elapsed	3550
total_timesteps	1488896
train/	

approx_kl	0.1266432	
clip_fraction	0.391	
clip_range	0.2	
entropy_loss	-15.2	
explained_variance	0.934	
learning_rate	0.0003	
loss	102	
n_updates	7260	
policy_gradient_loss	-0.0573	
std	0.604	
value_loss	287	

rollout/		
ep_len_mean	98.6	
ep_rew_mean	495	
time/		
fps	419	
iterations	728	
time_elapsed	3554	
total_timesteps	1490944	
train/		
approx_kl	0.12413794	
clip_fraction	0.412	
clip_range	0.2	
entropy_loss	-15.2	
explained_variance	0.937	
learning_rate	0.0003	
loss	111	
n_updates	7270	
policy_gradient_loss	-0.0526	
std	0.603	
value_loss	290	

rollout/		
ep_len_mean	98.9	
ep_rew_mean	499	
time/		
fps	419	
iterations	729	
time_elapsed	3559	
total_timesteps	1492992	
train/		
approx_kl	0.12571526	
clip_fraction	0.426	
clip_range	0.2	
entropy_loss	-15.1	



	explained_variance	0.932	
	learning_rate	0.0003	
	loss	139	
	n_updates	7280	
	policy_gradient_loss	-0.056	
	std	0.603	
	value_loss	305	

	rollout/		
	ep_len_mean	98.3	
	ep_rew_mean	497	
	time/		
	fps	419	
	iterations	730	
	time_elapsed	3563	
	total_timesteps	1495040	
	train/		
	approx_kl	0.08952077	
	clip_fraction	0.369	
	clip_range	0.2	
	entropy_loss	-15.1	
	explained_variance	0.921	
	learning_rate	0.0003	
	loss	110	
	n_updates	7290	
	policy_gradient_loss	-0.0578	
	std	0.602	
	value_loss	367	

	rollout/		
	ep_len_mean	99.9	
	ep_rew_mean	508	
	time/		
	fps	419	
	iterations	731	
	time_elapsed	3567	
	total_timesteps	1497088	
	train/		
	approx_kl	0.12650272	
	clip_fraction	0.41	
	clip_range	0.2	
	entropy_loss	-15.1	
	explained_variance	0.944	
	learning_rate	0.0003	
	loss	115	
	n_updates	7300	

	policy_gradient_loss		-0.0545	
	std		0.603	
	value_loss		297	

	rollout/			
	ep_len_mean		99.8	
	ep_rew_mean		508	
	time/			
	fps		419	
	iterations		732	
	time_elapsed		3572	
	total_timesteps		1499136	
	train/			
	approx_kl		0.16470687	
	clip_fraction		0.434	
	clip_range		0.2	
	entropy_loss		-15.1	
	explained_variance		0.945	
	learning_rate		0.0003	
	loss		70.6	
	n_updates		7310	
	policy_gradient_loss		-0.047	
	std		0.602	
	value_loss		251	

	rollout/			
	ep_len_mean		98.5	
	ep_rew_mean		502	
	time/			
	fps		419	
	iterations		733	
	time_elapsed		3576	
	total_timesteps		1501184	
	train/			
	approx_kl		0.13467768	
	clip_fraction		0.419	
	clip_range		0.2	
	entropy_loss		-15.1	
	explained_variance		0.931	
	learning_rate		0.0003	
	loss		126	
	n_updates		7320	
	policy_gradient_loss		-0.0501	
	std		0.602	
	value_loss		315	

## 1.4 Evaluate the Reinforcement Learning Performance

```
[4]: def evaluate_ppo(model, num_episodes=100, max_steps=1000):  
    env = gym.make("Humanoid-v5") #creates the humanoid environment  
    rewards = [] #stores total reward for each PPO attempt  
    lengths = [] #store # steps the humanoid survived  
  
    for ep in range(num_episodes): #loops through the PPO agent for as many as  
        ↪we set it to  
        obs, info = env.reset() #starts the environment from the beginning each  
        ↪time  
        total_reward = 0.0 #keeps track of the points earned each episode  
        length = 0 #counts the amount of steps the episode lasts  
  
        for t in range(max_steps): #loops through each episode step  
            action, _ = model.predict(obs, deterministic=True) #the trained  
            ↪ppo model pics the action  
            result = env.step(action) #stores the information for what happens  
            ↪each step  
            obs = result[0] #the state of the humanoid after doing the action  
            reward = result[1] #the amount of points the action earned  
            terminated = result[2] #true if the humanoid dint fal  
            truncated = result[3] #true if the episode hit the max steps  
            info = result[4] #stores extra info  
            total_reward += reward  
            length += 1  
  
            if terminated or truncated: #stops if we hit a time limit or the  
            ↪humanoid fails  
                break  
  
        rewards.append(total_reward) #stores the episode reward  
        lengths.append(length) #stores teh timelength of the episode  
  
    env.close() #closes the humanoid environment  
  
    return rewards, lengths #returns these values for plotting
```

```
[5]: import matplotlib.pyplot as plt  
  
def plot_rewards(rewards):  
    plt.figure(figsize=(10, 5))  
    plt.plot(rewards, color = 'green') #plots the rewards across all the  
    ↪episodes  
    plt.title('PPO Total Rewards per Episode')  
    plt.xlabel('Episode Number')  
    plt.ylabel('Total Reward')
```

```
plt.grid(True)


plt.savefig(r"C:\Users\Bri\Pictures\Humanoid_1,500,000.png", dpi=300)
↳ #saves graph to computer

plt.show()
```

```
[6]: model = PP0.load("ppo_humanoid_v5", env=vec_env)  #loads the trained ppo model

ppo_rewards, ppo_lengths = evaluate_ppo(model)  #runs the ppo model for the
↳ amount of episodes we set

plot_rewards(ppo_rewards)  #makes a plot of the rewards across episodes
```



## 1.5 Saves the Humanoid Video

```
[7]: def record_ppo_video(model, filename="humanoid_ppo.gif",
                             max_steps=2000, episode_horizon=1000):
    env = gym.make("Humanoid-v5", render_mode="rgb_array")  #makes the humanoid
    ↳ environment for RGB frames
    obs, info = env.reset()
    frames = []  #stores the video frames
    steps = 0

    while steps < max_steps:
```

```

    action, _ = model.predict(obs, deterministic=True) #these steps move
    ↪ through the models actions
    obs, reward, terminated, truncated, info = env.step(action)

    frame = env.render() #gets each frame at a time and stores it
    frames.append(frame)

    steps += 1
    if terminated or truncated: #if the humanoid fails it stops at this
    ↪ motion or frame
        obs, info = env.reset()

    env.close()
    print(f"Saving {len(frames)} frames to {filename} ...")
    imageio.mimsave(filename, frames, fps=30) #saves the environment as a GIF

model = PPO.load("ppo_humanoid_v5", env=vec_env) #takes the saved model and
    ↪ loads to memory
filename=r"C:\Users\Bri\Videos\humanoid_ppo_long.gif"
record_ppo_video(model, r"C:\Users\Bri\Videos\humanoid_ppo_long.gif") #saves
    ↪ the video to my computer as a GIF

```

Saving 2000 frames to C:\Users\Bri\Videos\humanoid\_ppo\_long.gif ...